

# **Neural Networks**

## **Learning the network: Part 1**

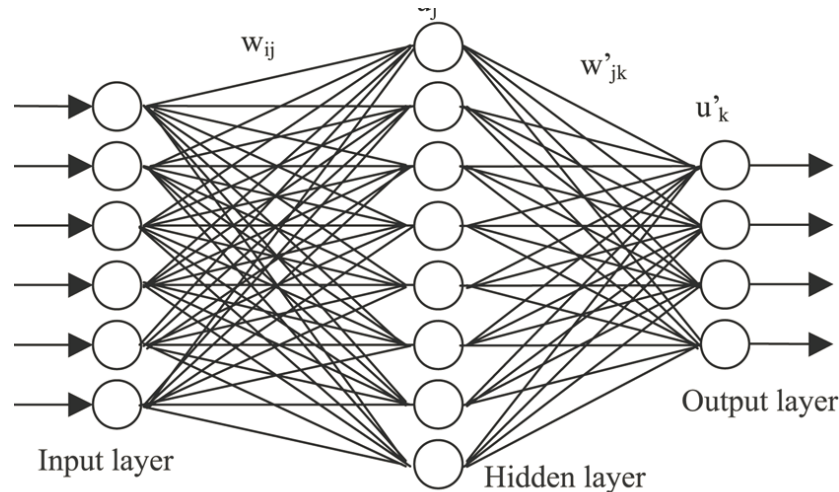
11-785, Spring 2021

Lecture 3

# Topics for the day

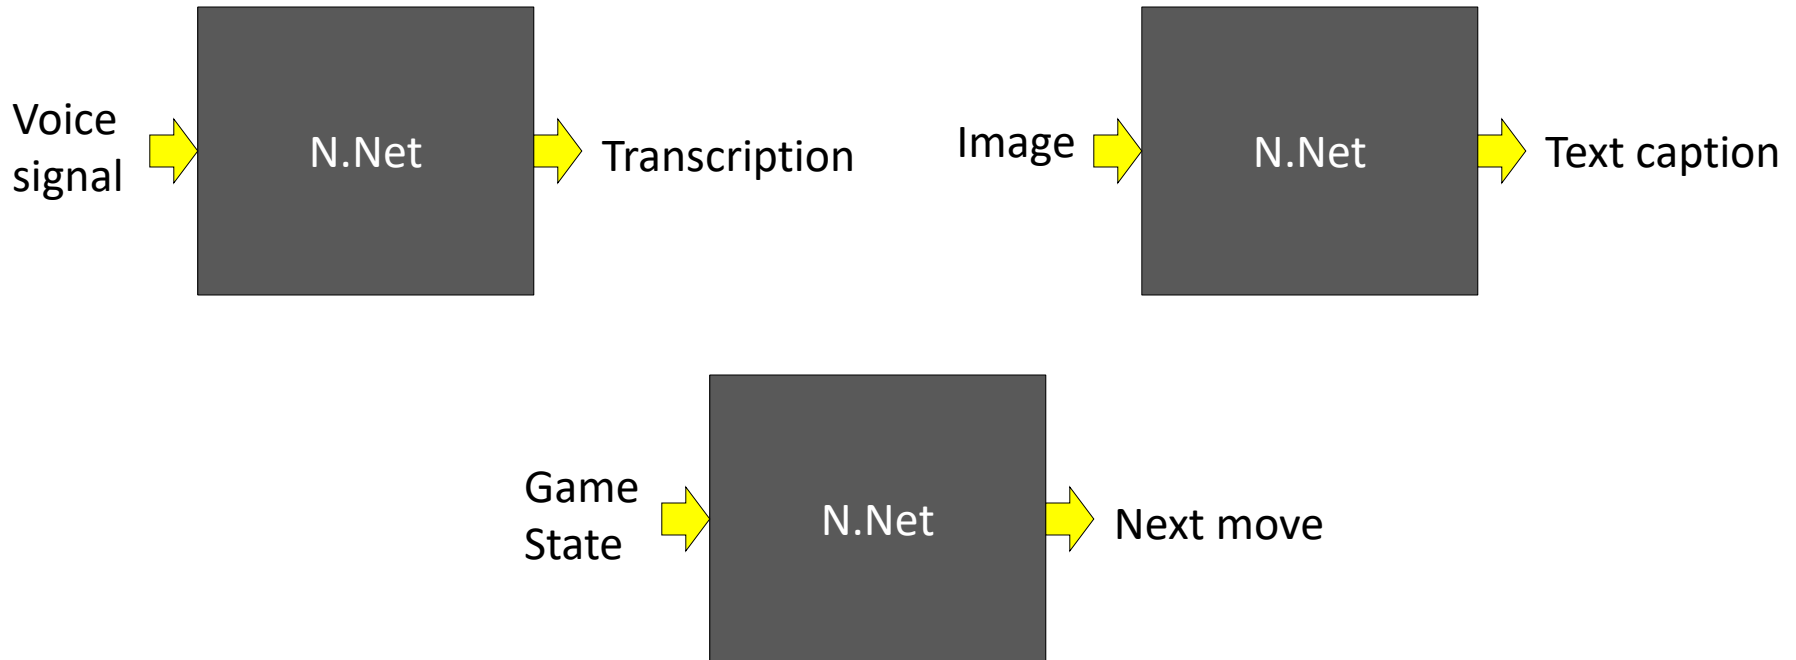
- The problem of learning
- The perceptron rule for perceptrons
  - And its inapplicability to multi-layer perceptrons
- Greedy solutions for classification networks: ADALINE and MADALINE
- Learning through Empirical Risk Minimization
- Intro to function optimization and gradient descent

# Recap



- **Neural networks are universal function approximators**
  - Can model any Boolean function
  - Can model any classification boundary
  - Can model any continuous valued function
- *Provided the network satisfies minimal architecture constraints*
  - Networks with fewer than the required number of parameters can be very poor approximators

# These boxes are functions



- Take an input
- Produce an output
- Can be modeled by a neural network!

# Questions



- Preliminaries:
  - How do we represent the input?
  - How do we represent the output?
- How do we compose the network that performs the requisite function?

# Questions



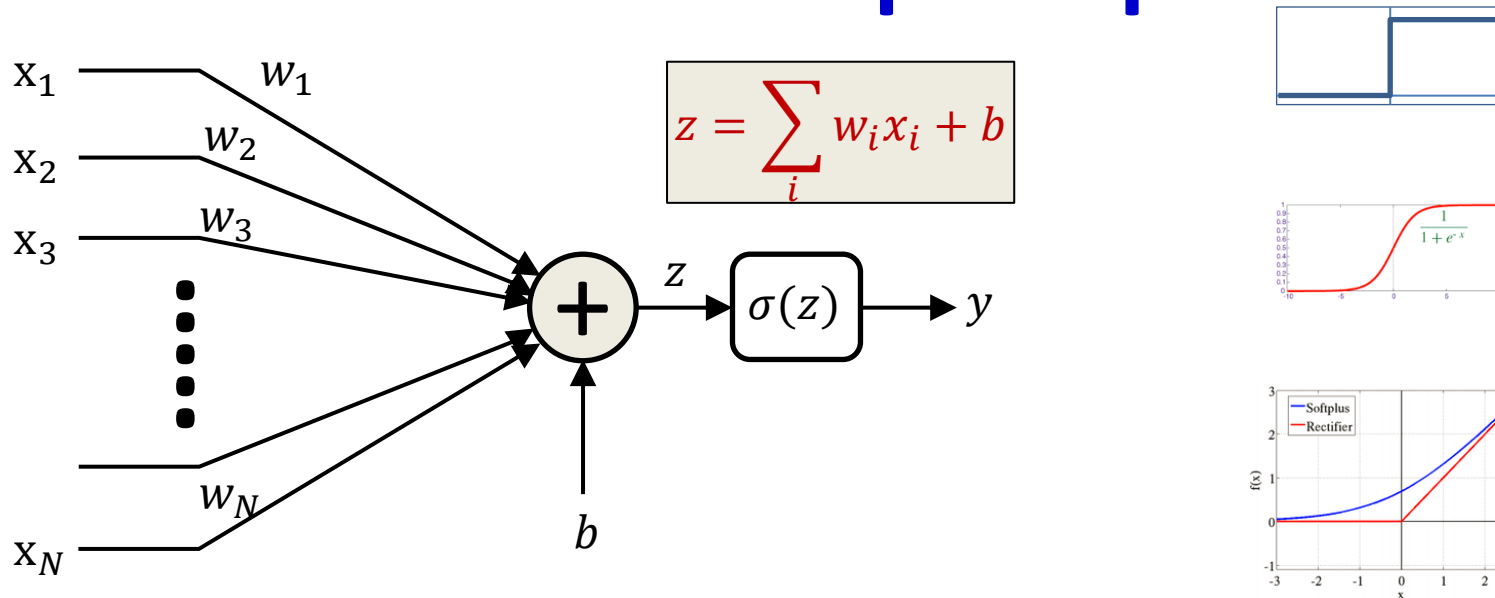
- Preliminaries:

- How do we represent the input?
- How do we represent the output?

A bit later in the program

- *How do we compose the network that performs the requisite function?* ←

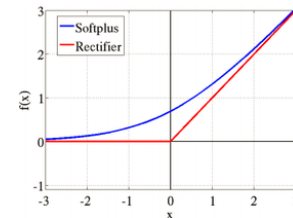
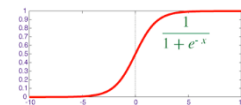
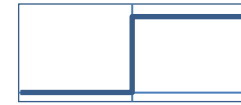
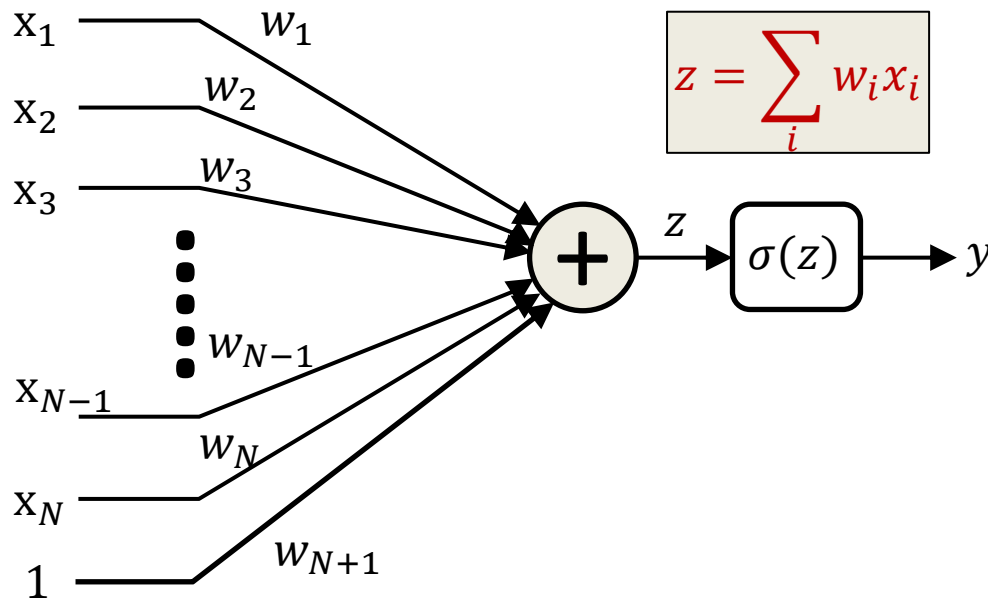
# Preliminaries: The units in the network – the perceptron



Activation functions  $\sigma(z)$

- Perceptron
  - General setting, inputs are real valued
  - A *bias*  $b$  representing a threshold to trigger the perceptron
  - Activation functions are not necessarily threshold functions
- The parameters of the perceptron (which determine how it behaves) are its weights and bias

# Preliminaries: Redrawing the neuron

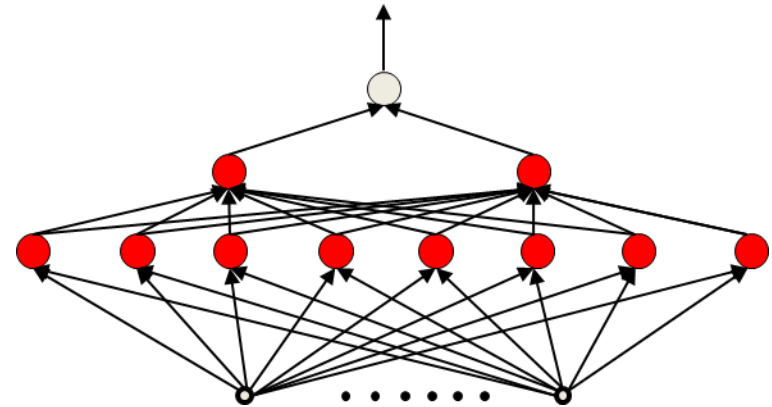
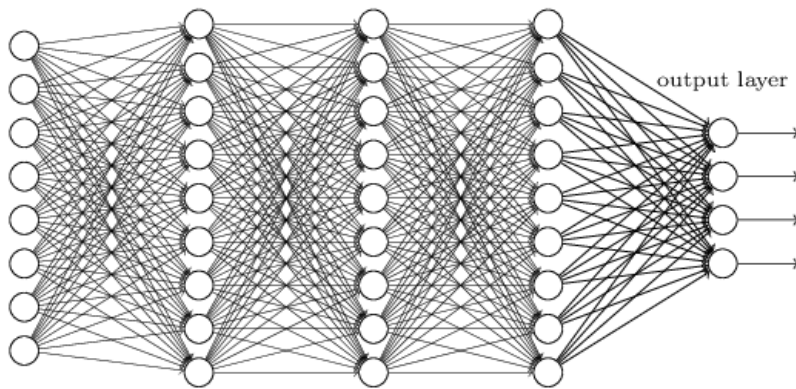


*Activation functions  $\sigma(z)$*

- The bias can also be viewed as the weight of another input component that is always set to 1
  - If the bias is not explicitly mentioned, we will implicitly be assuming that every perceptron has an additional input that is always fixed at 1

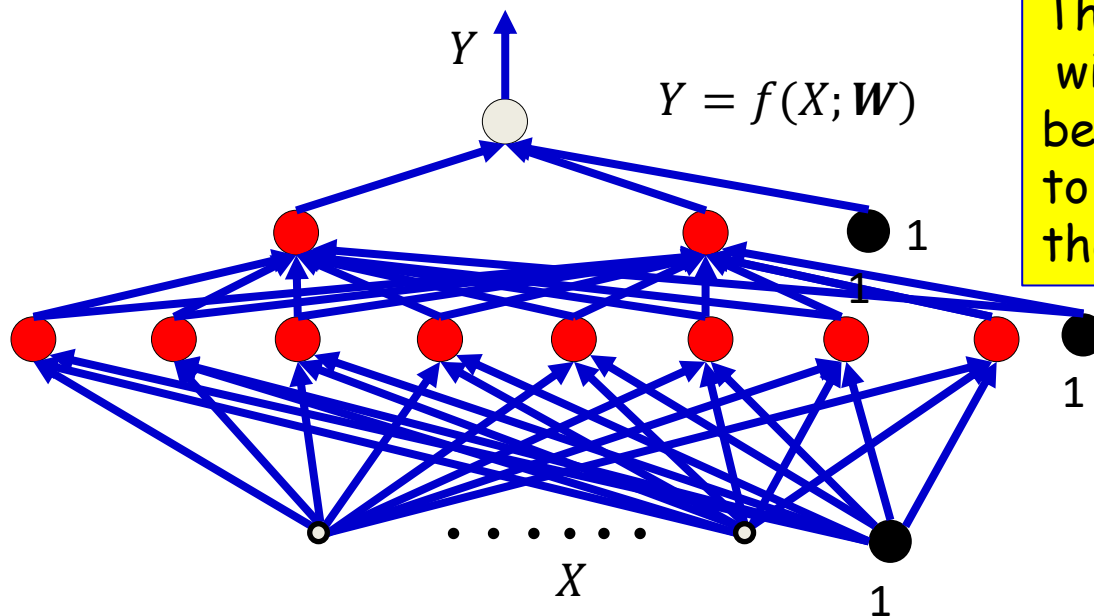


# First: the structure of the network



- We will assume a *feed-forward* network
  - No loops: Neuron outputs do not feed back to their inputs directly or indirectly
  - Loopy networks are a future topic
- **Part of the design of a network: The architecture**
  - How many layers/neurons, which neuron connects to which and how, etc.
- For now, assume the architecture of the network is capable of representing the needed function

# What we learn: The parameters of the network

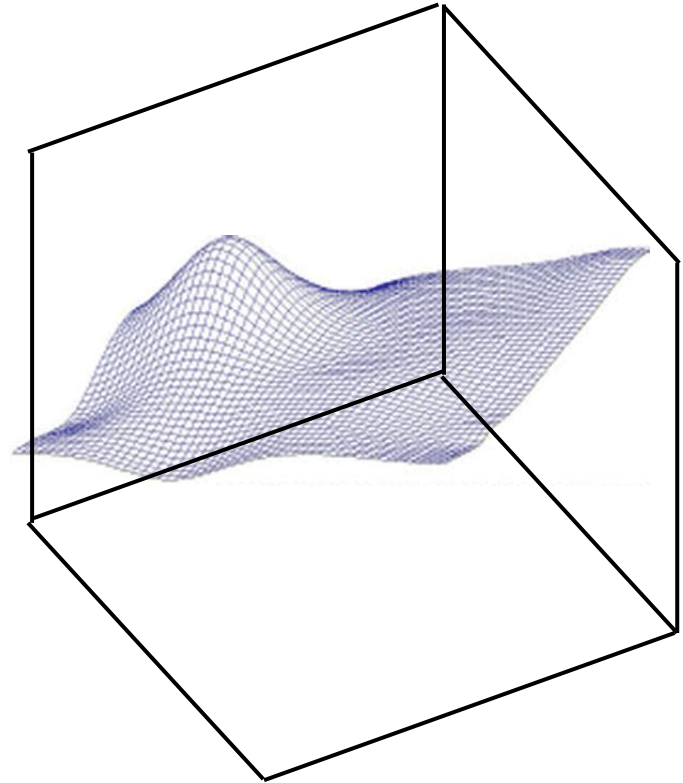
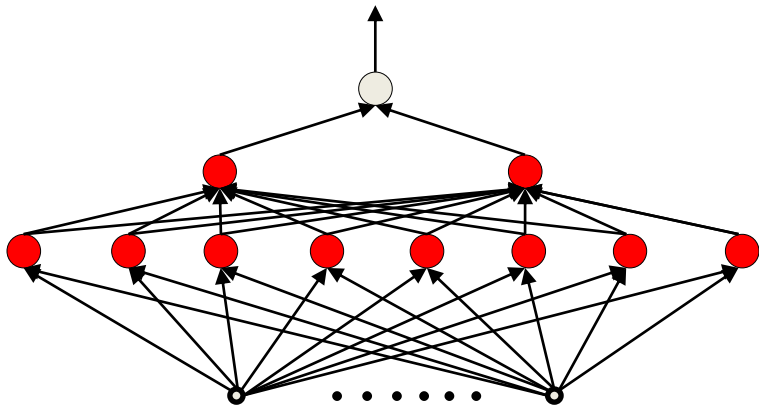


The network is a function  $f()$  with parameters  $W$  which must be set to the appropriate values to get the desired behavior from the net

- **Given:** the architecture of the network
- **The parameters of the network:** The weights and biases
  - The weights associated with the blue arrows in the picture
- **Learning the network :** Determining the values of these parameters such that the network computes the desired function

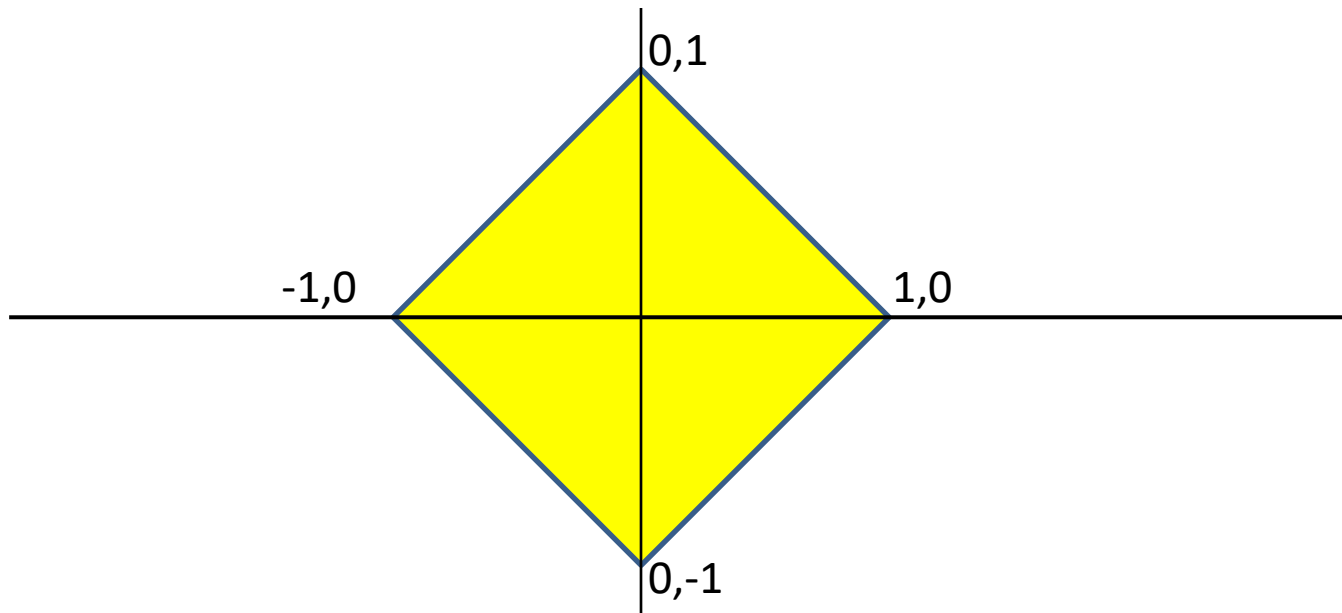
- Moving on..

# The MLP *can* represent anything



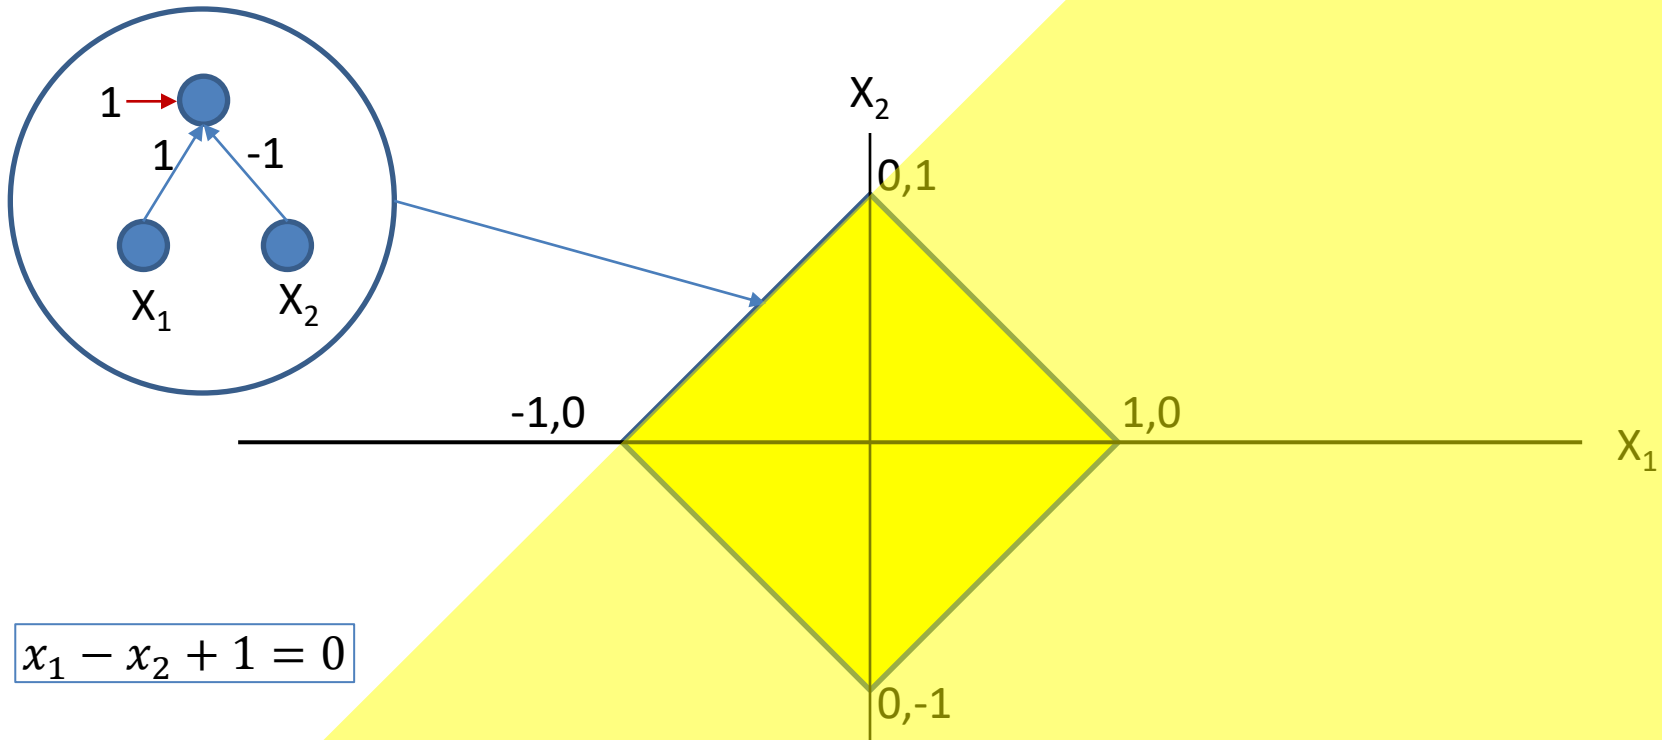
- The MLP *can be constructed* to represent anything
- But *how* do we construct it?

# Option 1: Construct by hand



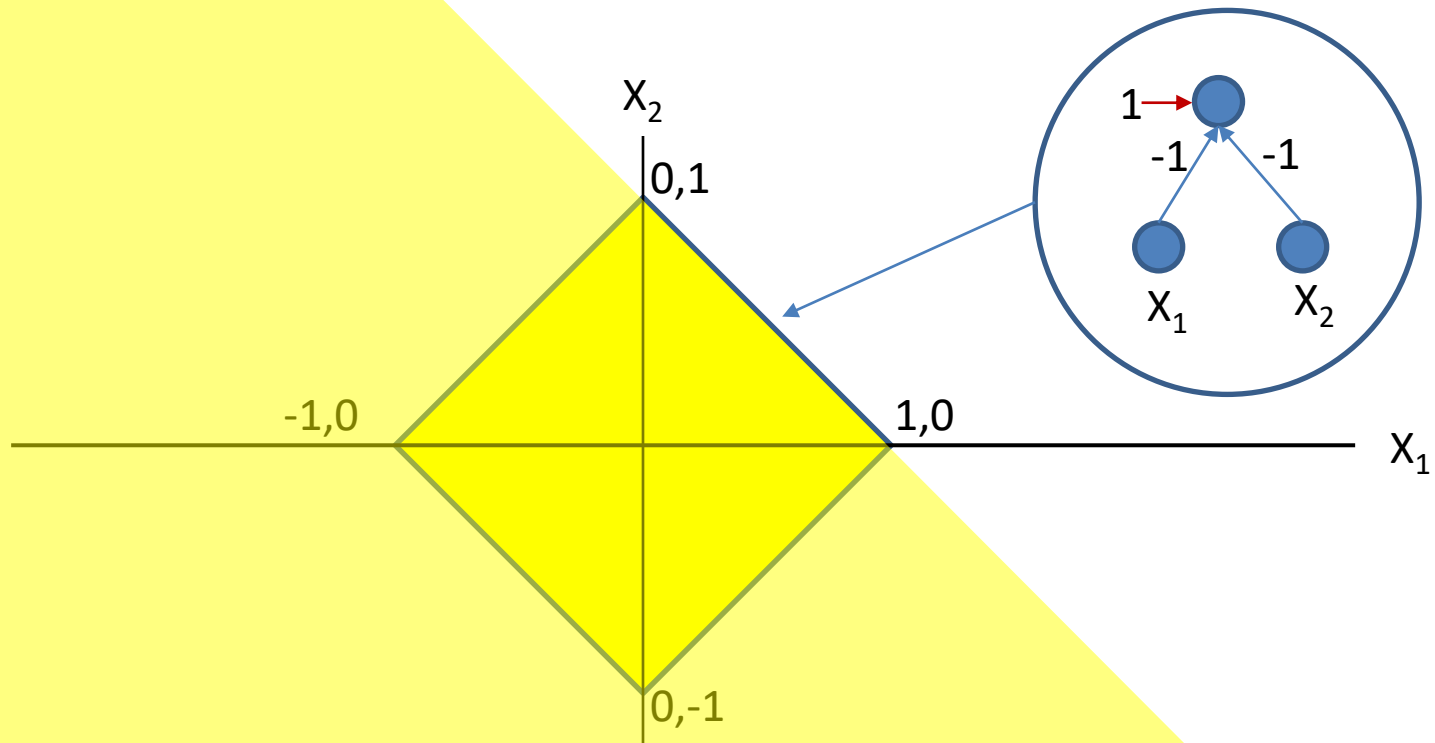
- Given a function, *handcraft* a network to satisfy it
- E.g.: Build an MLP to classify this decision boundary

# Option 1: Construct by hand



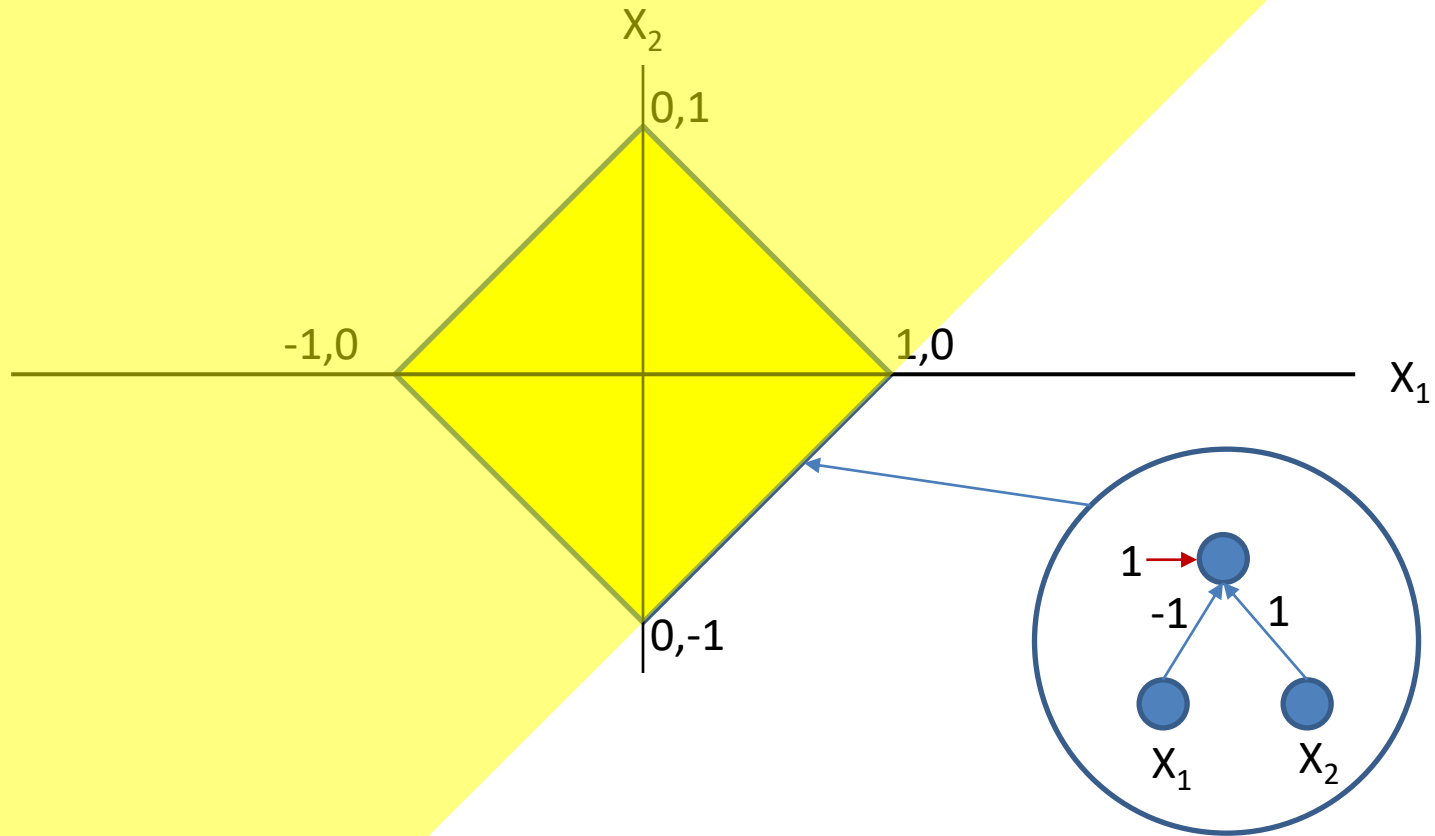
Assuming simple perceptrons:  
output = 1 if  $\sum_i w_i x_i + b_i \geq 0$ , else 0

# Option 1: Construct by hand



Assuming simple perceptrons:  
output = 1 if  $\sum_i w_i x_i + b_i \geq 0$ , else 0

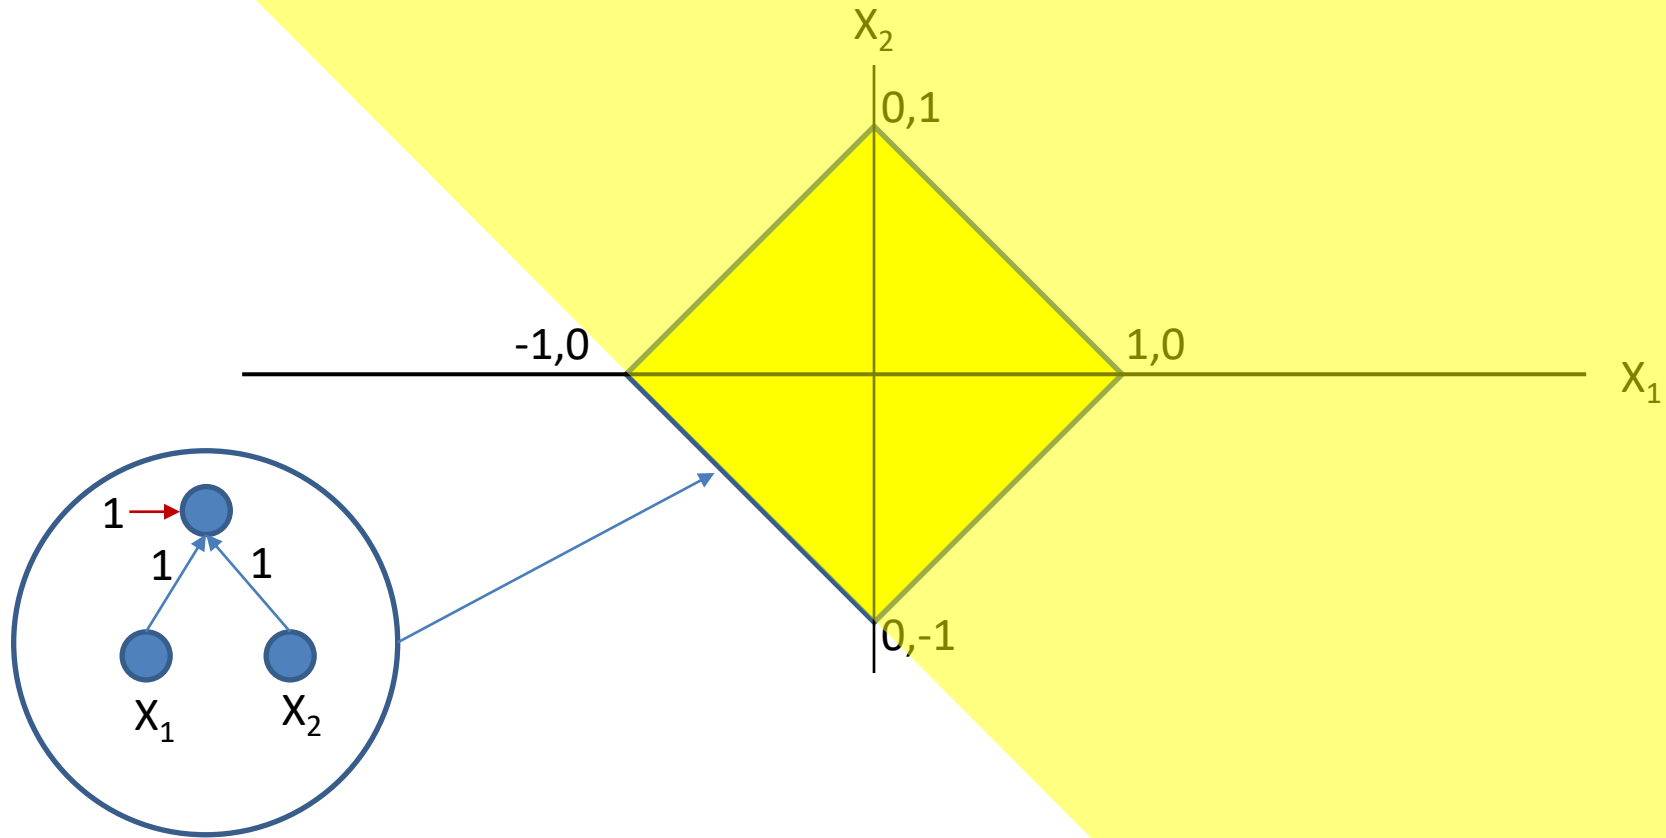
# Option 1: Construct by hand



Assuming simple perceptrons:  
output = 1 if  $\sum_i w_i x_i + b_i \geq 0$ , else 0

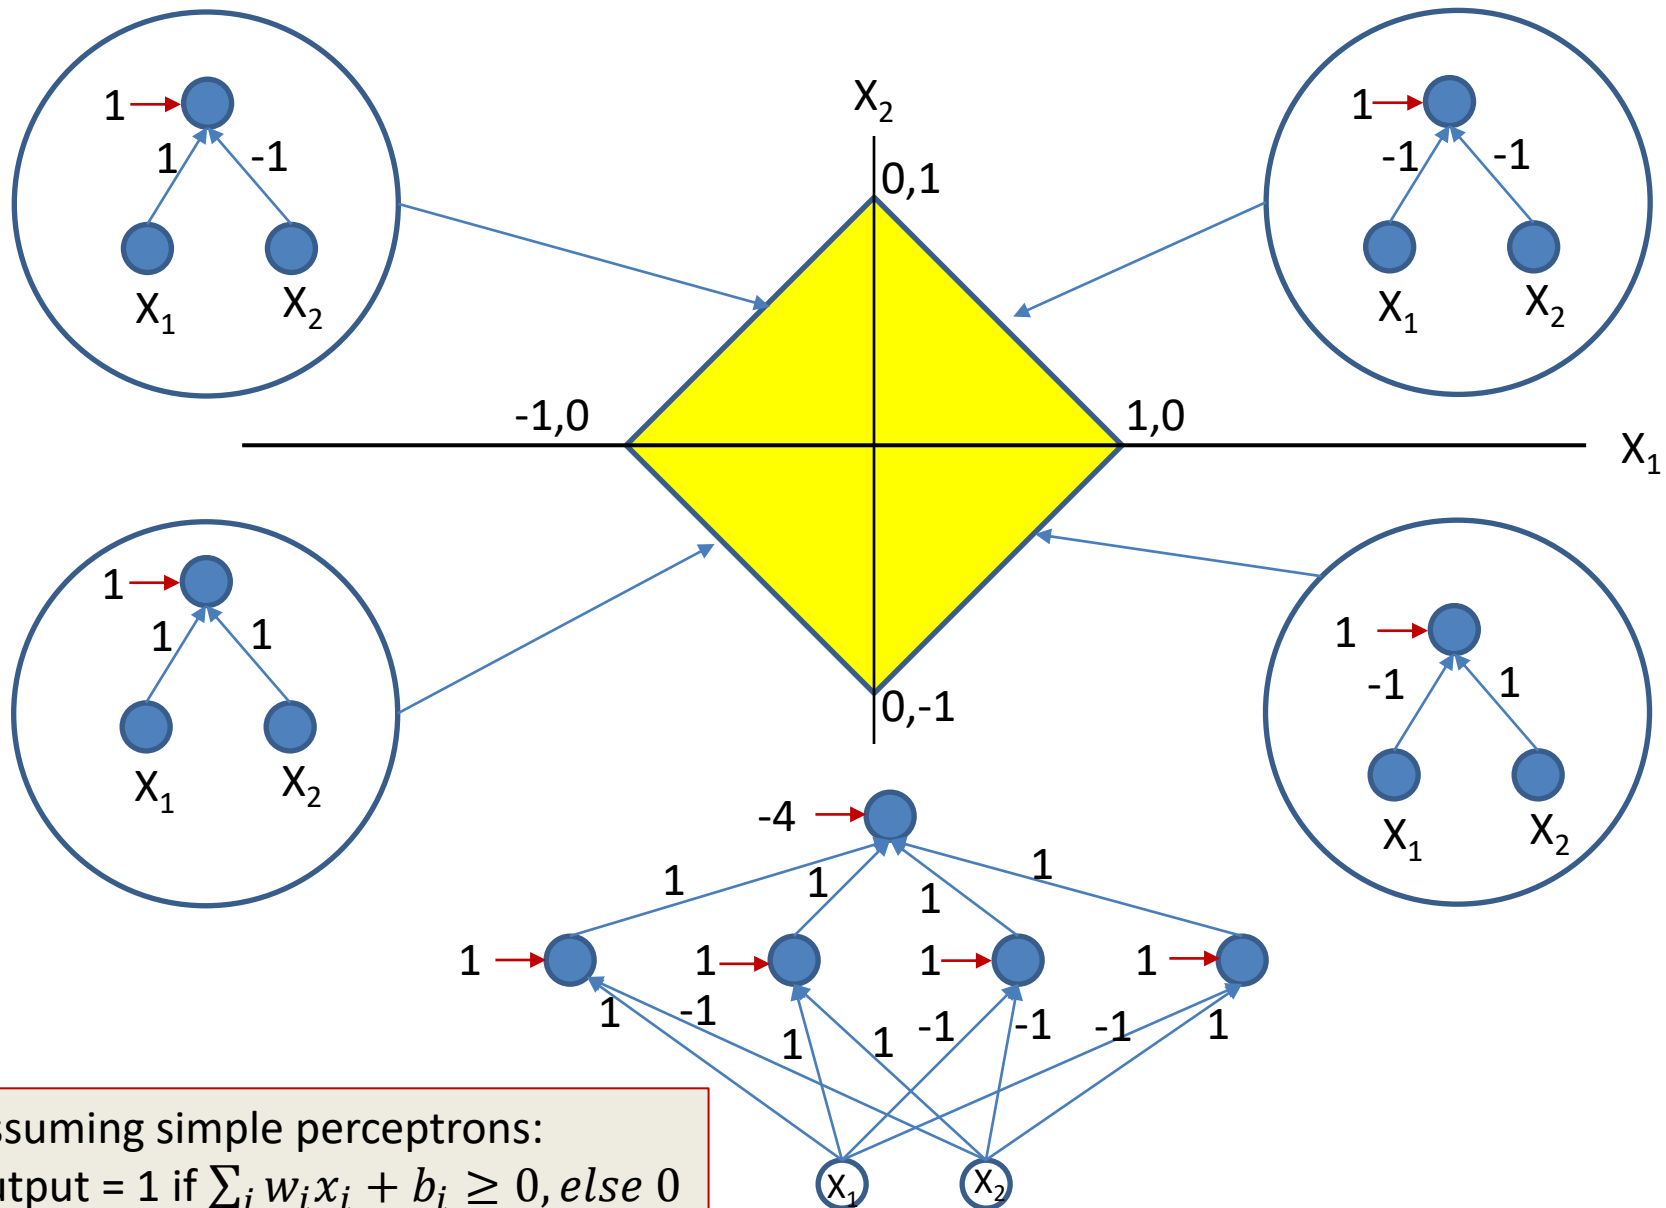


# Option 1: Construct by hand

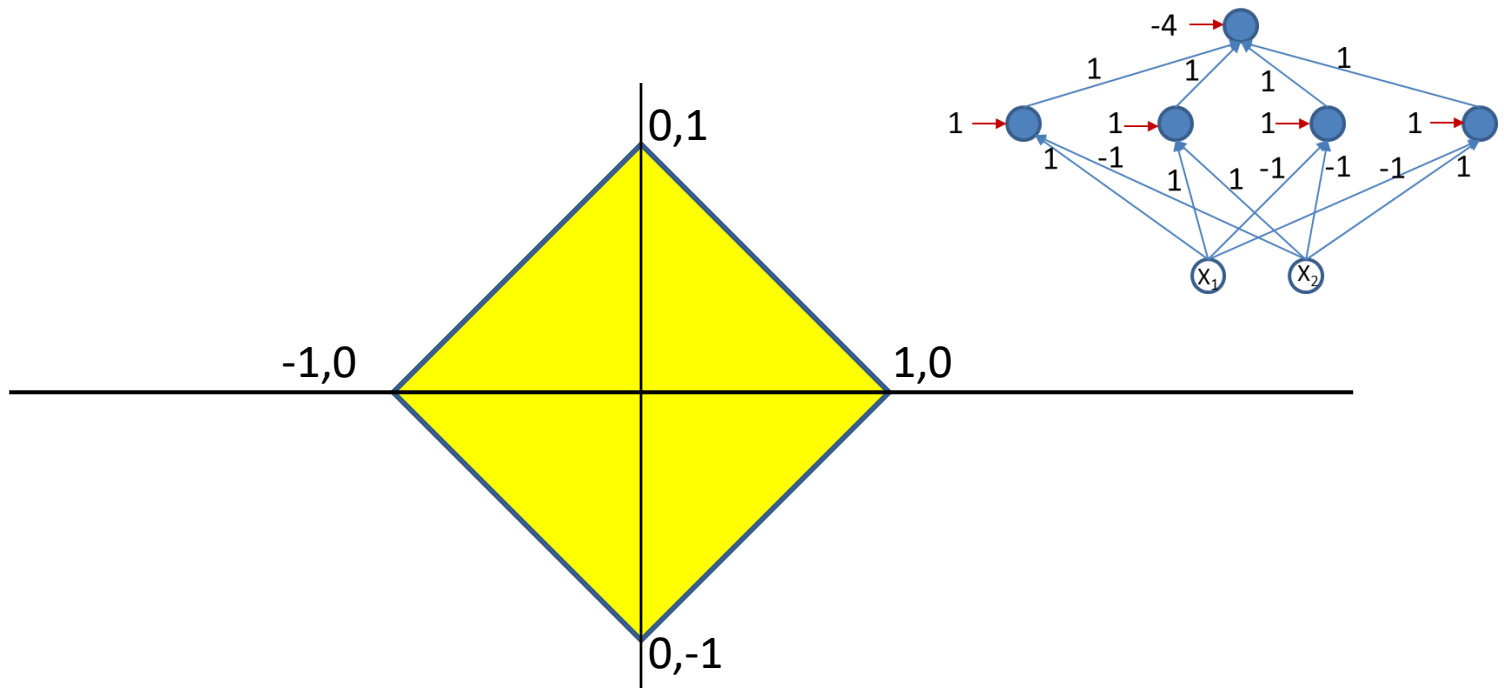


Assuming simple perceptrons:  
output = 1 if  $\sum_i w_i x_i + b_i \geq 0$ , else 0

# Option 1: Construct by hand

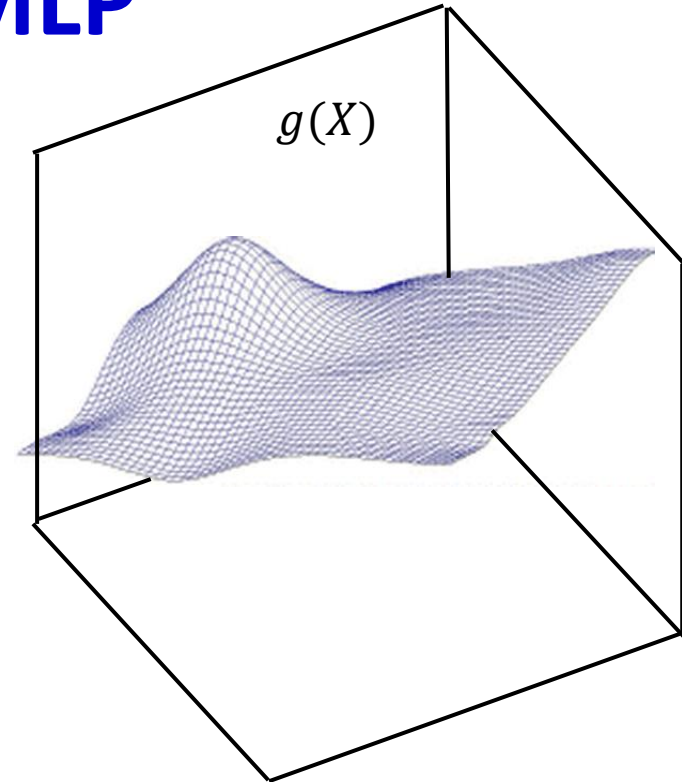
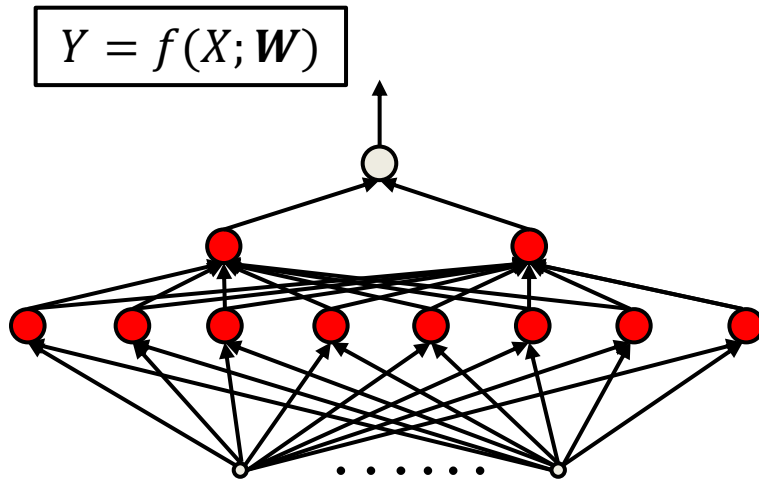


# Option 1: Construct by hand



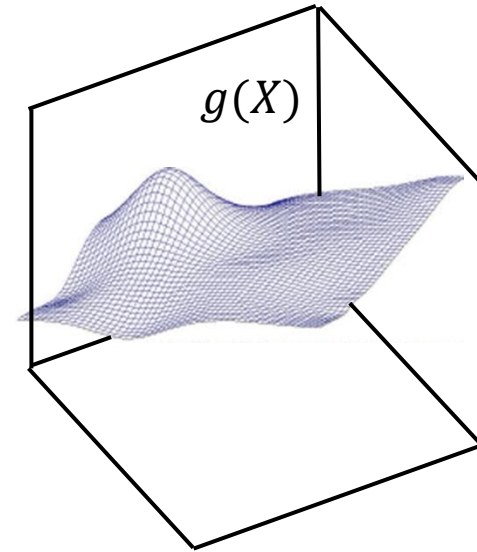
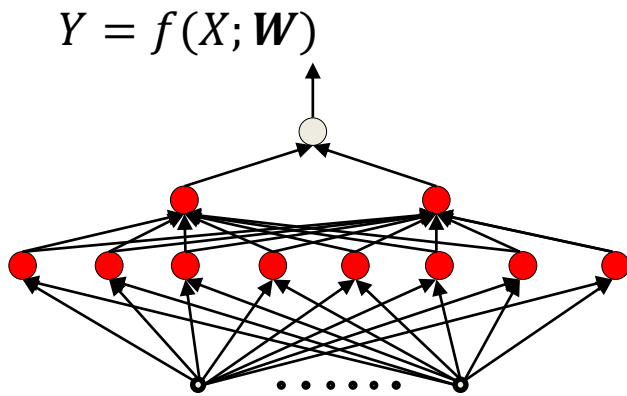
- Given a function, *handcraft* a network to satisfy it
- E.g.: Build an MLP to classify this decision boundary
- Not possible for all but the simplest problems..

## Option 2: Automatic estimation of an MLP



- More generally, *given* the function  $g(X)$  to model, we can *derive* the parameters of the network to model it, through computation

# How to learn a network?

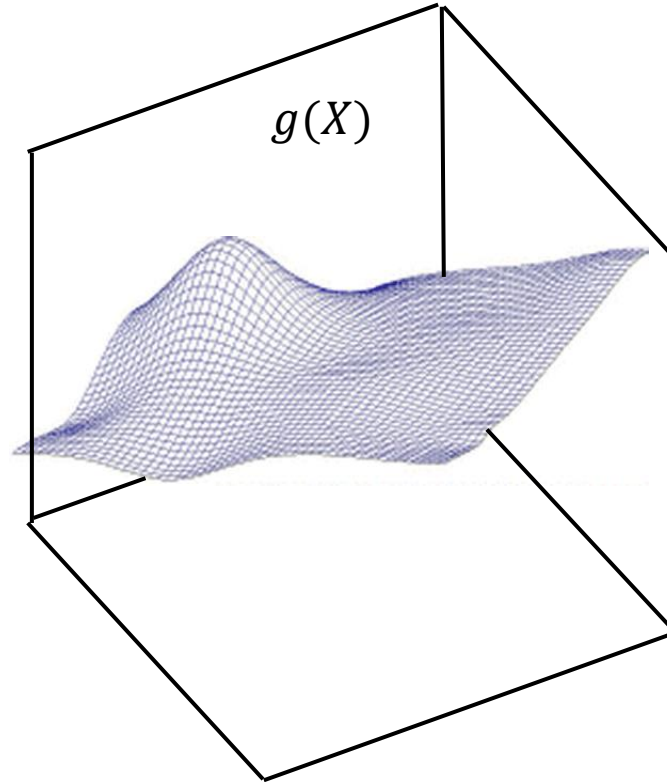


- When  $f(X; \mathbf{W})$  has the capacity to exactly represent  $g(X)$

$$\widehat{\mathbf{W}} = \operatorname{argmin}_{\mathbf{W}} \int_X \operatorname{div}(f(X; \mathbf{W}), g(X)) dX$$

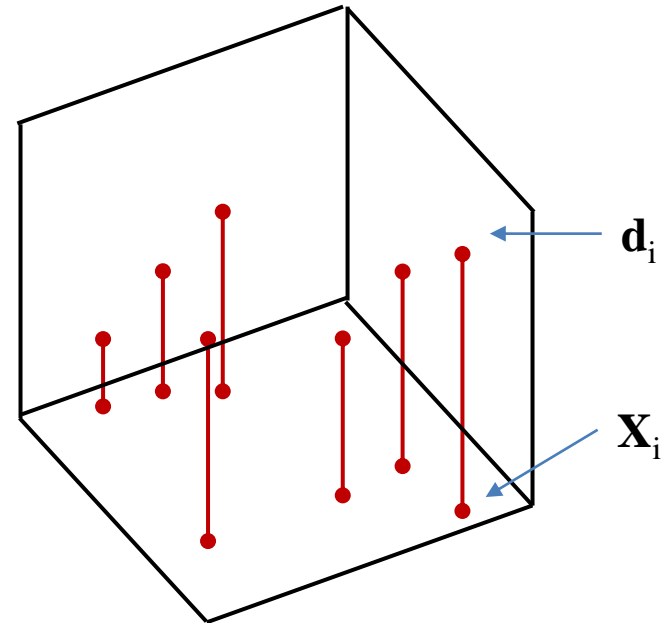
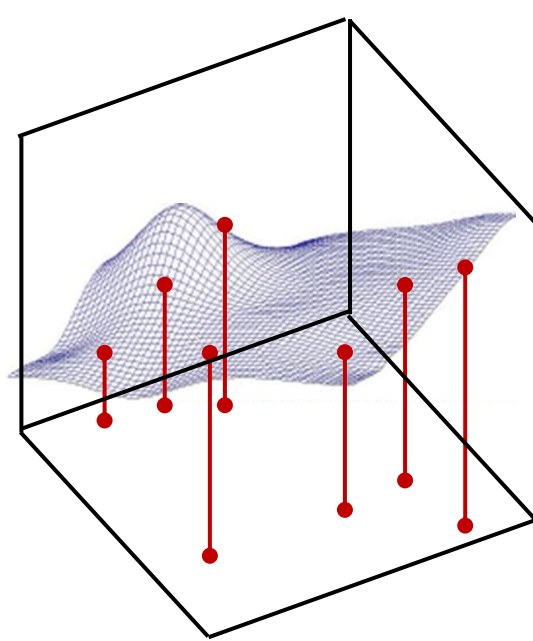
- $\operatorname{div}()$  is a *divergence* function that goes to zero when  $f(X; \mathbf{W}) = g(X)$

# Problem $g(X)$ is unknown



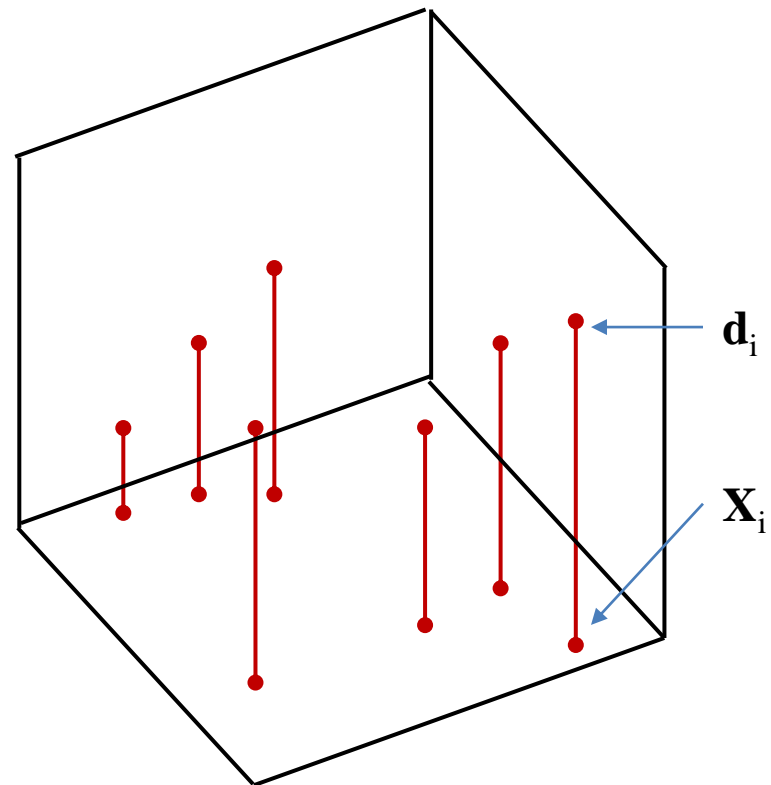
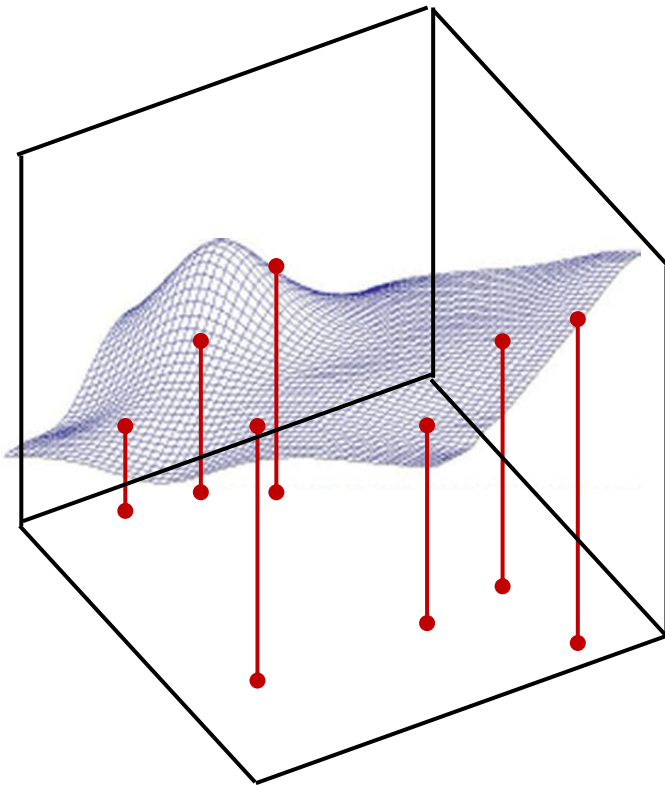
- Function  $g(X)$  must be fully specified
  - Known *everywhere*, i.e. for *every* input  $X$
- **In practice we will not have such specification**

# Sampling the function



- *Sample  $g(X)$* 
  - Basically, get input-output pairs for a number of samples of input  $X_i$ 
    - Many samples  $(X_i, d_i)$ , where  $d_i = g(X_i) + noise$
- Very easy to do in most problems: just gather training data
  - E.g. set of images and their class labels
  - E.g. speech recordings and their transcription

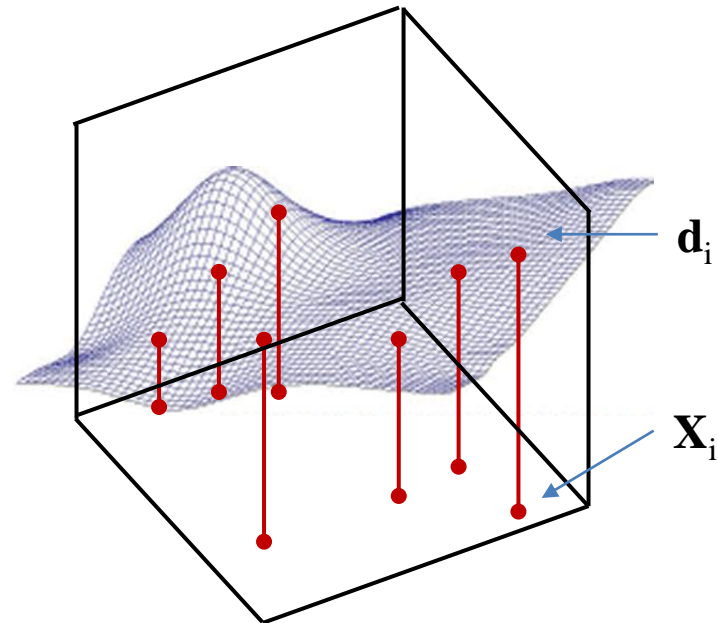
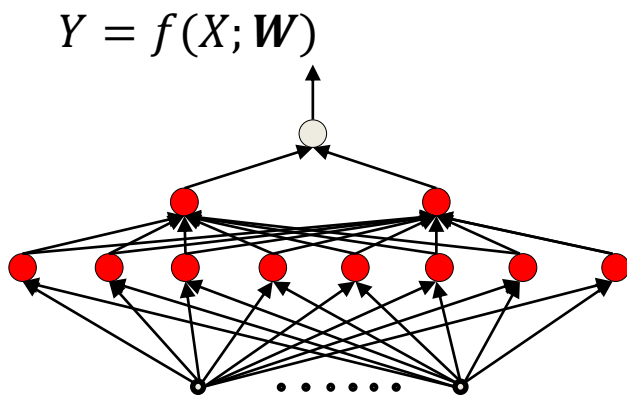
# *Drawing samples*



- We must *learn* the *entire* function from these few examples
  - The “training” samples



# Learning the function



- Estimate the network parameters to “fit” the training points exactly
  - Assuming network architecture is sufficient for such a fit
  - Assuming unique output  $d$  at any  $\mathbf{X}$ 
    - And hopefully the resulting function is also correct where we *don't* have training samples

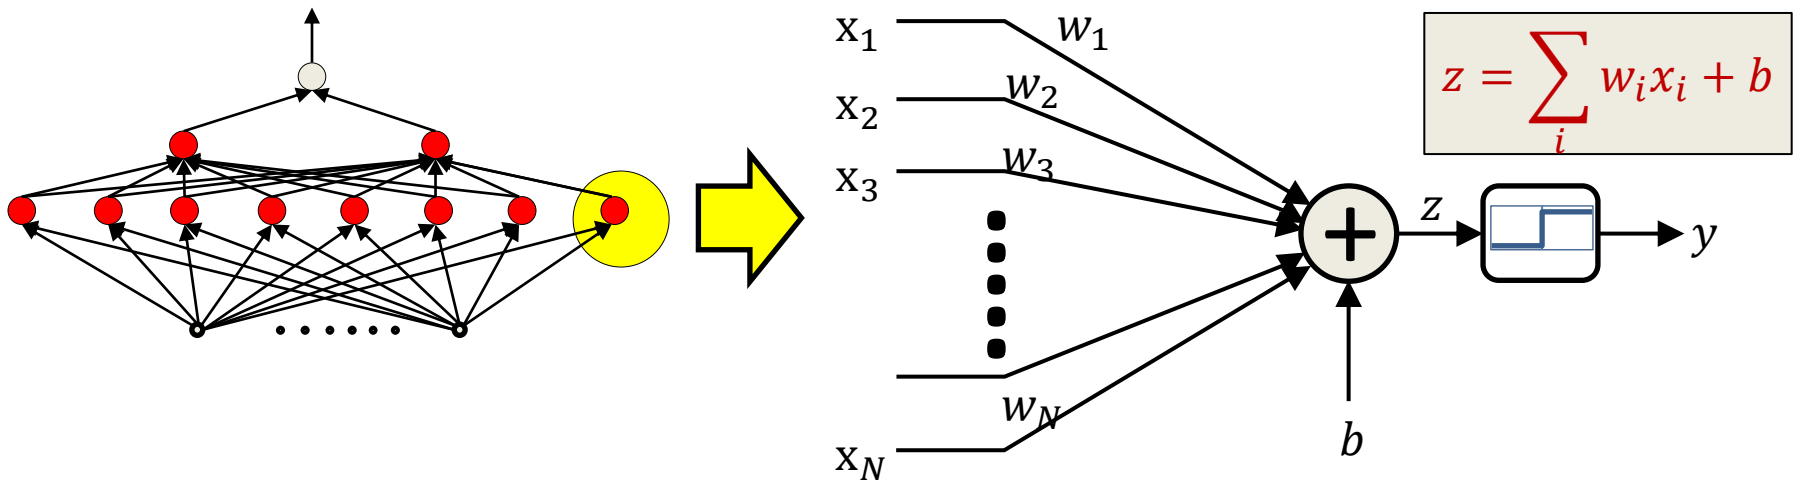
# Story so far

- “Learning” a neural network == determining the parameters of the network (weights and biases) required for it to model a desired function
  - The network must have sufficient capacity to model the function
- Ideally, we would like to optimize the network to represent the desired function everywhere
- However this requires knowledge of the function everywhere
- Instead, we draw “input-output” *training* instances from the function and estimate network parameters to “fit” the input-output relation at these instances
  - And hope it fits the function elsewhere as well

# Let's begin with a simple task

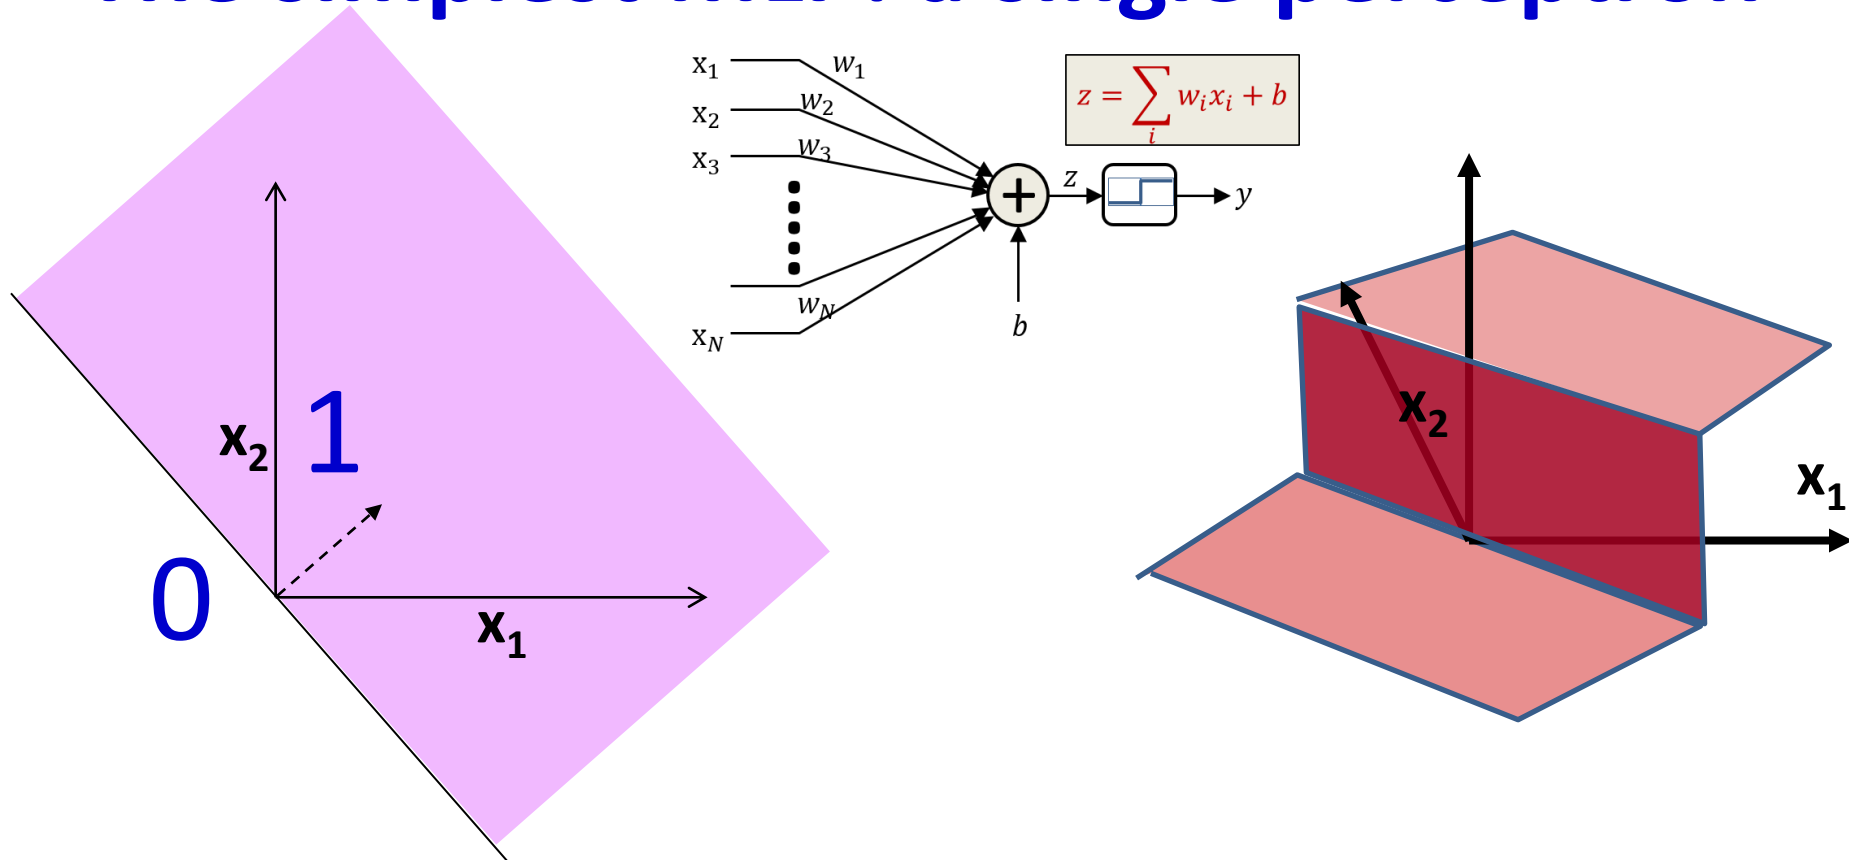
- Learning a *classifier*
  - Simpler than regressions
- This was among the earliest problems addressed using MLPs
- Specifically, consider *binary* classification
  - Generalizes to multi-class

# History: The original MLP



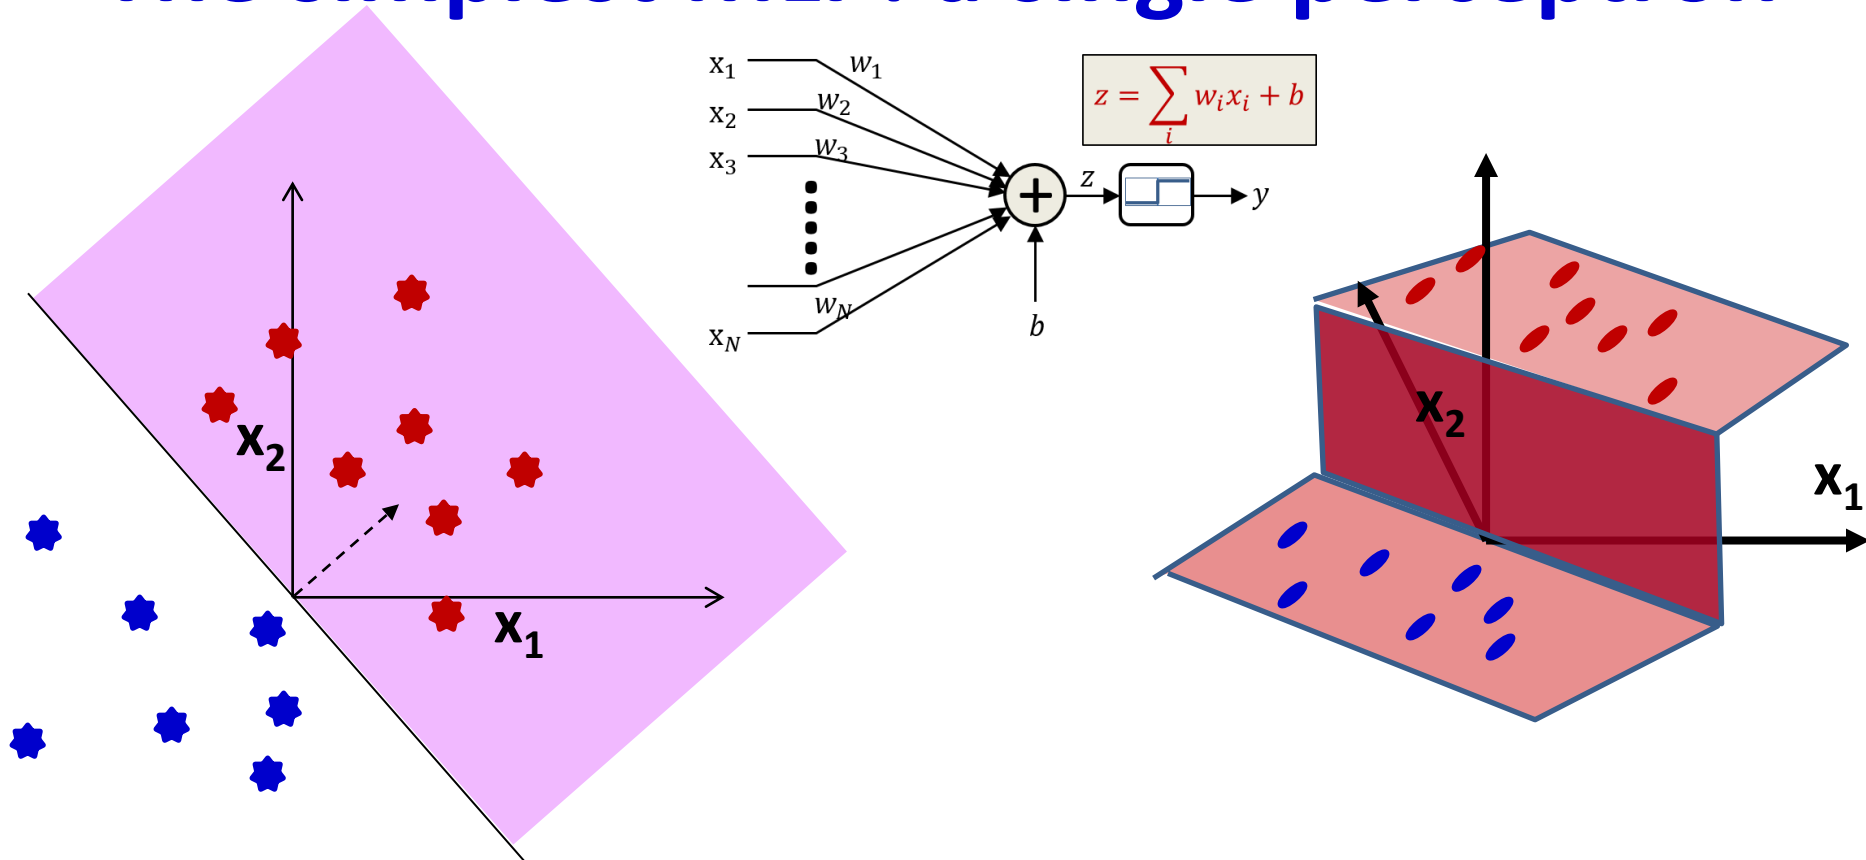
- The original MLP as proposed by Minsky: a network of threshold units
  - But how do you train it?
    - Given only “training” instances of input-output pairs

# The simplest MLP: a single perceptron



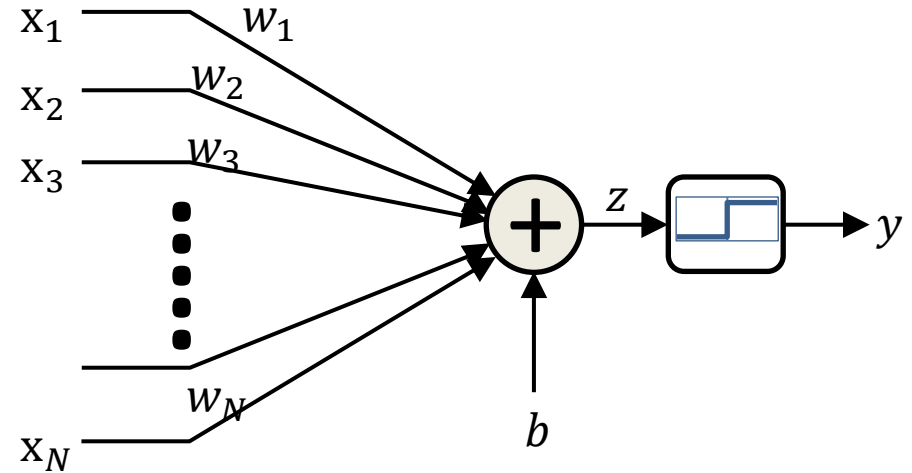
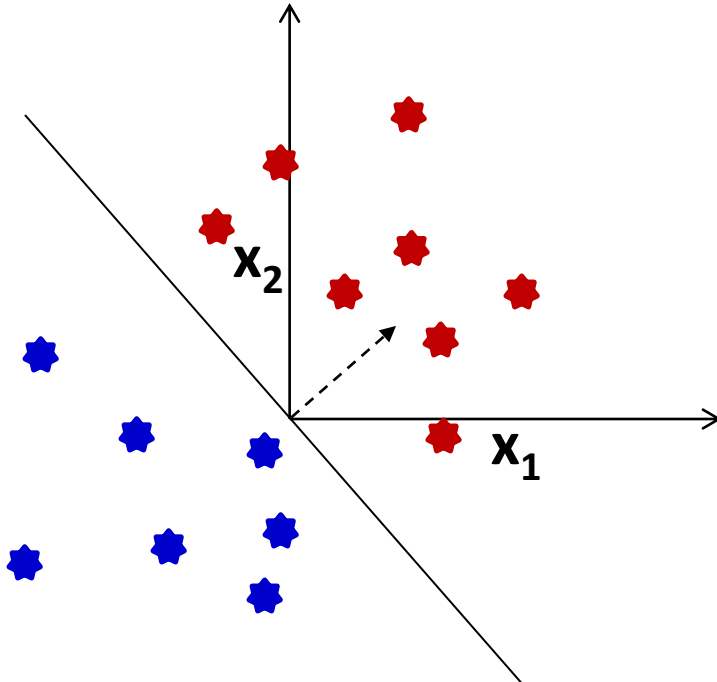
- Learn this function
  - A step function across a hyperplane

# The simplest MLP: a single perceptron



- Learn this function
  - A step function across a hyperplane
  - Given only samples from it

# Learning the perceptron



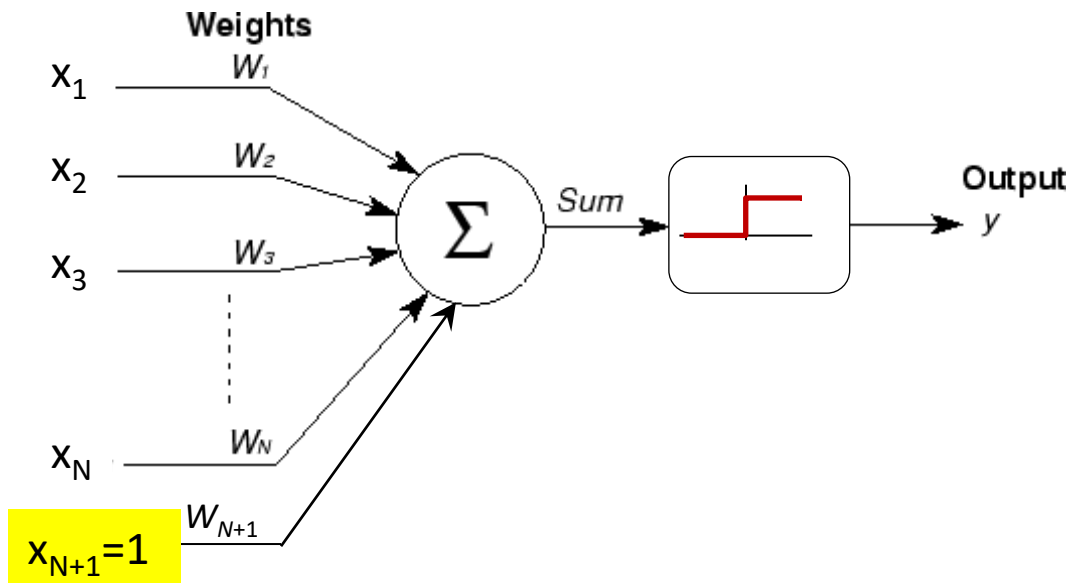
- Given a number of input output pairs, learn the weights and bias

$$y = \begin{cases} 1 & \text{if } \sum_{i=1}^N w_i X_i + b \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\text{Boundary: } \sum_{i=1}^N w_i X_i + b = 0$$

- Learn  $W = [w_1 \dots w_N]^T$  and  $b$ , given several  $(X, y)$  pairs

# Restating the perceptron



- Restating the perceptron equation by adding another dimension to  $X$

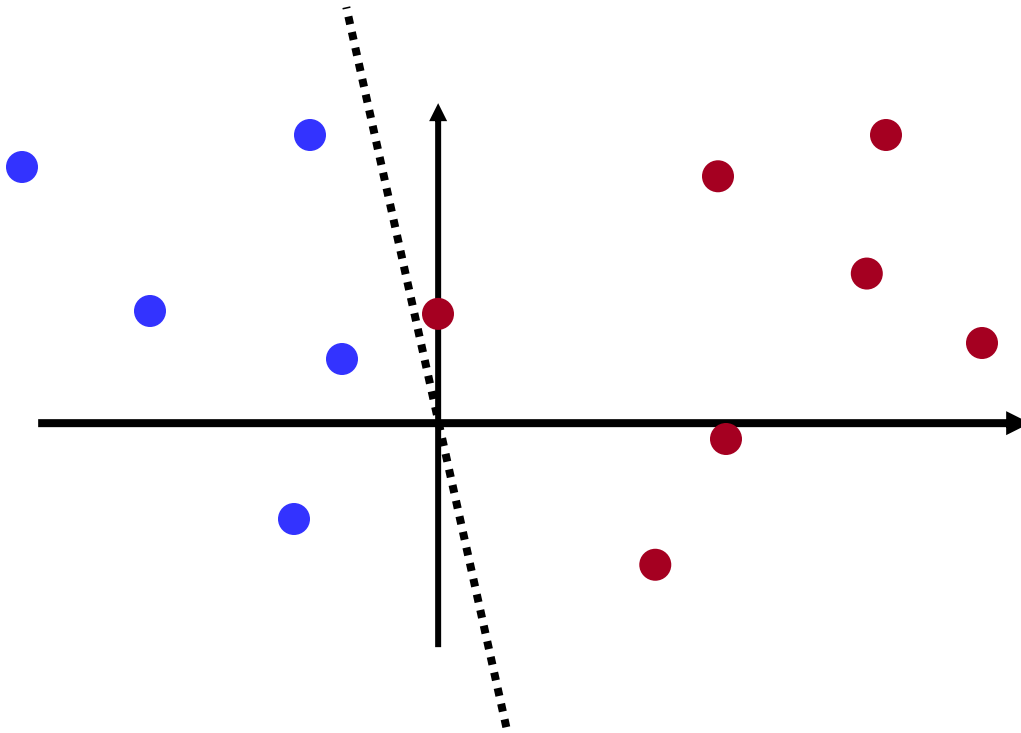
$$y = \begin{cases} 1 & \text{if } \sum_{i=1}^{N+1} w_i X_i \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

where  $X_{N+1} = 1$

- Note that the boundary  $\sum_{i=1}^{N+1} w_i X_i = 0$  is now a hyperplane through origin

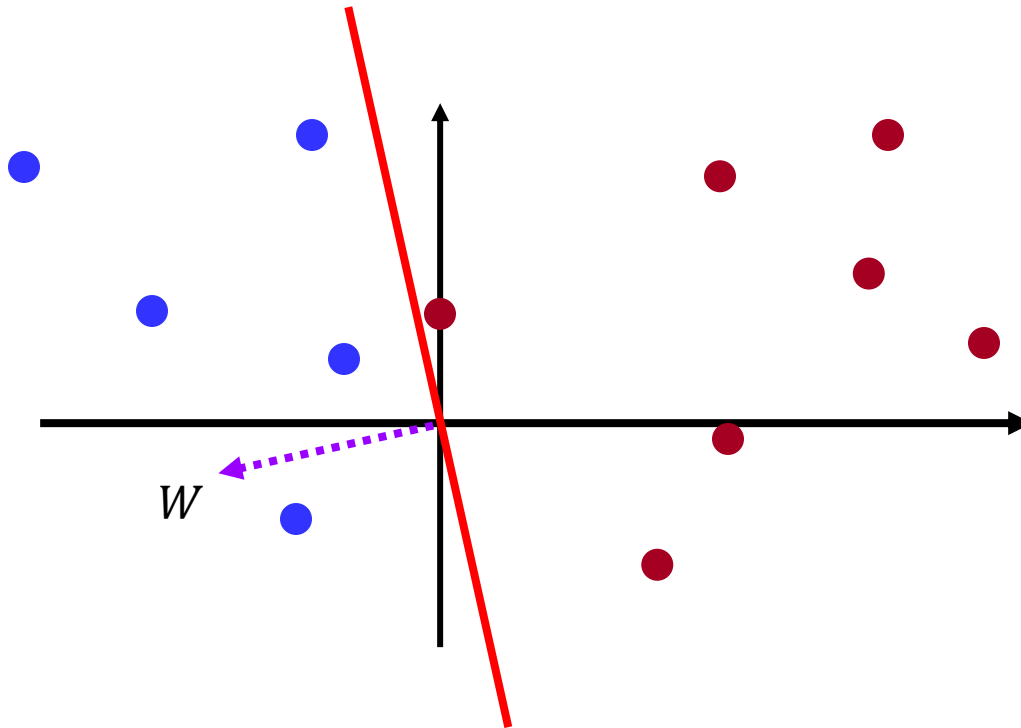


# The Perceptron Problem



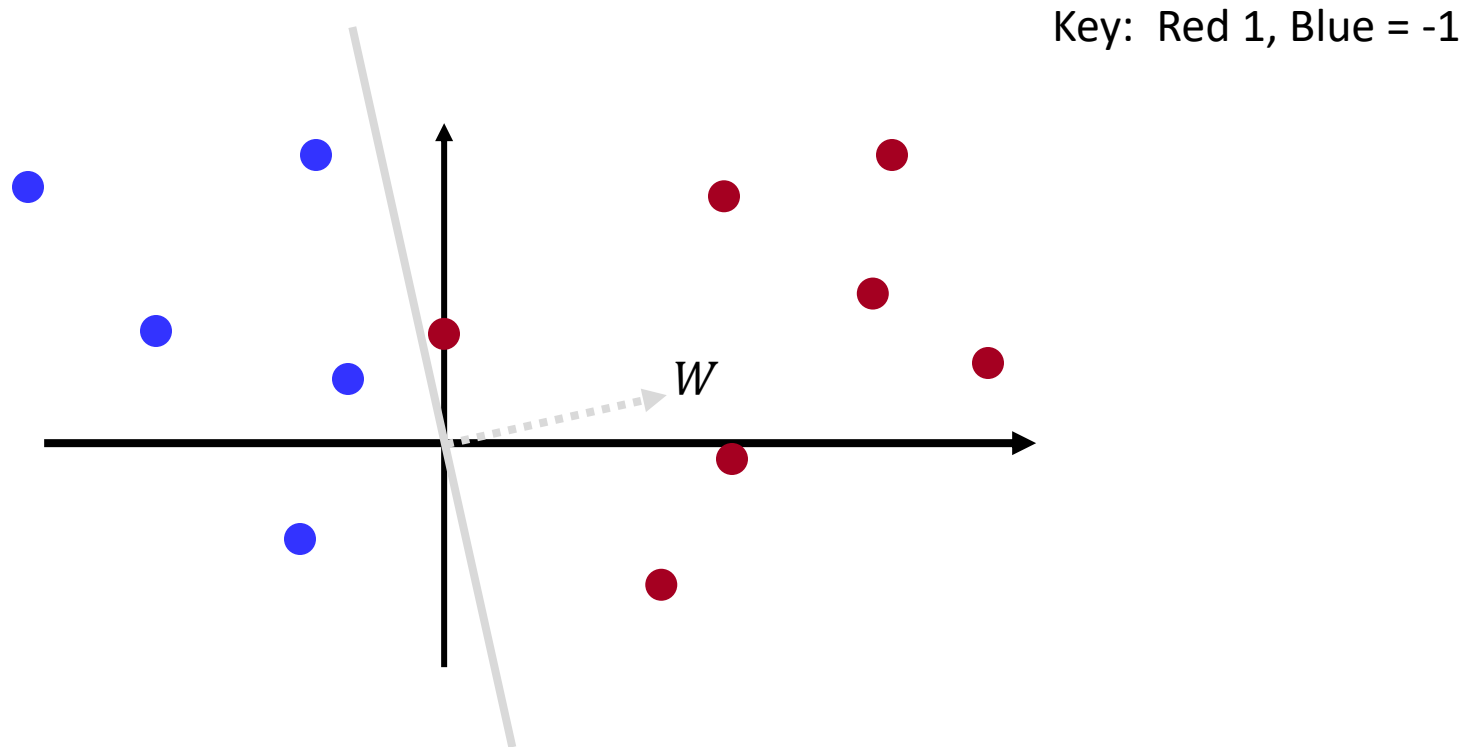
- Find the hyperplane  $\sum_{i=1}^{N+1} w_i X_i = 0$  that perfectly separates the two groups of points

# The Perceptron Problem



- Find the hyperplane  $\sum_{i=1}^{N+1} w_i X_i = 0$  that perfectly separates the two groups of points
  - Let vector  $W = [w_1, w_2, \dots, w_{N+1}]^T$  and vector  $X = [x_1, x_2, \dots, x_N, 1]^T$
  - $\sum_{i=1}^{N+1} w_i X_i = W^T X$  is an inner product
  - $W^T X = 0$  is the hyperplane comprising all  $X$ s orthogonal to vector  $W$ 
    - Learning the perceptron = finding the weight vector  $W$  for the separating hyperplane
    - $W$  points in the direction of the positive class

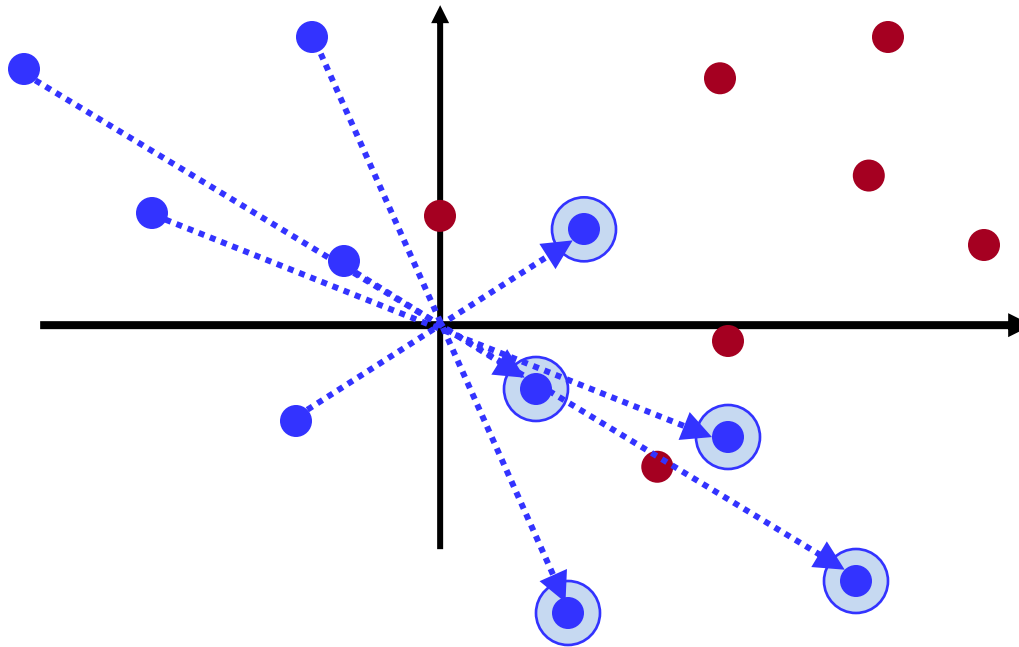
# The Perceptron Problem



- Learning the perceptron: Find the weights vector  $W$  such that the plane described by  $W^T X = 0$  perfectly separates the classes
  - $W^T X$  is positive for all red dots and negative for all blue ones

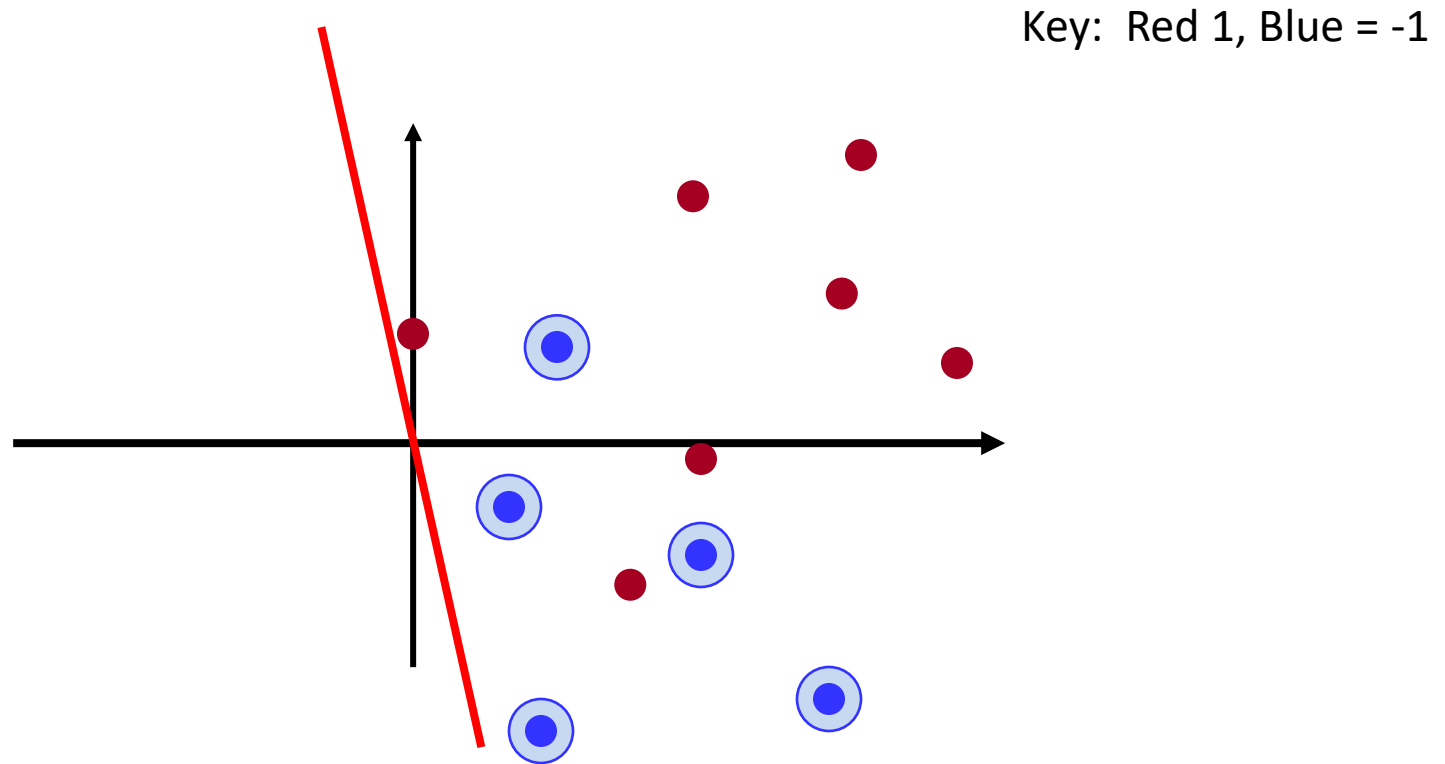
# A simple solution

Key: Red 1, Blue = -1



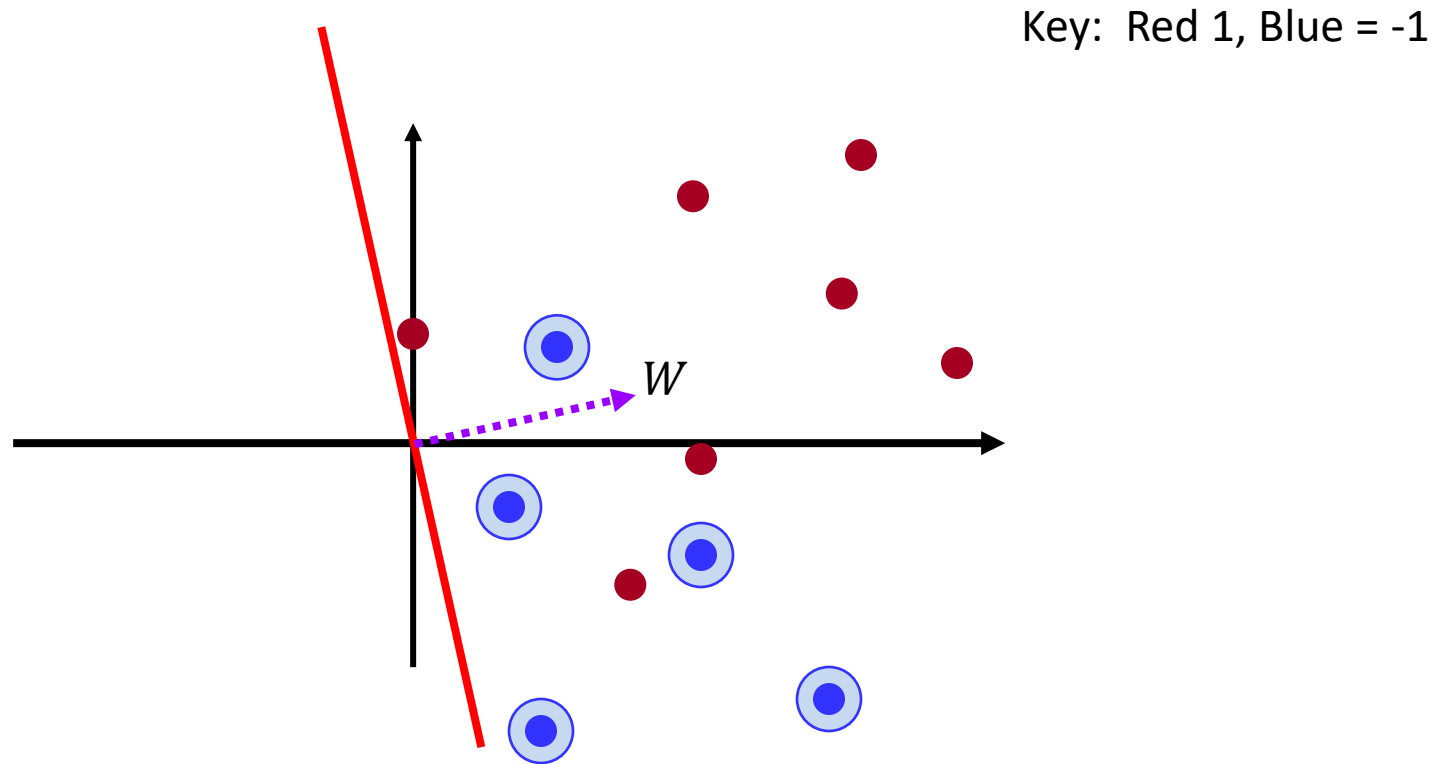
- *Reflect* all the negative instances across the origin
  - Negate every component of vector  $X$
- If we use class  $y \in \{+1, -1\}$  notation for the labels (instead of  $y \in \{0,1\}$ ), we can simply write the “reflected” values as  $X' = yX$ 
  - Will retain the features  $X$  for the positive class, but reflect/negate them for the negative class

# The Perceptron Solution



- Learning the perceptron: Find a plane such that all the modified ( $X'$ ) features lie on one side of the plane
  - Such a plane can always be found if the classes are linearly separable

# The Perceptron Solution: Linearly separable case



- When classes are linearly separable: a trivial solution

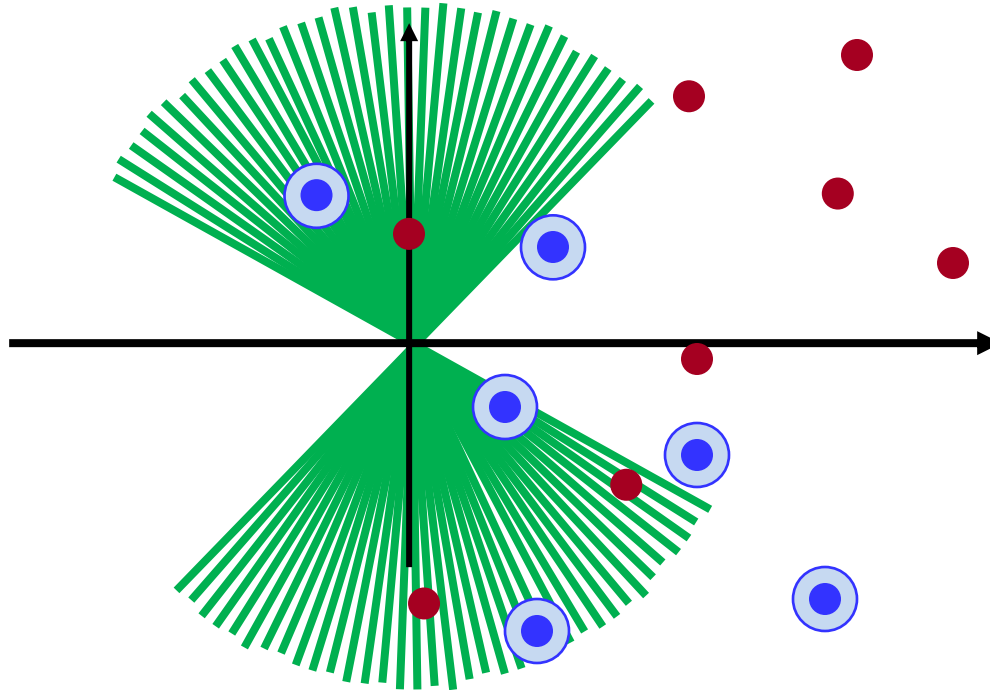
$$W = \frac{1}{N} \sum_i X'_i = \frac{1}{N} \sum_i y_i X_i$$

- Other solutions are also possible, e.g. max-margin solution

# The Perceptron Solution:

## when classes are not linearly separable

Key: Red 1, Blue = -1



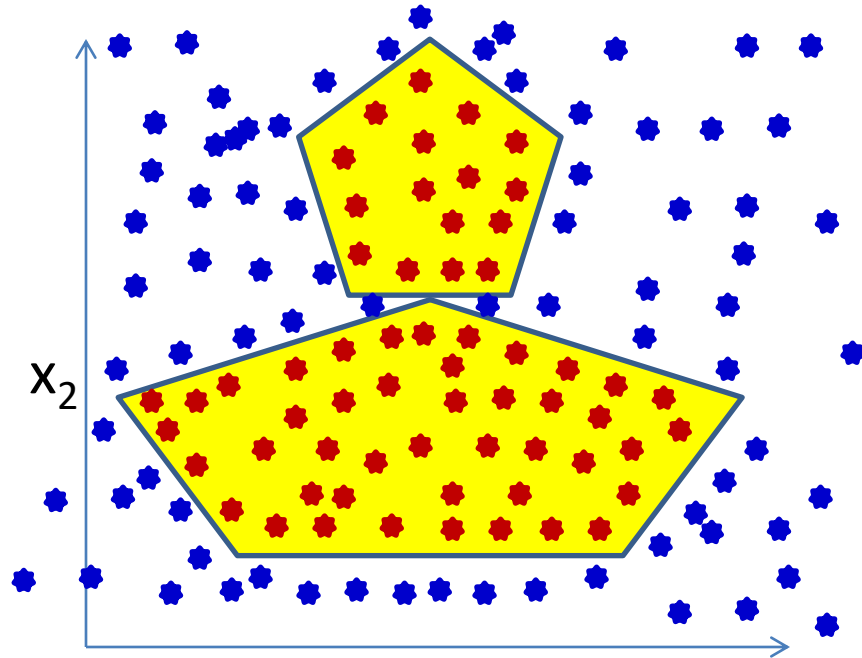
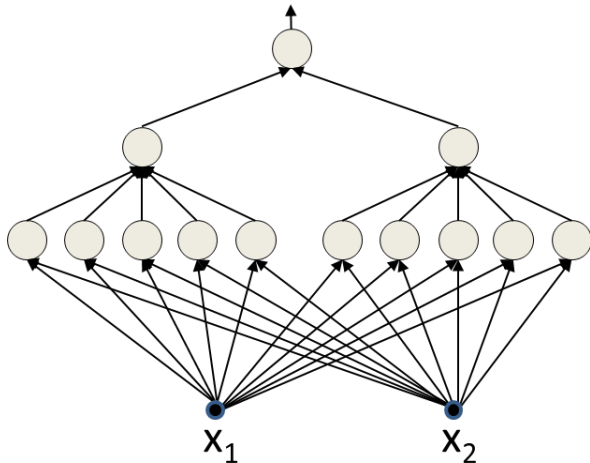
- When classes are not linearly separable, not possible to find a “support” plane
  - Some points will always lie on the other side
  - Model does not support perfect classification of this data

# The *online* perceptron solution

- The more popular solution, originally proposed by Rosenblatt is an *online* algorithm
  - The famous “perceptron” algorithm
- Initializes  $W$  and incrementally updates it each time we encounter an instance that is incorrectly classified
  - Guaranteed to find the correct solution for linearly separable data
  - On following slides, but will not cover in class

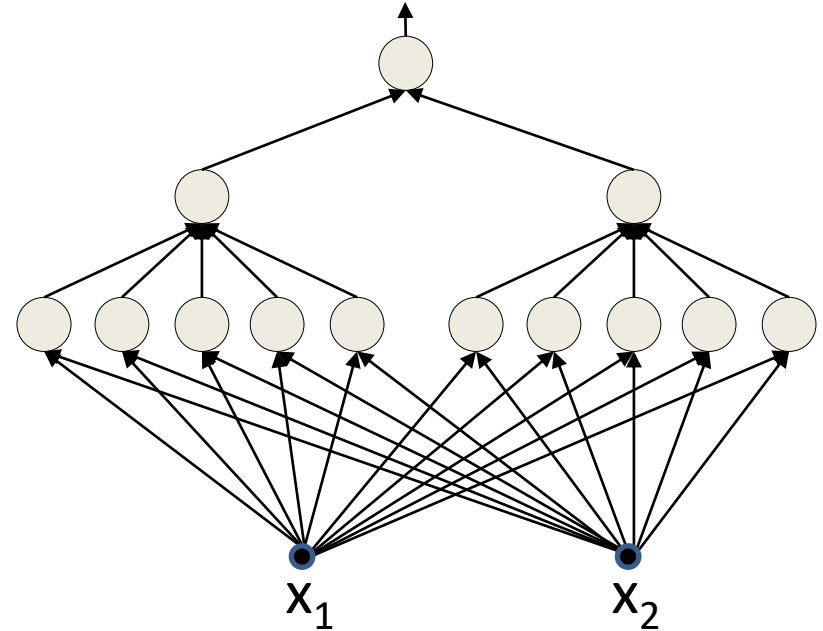
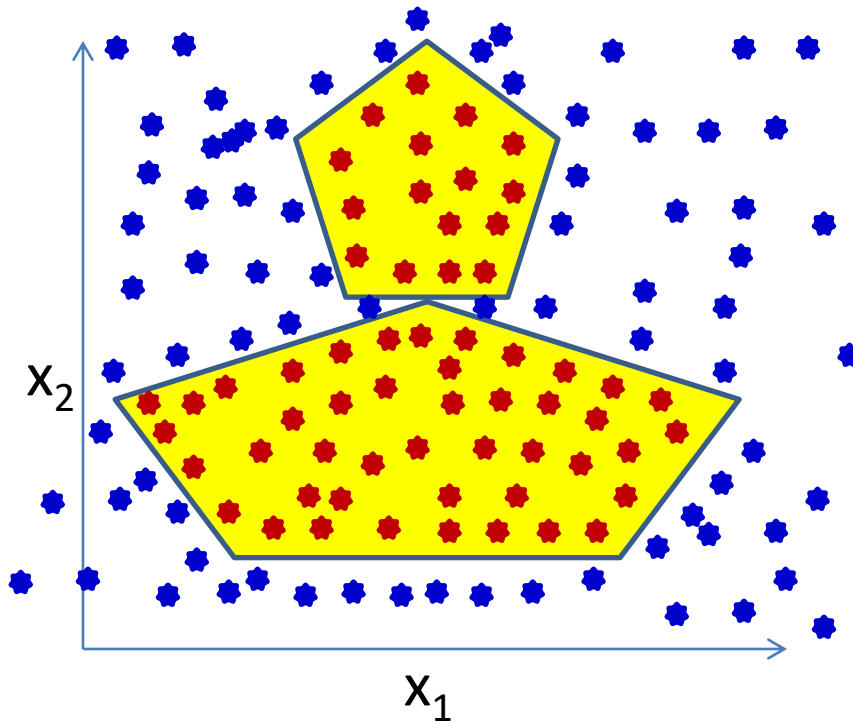


# History: A more complex problem



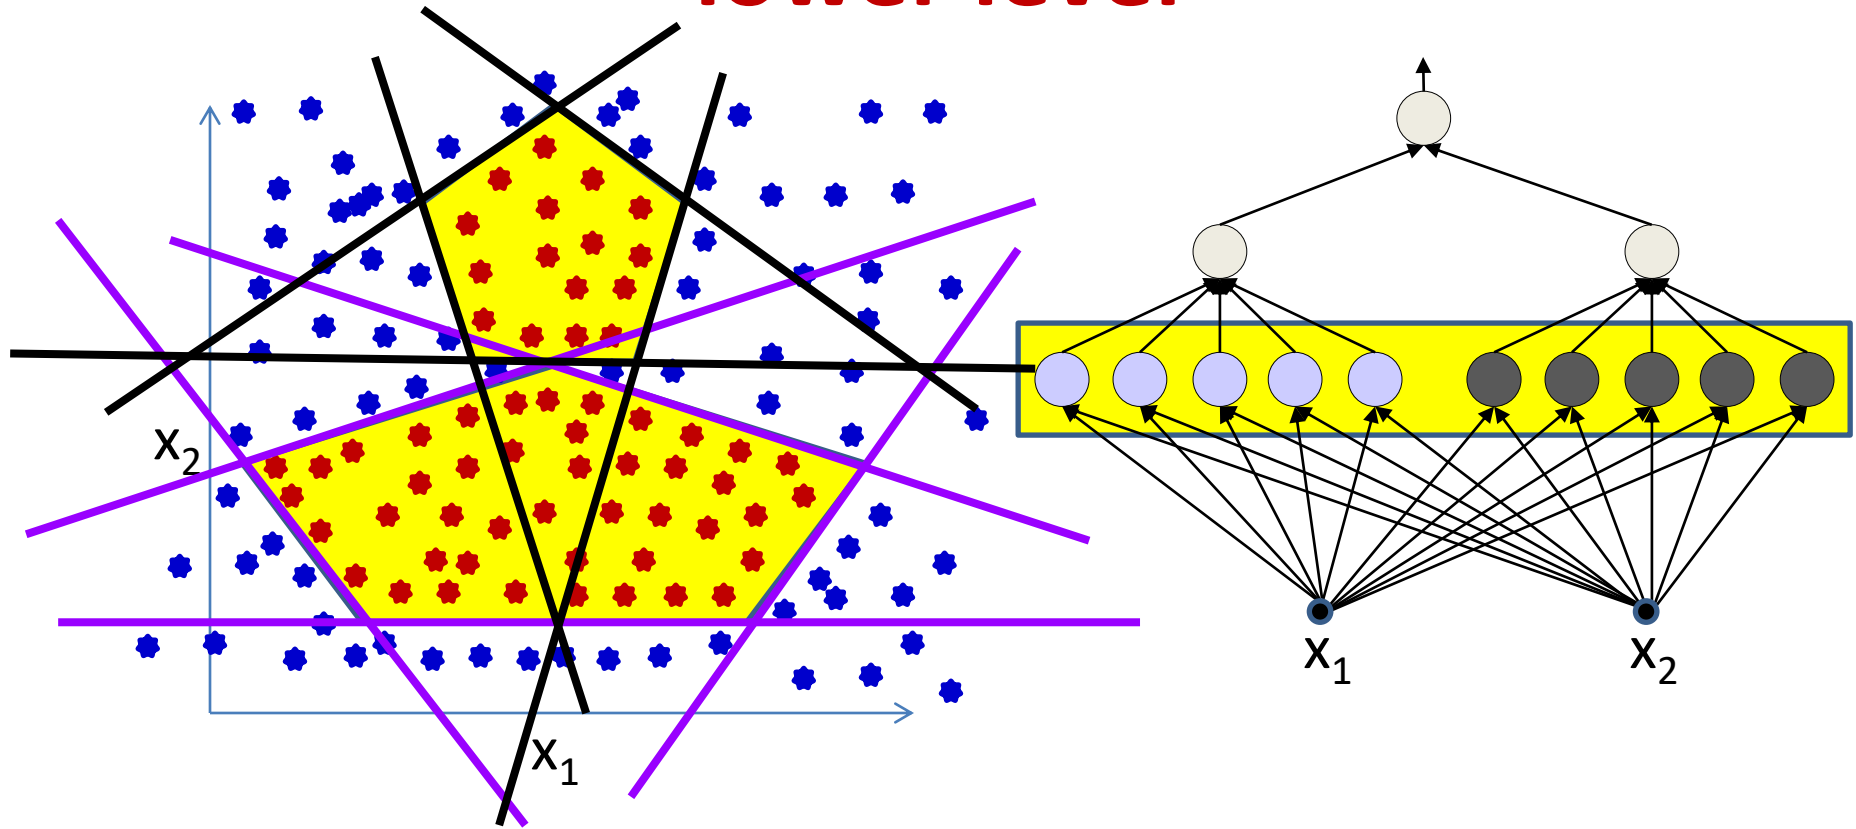
- Learn an *MLP* for this function
  - 1 in the yellow regions, 0 outside
- Using just the samples
- We know this can be perfectly represented using an MLP

# More complex decision boundaries



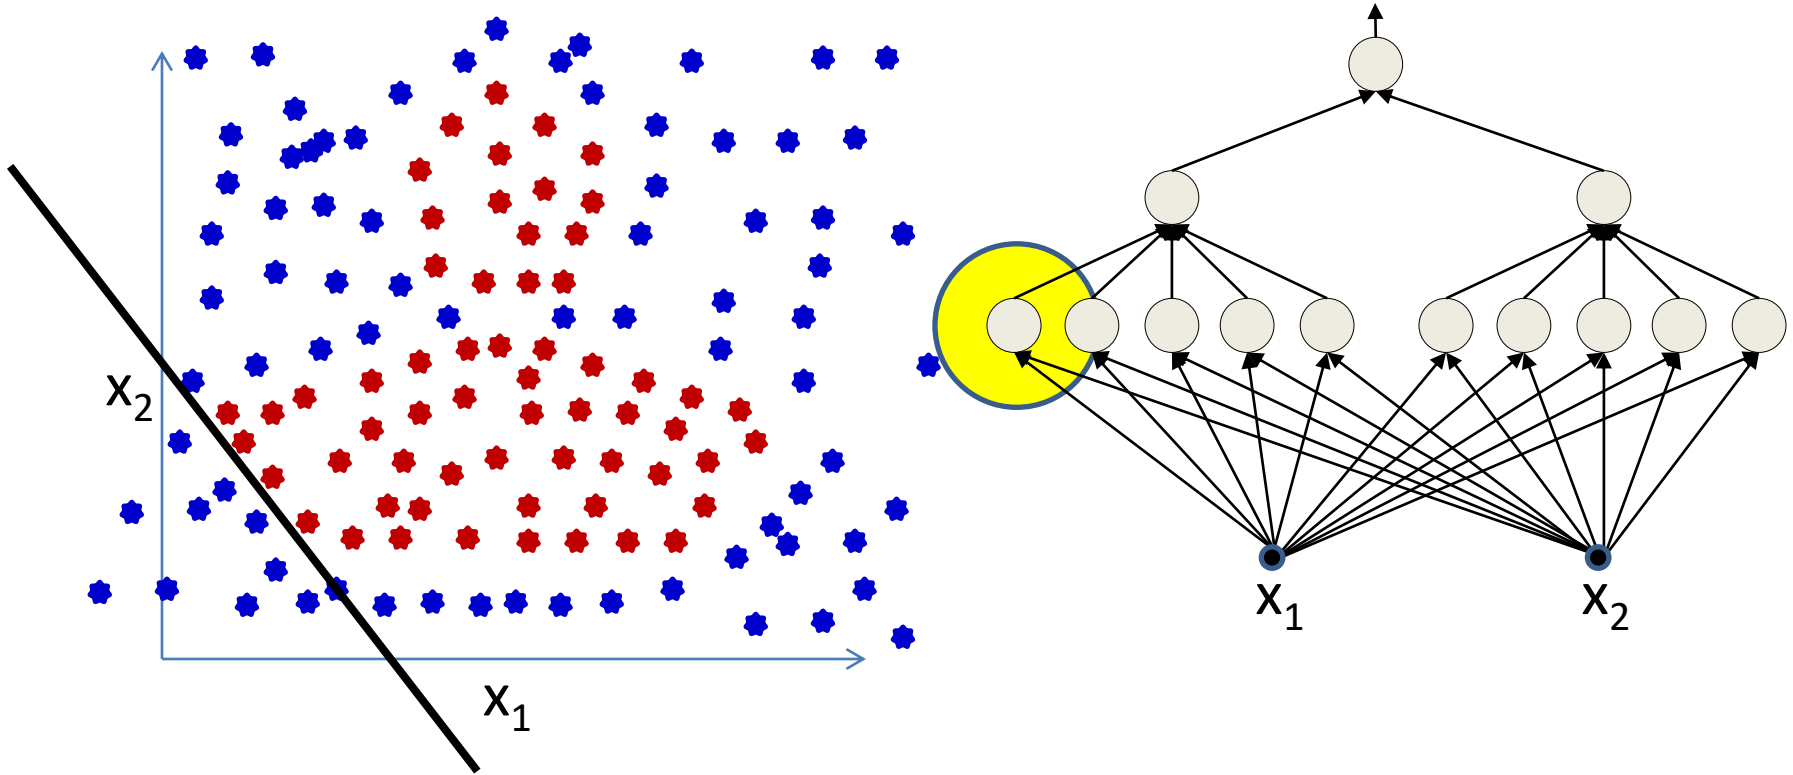
- Even using the perfect architecture...
- ... can we use perceptron learning rules to learn this classification function?

# The pattern to be learned at the lower level



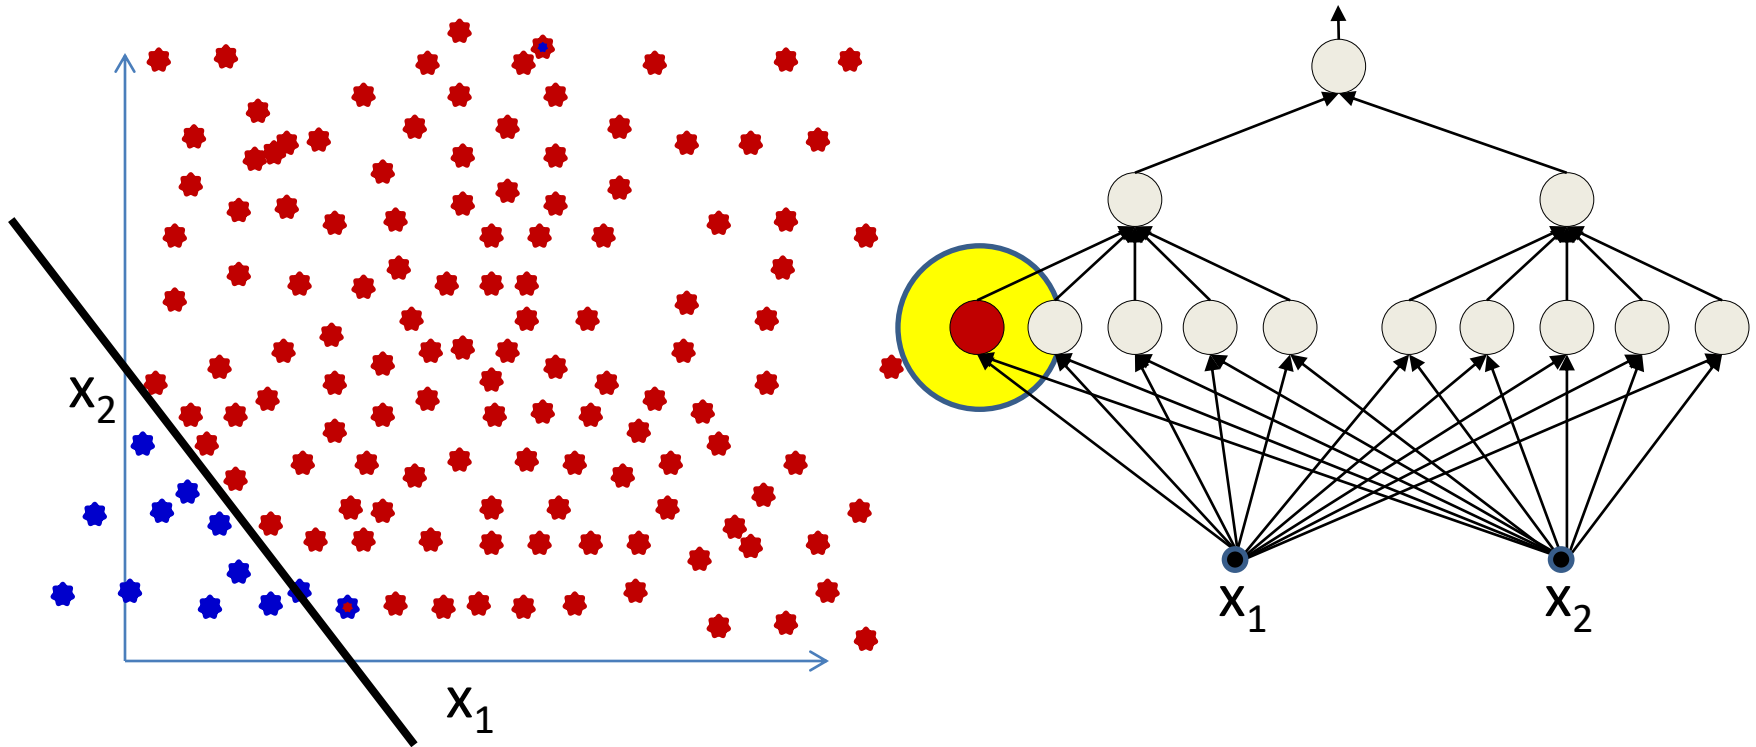
- The lower-level neurons are linear classifiers
  - They require linearly separated labels to be learned
  - The actually provided labels are not linearly separated
  - Challenge: *Must also learn the labels for the lowest units!* 59

# The pattern to be learned at the lower level



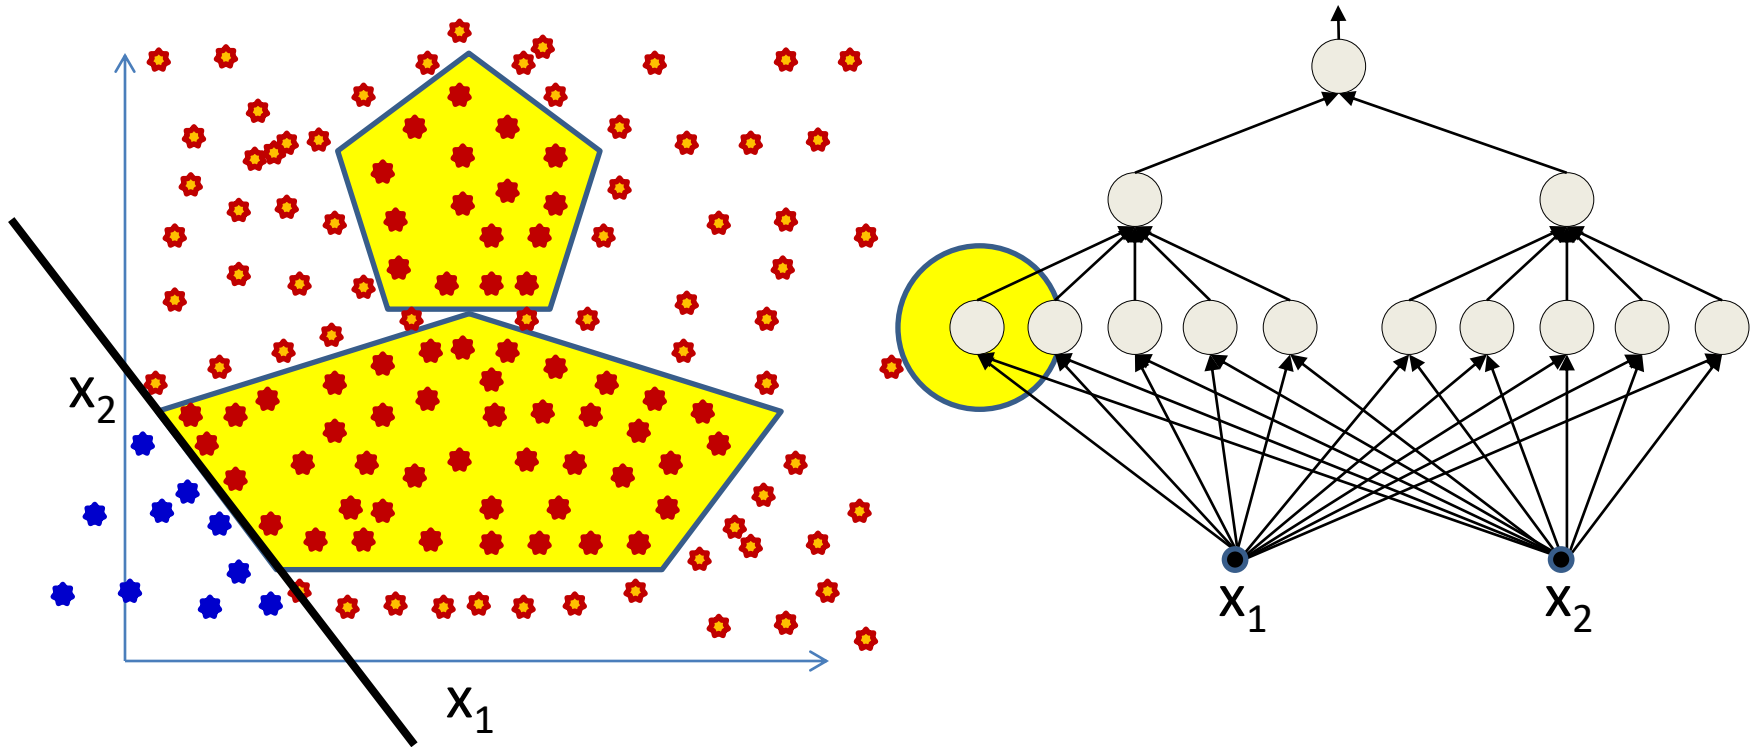
- Consider a single linear classifier that must be learned from the training data
  - Can it be learned from this data?

# The pattern to be learned at the lower level



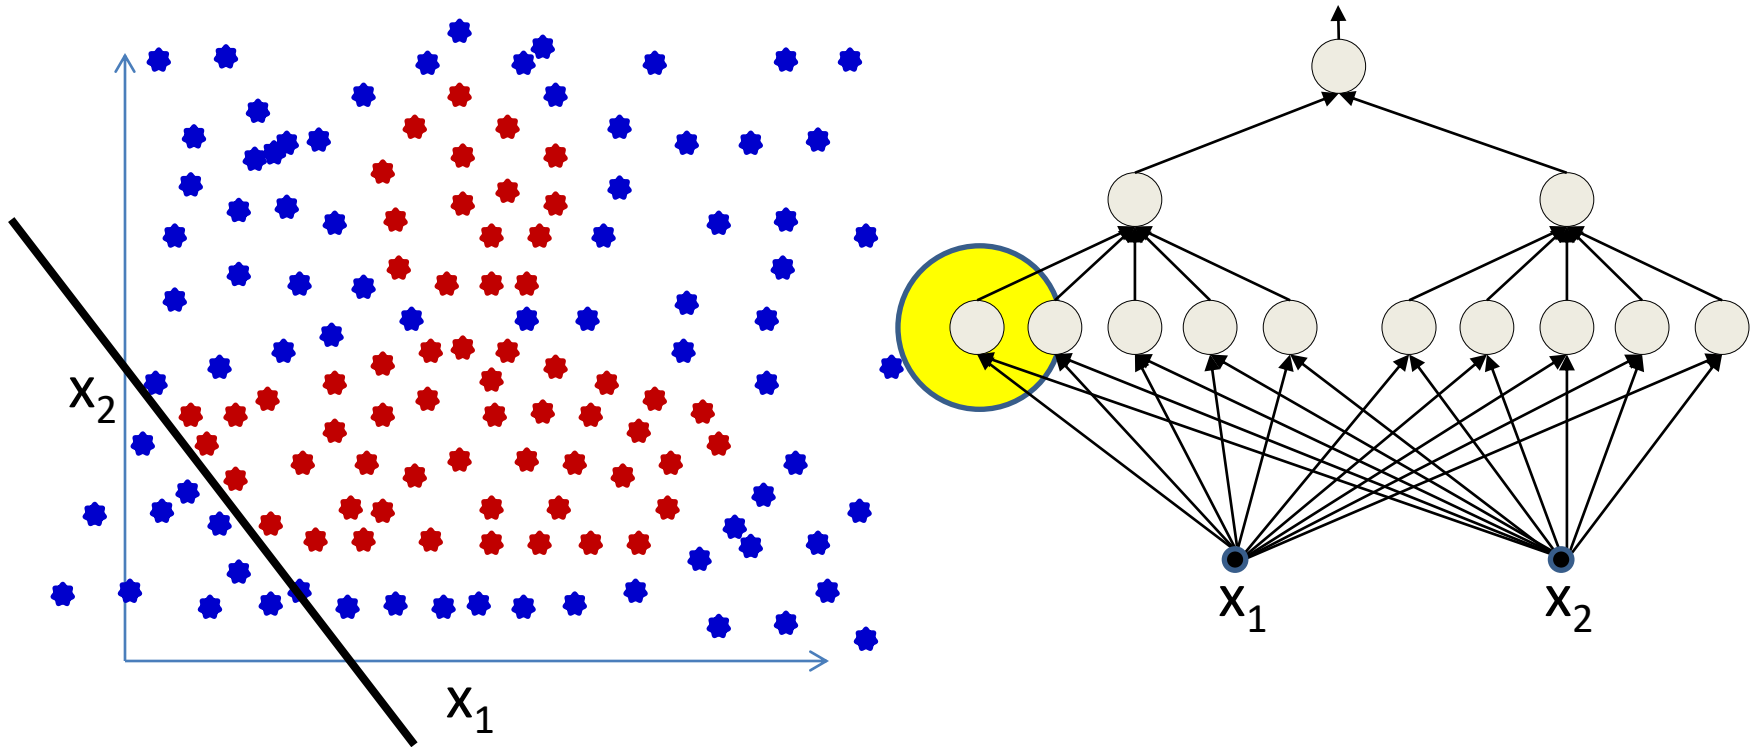
- Consider a single linear classifier that must be learned from the training data
  - Can it be learned from this data?
  - The individual classifier actually requires the kind of labelling shown here
    - Which is *not* given!!

# The pattern to be learned at the lower level



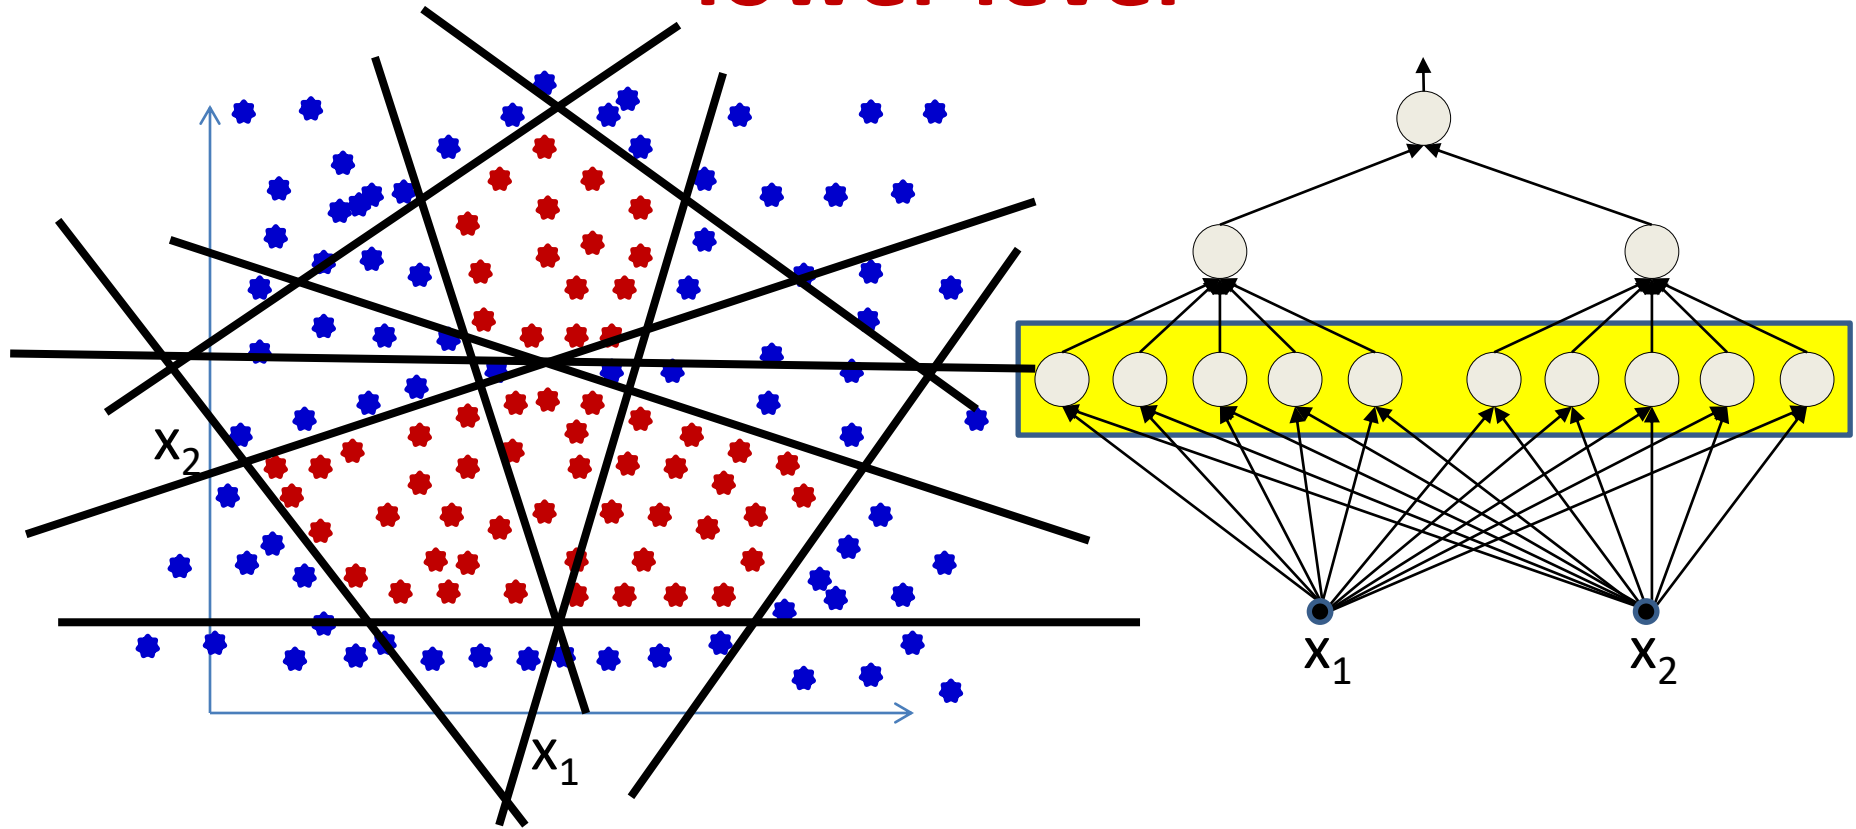
- The lower-level neurons are linear classifiers
  - They require linearly separated labels to be learned
  - The actually provided labels are not linearly separated
  - *Challenge: Must also learn the labels for the lowest units!*

# The pattern to be learned at the lower level



- For a single line:
  - Try out *every possible way of relabeling the blue dots such that we can learn a line that keeps all the red dots on one side!*

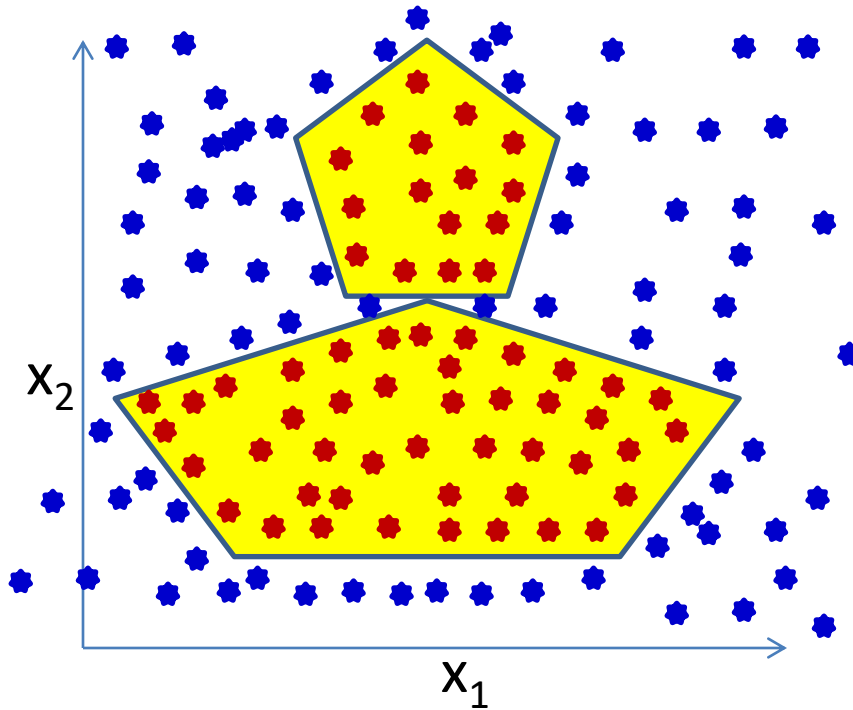
# The pattern to be learned at the lower level



- This must be done for *each* of the lines (perceptrons)
- Such that, when all of them are combined by the higher-level perceptrons we get the desired pattern
  - Basically an exponential search over inputs



Individual neurons represent one of the lines that compose the figure (linear classifiers)



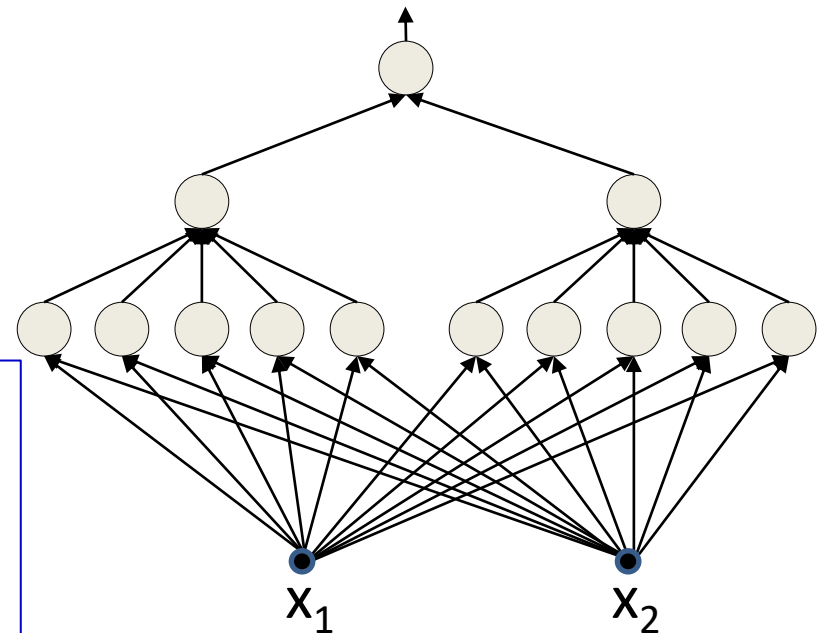
Must know the output of every neuron for *every* training instance, in order to learn this neuron

The outputs should be such that the neuron individually has a linearly separable task

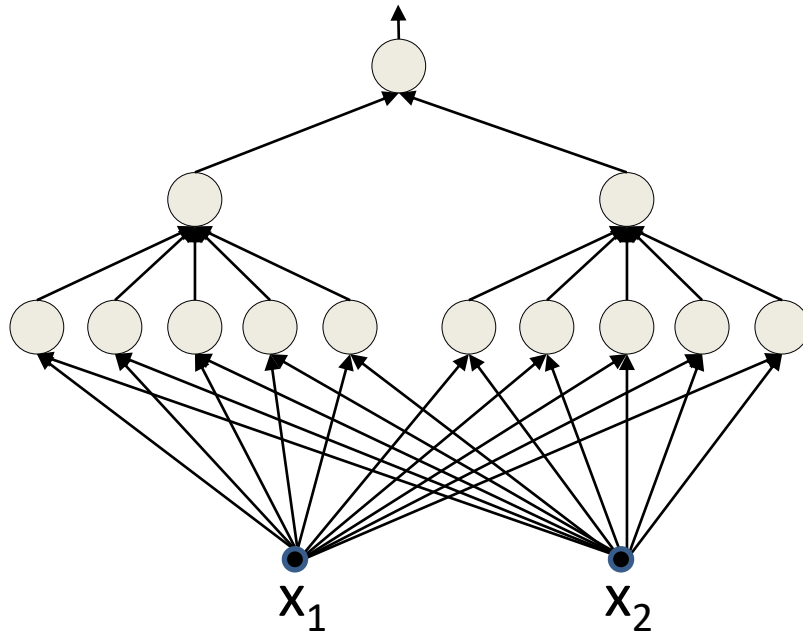
The linear separators must combine to form the desired boundary

This must be done for *every* neuron

Getting any of them wrong will result in incorrect output!



# Learning a *multilayer* perceptron



Training data only specifies  
input and output of network

Intermediate outputs (outputs  
of individual neurons) are not specified

- Training this network using the perceptron rule is a combinatorial optimization problem
- We don't know the outputs of the individual intermediate neurons in the network for any training input
- **Must also determine the correct output for *each* neuron for *every* training instance**
- **At least exponential (in inputs) time complexity!!!!!!**

# Greedy algorithms: Adaline and Madaline

- Perceptron learning rules cannot directly be used to learn an MLP
  - Exponential complexity of assigning intermediate labels
    - Even worse when classes are not actually separable
- Can we use a *greedy* algorithm instead?
  - Adaline / Madaline
  - On slides, will skip in class (check the quiz)

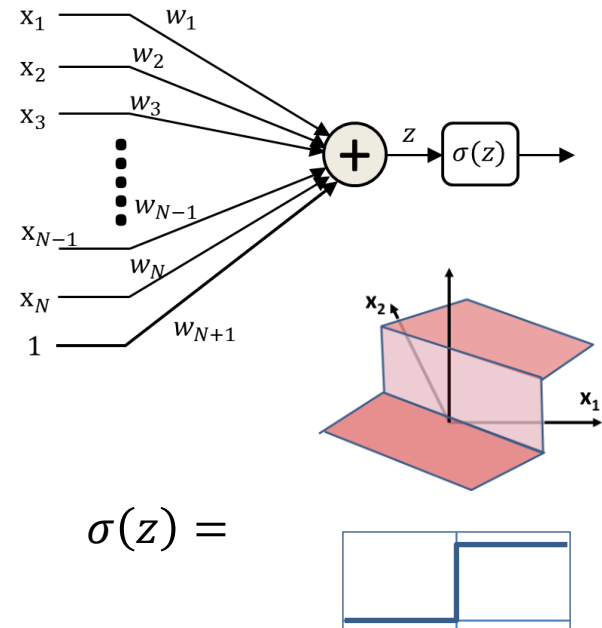
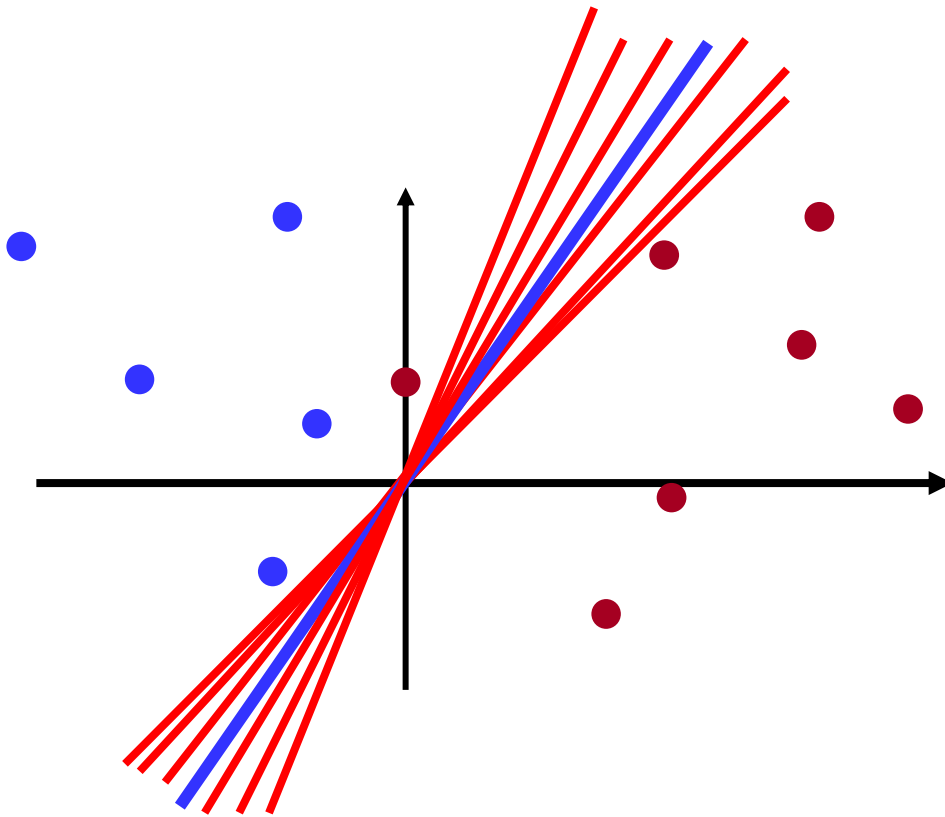
# Story so far

- “Learning” a network = learning the weights and biases to compute a target function
  - Will require a network with sufficient “capacity”
- In practice, we learn networks by “fitting” them to match the input-output relation of “training” instances drawn from the target function
- A linear decision boundary can be learned by a single perceptron (with a threshold-function activation) in linear time if classes are linearly separable
- Non-linear decision boundaries require networks of perceptrons
- Training an MLP with threshold-function activation perceptrons will require knowledge of the input-output relation for every training instance, for *every* perceptron in the network
  - These must be determined as part of training
  - For threshold activations, this is an NP-complete combinatorial optimization problem

# History..

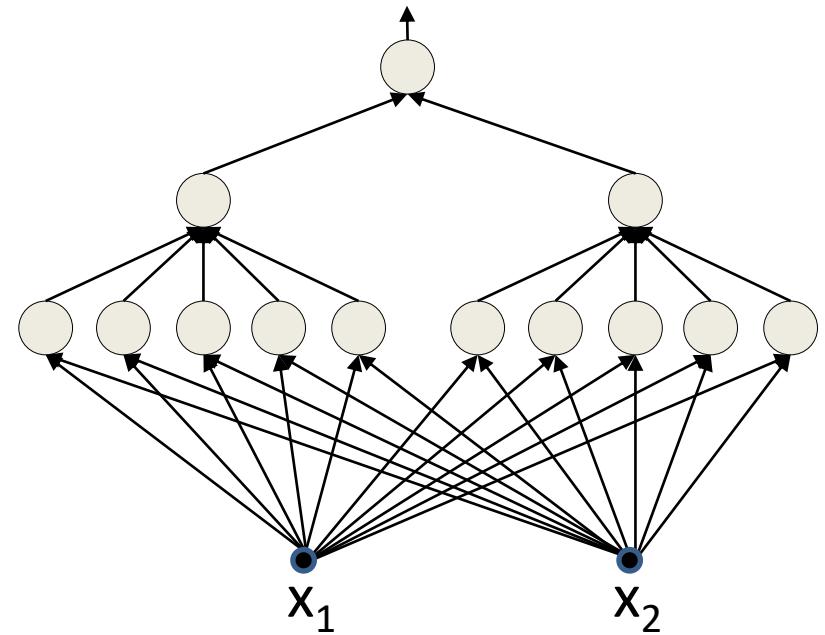
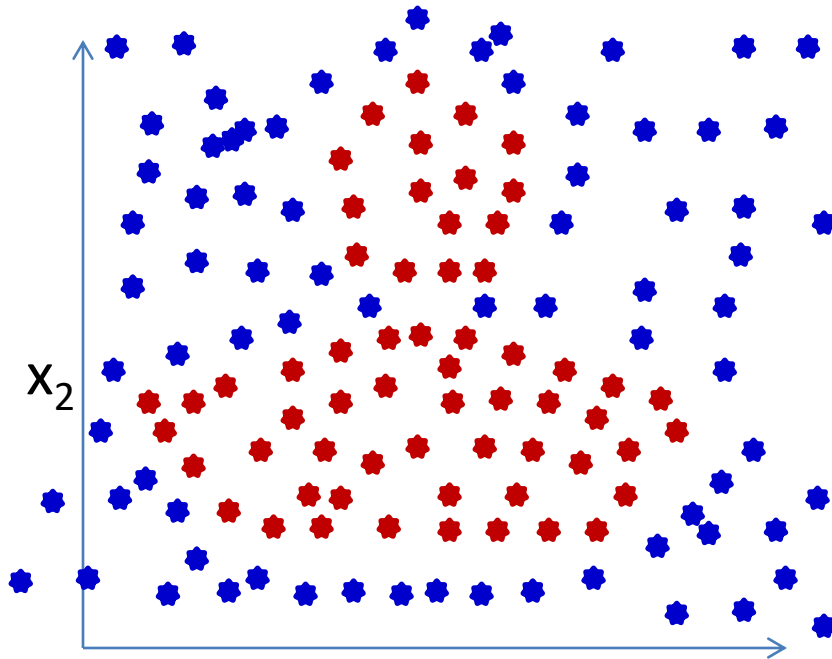
- The realization that training an entire MLP was a combinatorial optimization problem stalled development of neural networks for well over a decade!

# Why this problem?



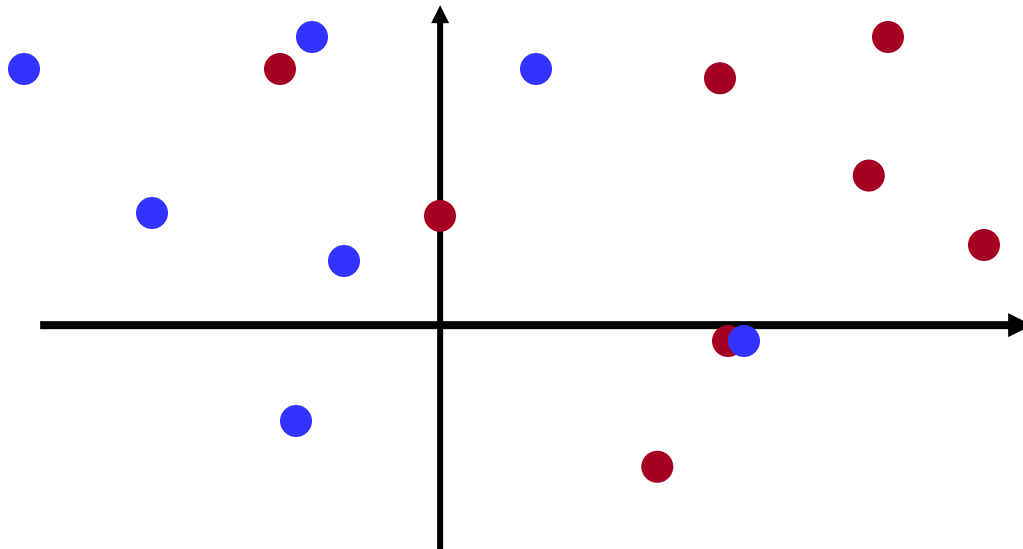
- The perceptron is a flat function with zero derivative everywhere, except at 0 where it is non-differentiable
  - You can vary the weights a *lot* without changing the error
  - There is no indication of which direction to change the weights to reduce error

# This only compounds on larger problems



- Individual neurons' weights can change significantly without changing overall error
- The simple MLP is a flat, non-differentiable function
  - Actually a function with 0 derivative nearly everywhere, and no derivatives at the boundaries

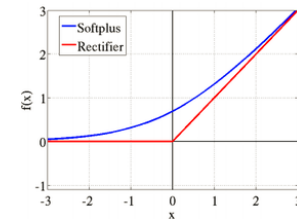
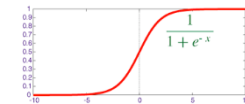
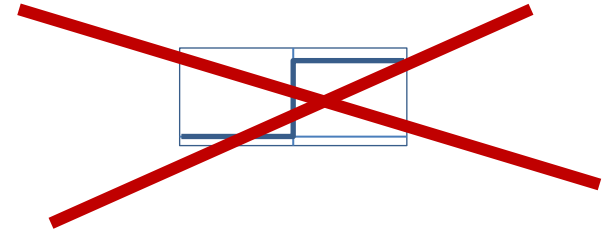
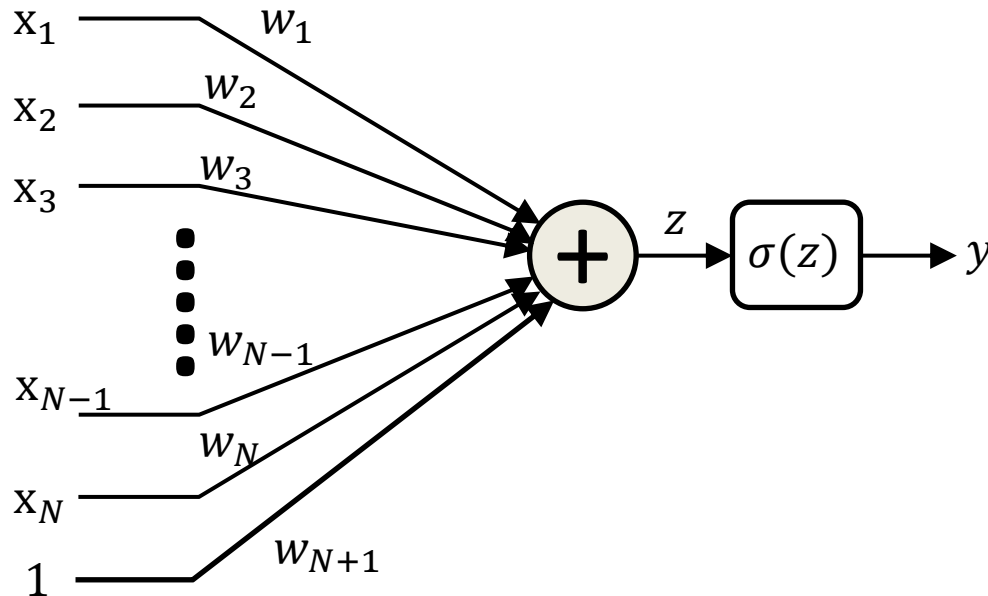
# A second problem: What we *actually* model



- Real-life data are rarely clean
  - Not linearly separable
  - Rosenblatt's perceptron wouldn't work in the first place



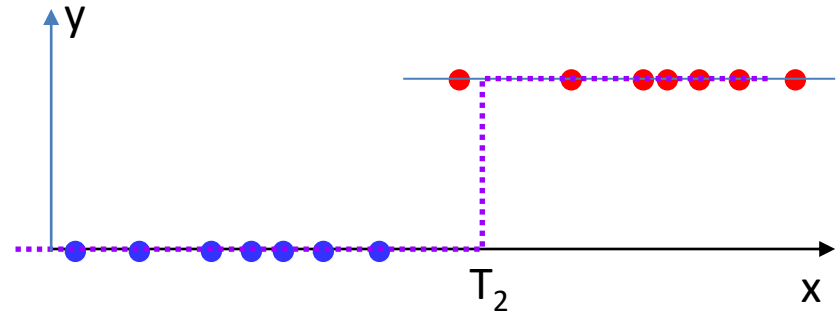
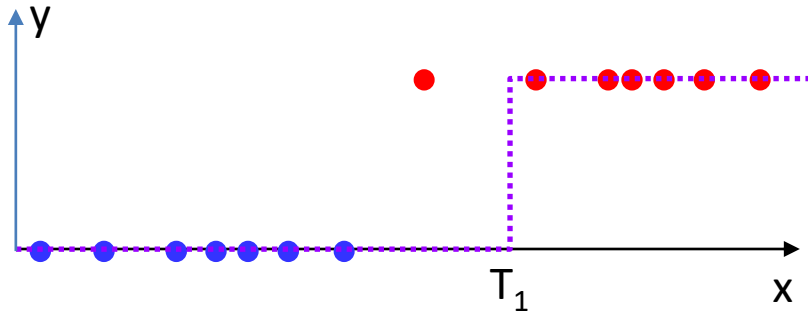
# Solution



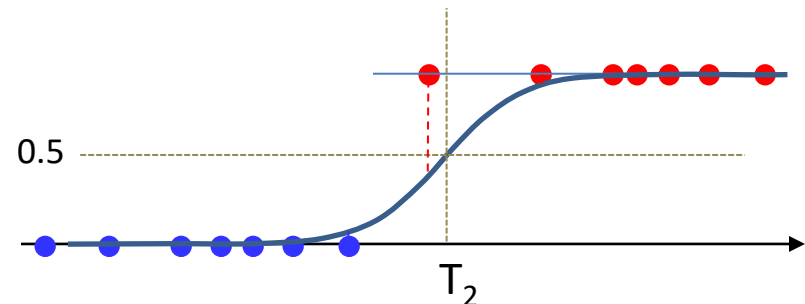
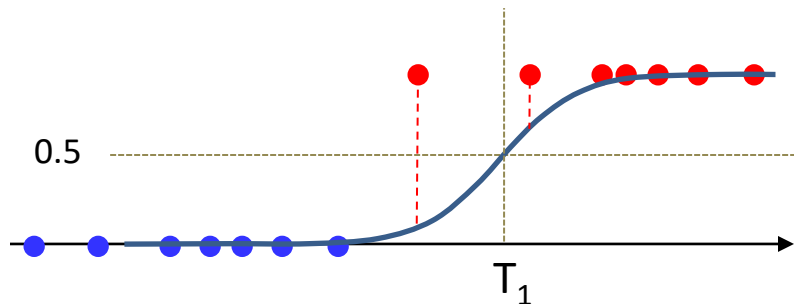
Activation functions  $\sigma(z)$

- Lets make the neuron differentiable, *with non-zero derivatives over much of the input space*
  - Small changes in weight can result in non-negligible changes in output
  - This enables us to estimate the parameters using gradient descent techniques..

# Differentiable activation function

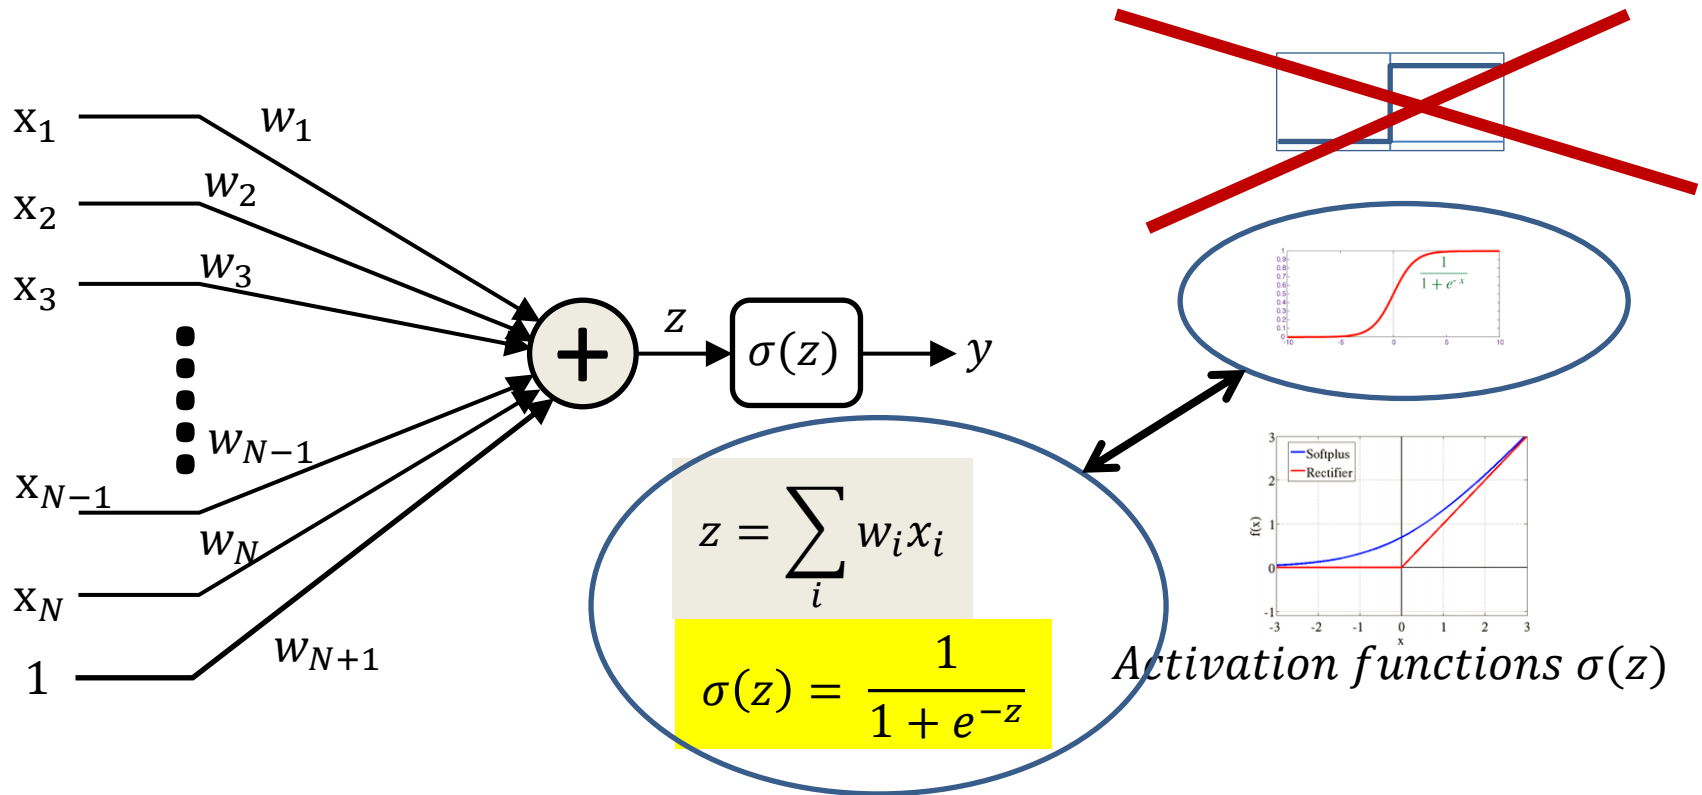


- Threshold activation: shifting the threshold from  $T_1$  to  $T_2$  does not change classification error
  - Does not indicate if moving the threshold left was good or not



- Smooth, continuously varying activation: Classification based on whether the output is greater than 0.5 or less
  - Can now quantify *how much* the output differs from the desired target value (0 or 1)
  - Moving the function left or right changes this quantity, even if the classification error itself doesn't change

# The sigmoid activation is special

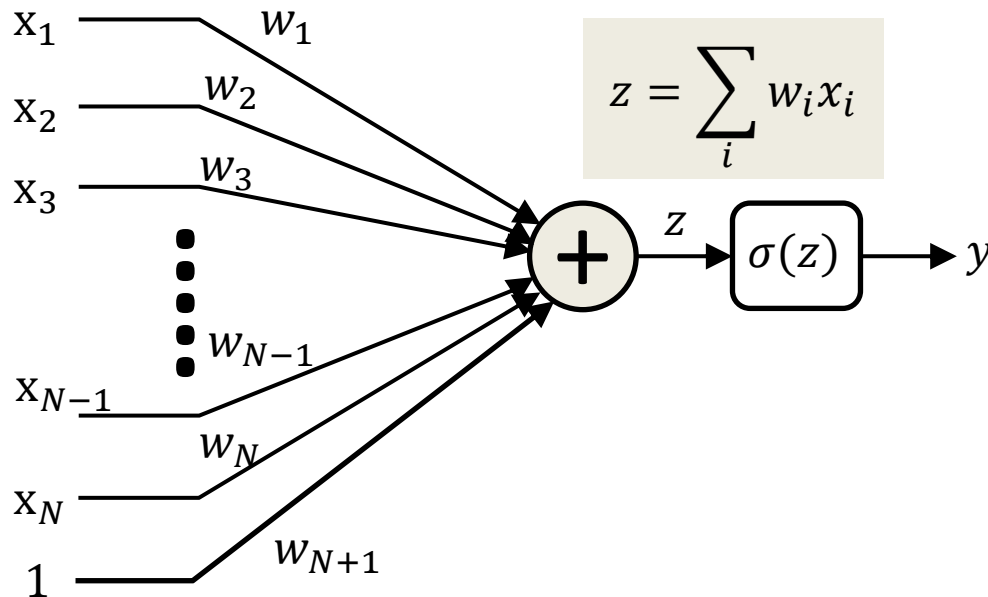


- This particular one has a nice interpretation
- It can be interpreted as  $P(y = 1|x)$

# Perceptrons and probabilities

- We will return to the fact that perceptrons with sigmoidal activations actually model class probabilities in a later lecture
- But for now moving on..

# Perceptrons with differentiable activation functions



$$\frac{dy}{dz} = \sigma'(z)$$

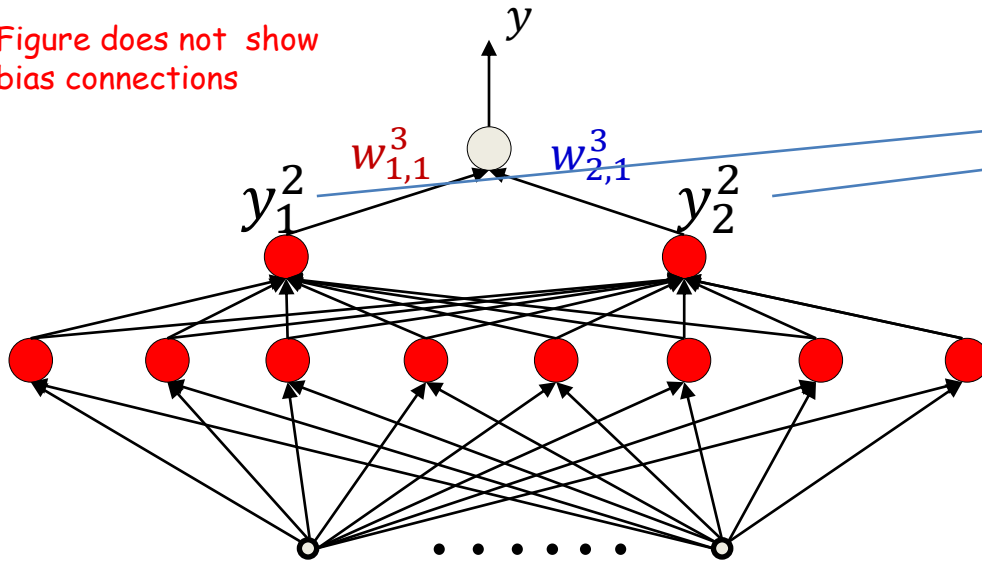
$$\frac{dy}{dw_i} = \frac{dy}{dz} \frac{dz}{dw_i} = \sigma'(z) x_i$$

$$\frac{dy}{dx_i} = \frac{dy}{dz} \frac{dz}{dx_i} = \sigma'(z) w_i$$

- $\sigma(z)$  is a differentiable function of  $z$ 
  - $\frac{d\sigma(z)}{dz}$  is well-defined and finite for all  $z$
- Using the chain rule,  $y$  is a differentiable function of both inputs  $x_i$  and weights  $w_i$
- This means that we can compute the change in the output for *small* changes in either the input or the weights

# Overall network is differentiable

Figure does not show bias connections



$$y = \sigma(w_{1,1}^3 y_1^2 + w_{2,1}^3 y_2^2 + w_{3,1}^3)$$

$y$  = output of overall network

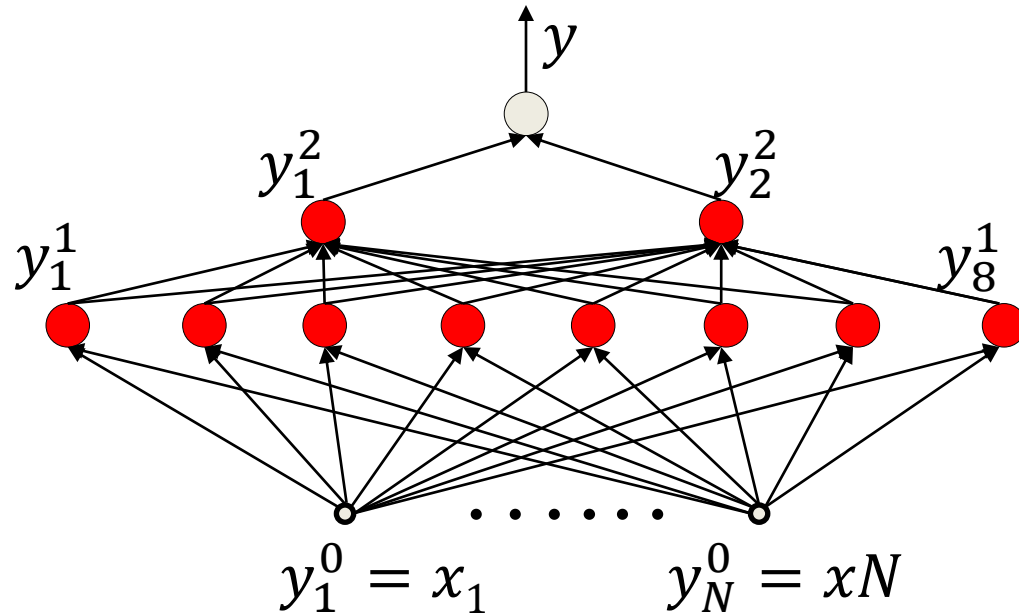
$w_{i,j}^k$  = weight connecting the  $i$ th unit of the  $(k-1)$ th layer to the  $j$ th unit of the  $k$ -th layer

$y_i^k$  = output of the  $i$ th unit of the  $k$ th layer

$\sigma()$  is differentiable w.r.t both  $w$  and  $y_i^k$

- Every individual perceptron is differentiable w.r.t its inputs and its weights (including “bias” weight)
- By the chain rule, the overall function is differentiable w.r.t every parameter (weight or bias)
  - Small changes in the parameters result in measurable changes in output

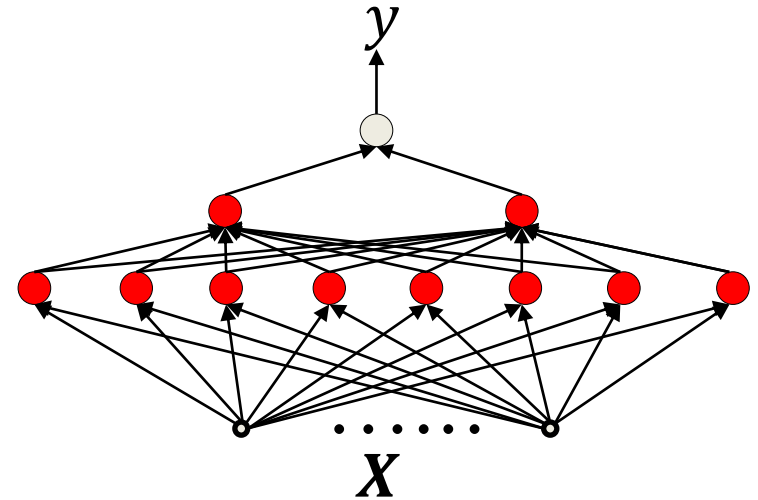
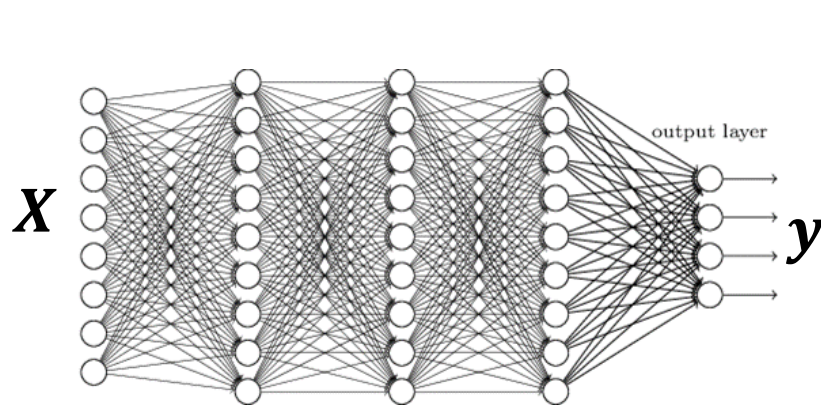
# Overall function is differentiable



$$y_j^k = \sigma \left( \sum_i w_{i,j}^{k-1} y_i^{k-1} \right)$$

- The overall function is differentiable w.r.t every parameter
  - We can compute how small changes in the parameters change the output
    - For non-threshold activations the derivative are finite and generally non-zero
  - We will derive the actual derivatives using the chain rule later

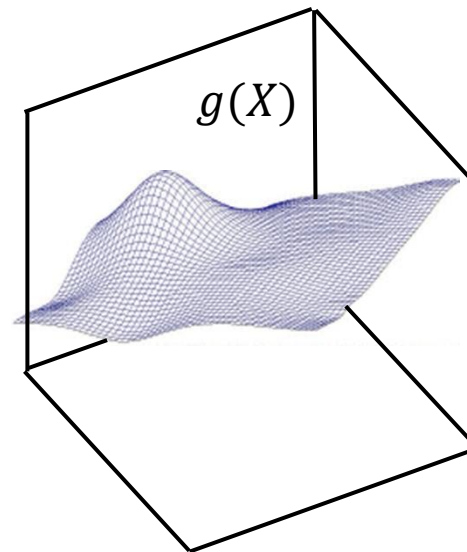
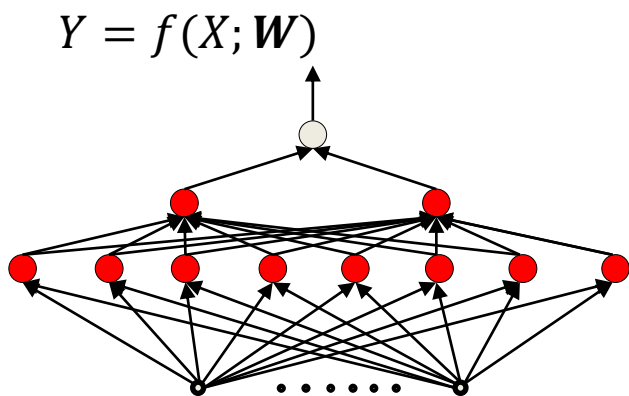
# Overall setting for “Learning” the MLP



- Given a training set of input-output pairs  $(X_1, d_1), (X_2, d_2), \dots, (X_N, d_N) \dots$ 
  - $d$  is the *desired output* of the network in response to  $X$
  - $X$  and  $d$  may both be vectors
- ...we must find the network parameters such that the network produces the desired output for each training input
  - Or a close approximation of it
  - **The *architecture* of the network must be specified by us**



# Recap: Learning the function

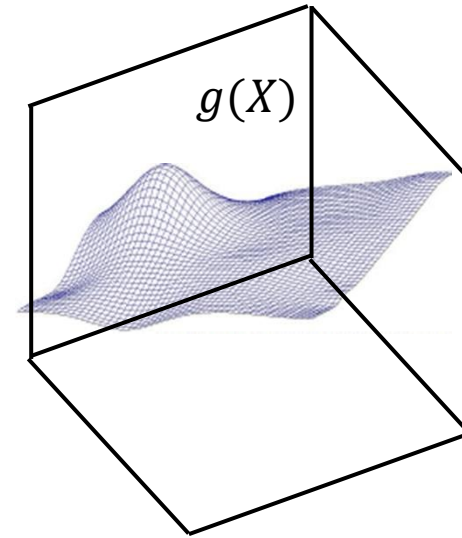
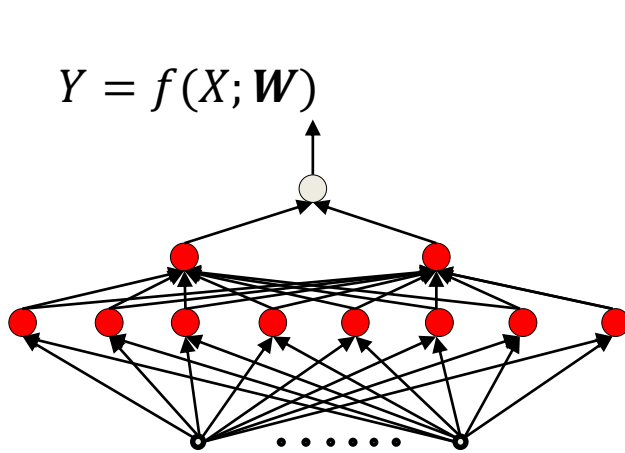


- When  $f(X; \mathbf{W})$  has the capacity to exactly represent  $g(X)$

$$\widehat{\mathbf{W}} = \operatorname{argmin}_{\mathbf{W}} \int_X \operatorname{div}(f(X; \mathbf{W}), g(X)) dX$$

- $\operatorname{div}()$  is a divergence function that goes to zero when  $f(X; \mathbf{W}) = g(X)$

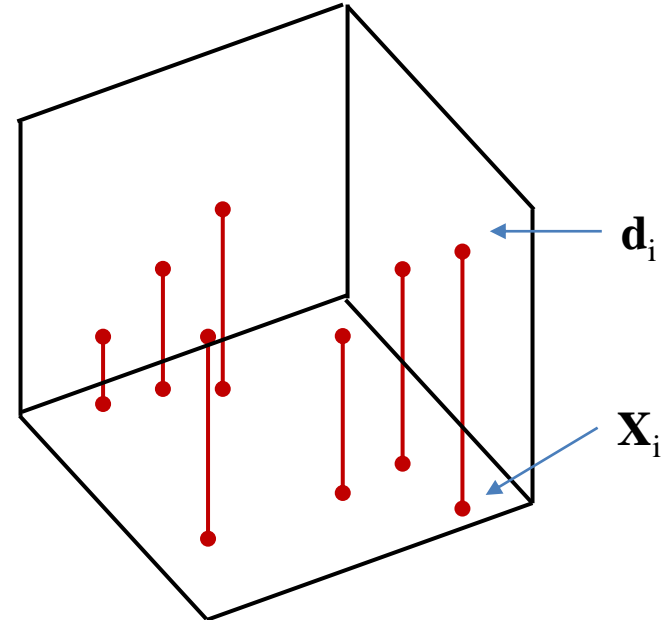
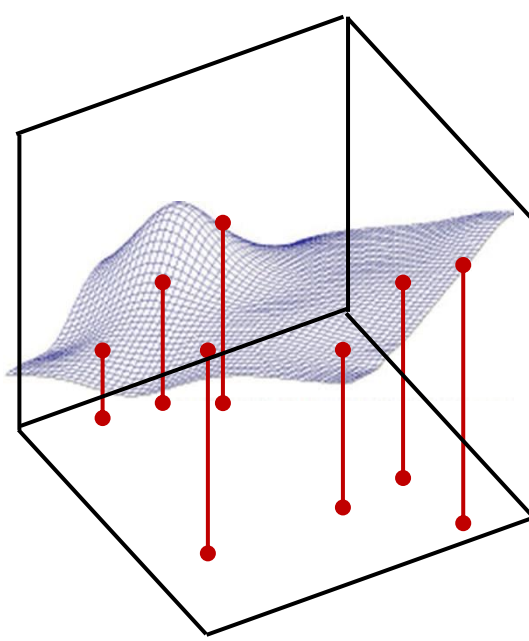
# Minimizing *expected* error



- More generally, assuming  $X$  is a random variable

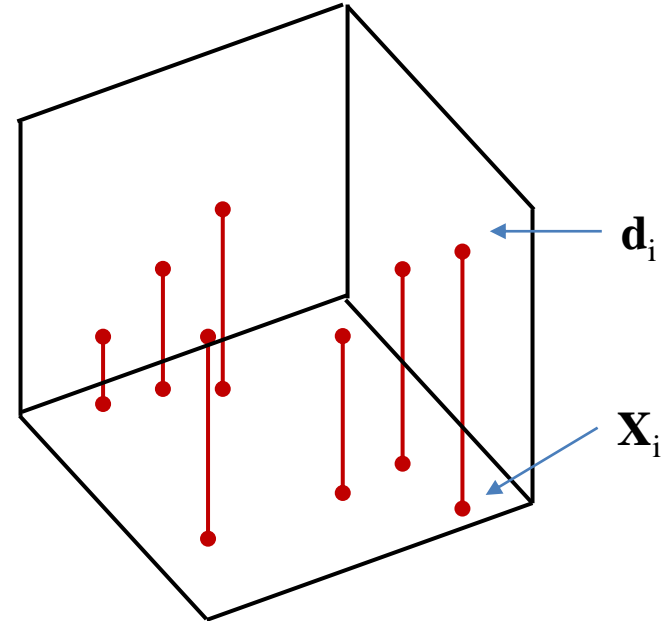
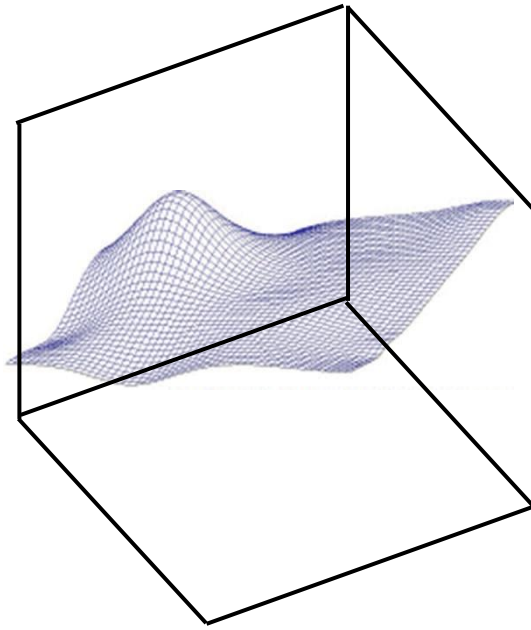
$$\begin{aligned}\widehat{\mathbf{W}} &= \operatorname{argmin}_{\mathbf{W}} \int_{\mathbf{X}} \operatorname{div}(f(\mathbf{X}; \mathbf{W}), g(\mathbf{X})) P(\mathbf{X}) d\mathbf{X} \\ &= \operatorname{argmin}_{\mathbf{W}} E[\operatorname{div}(f(\mathbf{X}; \mathbf{W}), g(\mathbf{X}))]\end{aligned}$$

# Recap: Sampling the function



- *We don't have  $g(X)$  so sample  $g(X)$* 
  - Obtain input-output pairs for a number of samples of input  $X_i$
  - Good sampling: the samples of  $X$  will be drawn from  $P(X)$
- Estimate function from the samples

# The *Empirical* risk



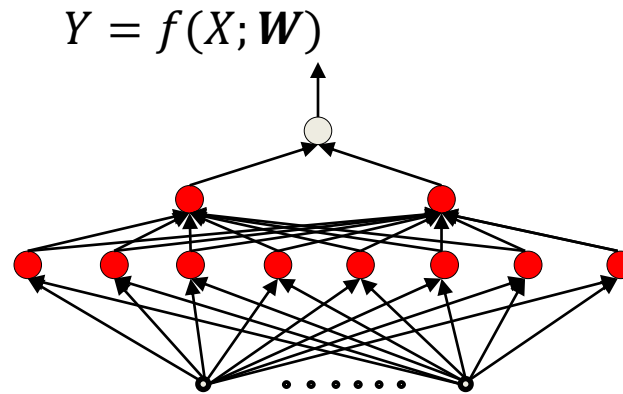
- The *expected* divergence (or risk) is the average divergence over the entire input space

$$E[\text{div}(f(X; W), g(X))] = \int_X \text{div}(f(X; W), g(X)) P(X) dX$$

- The *empirical estimate* of the expected risk is the *average* divergence over the samples

$$E[\text{div}(f(X; W), g(X))] \approx \frac{1}{N} \sum_{i=1}^N \text{div}(f(X_i; W), d_i)$$

# Empirical Risk Minimization



- Given a training set of input-output pairs  $(X_1, d_1), (X_2, d_2), \dots, (X_N, d_N)$ 
  - Quantification of error on the  $i^{\text{th}}$  instance:  $\text{div}(f(X_i; W), d_i)$
  - Empirical average divergence (Empirical Risk) on all training data:

$$\text{Loss}(W) = \frac{1}{N} \sum_i \text{div}(f(X_i; W), d_i)$$

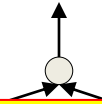
- Estimate the parameters to minimize the empirical estimate of expected divergence (empirical risk)

$$\widehat{W} = \underset{W}{\operatorname{argmin}} \text{Loss}(W)$$

- I.e. minimize the *empirical risk* over the drawn samples

# Empirical Risk Minimization

$$Y = f(X; W)$$



Note : Its really a measure of error, but using standard terminology, we will call it a "Loss"

Note 2: The empirical risk  $Loss(W)$  is only an empirical approximation to the true risk  $E[div(f(X; W), g(X))]$  which is our *actual* minimization objective

Note 3: For a given training set the loss is only a function of  $W$

$$Loss(W) = \frac{1}{N} \sum_i div(f(X_i; W), d_i)$$

- Estimate the parameters to minimize the empirical estimate of expected error

$$\hat{W} = \underset{W}{\operatorname{argmin}} Loss(W)$$

- I.e. minimize the *empirical error* over the drawn samples

# Problem Statement

- Given a training set of input-output pairs  $(X_1, d_1), (X_2, d_2), \dots, (X_N, d_N)$

- Minimize the following function

$$Loss(W) = \frac{1}{N} \sum_i div(f(X_i; W), d_i)$$

w.r.t  $W$

- This is problem of function minimization
  - An instance of optimization

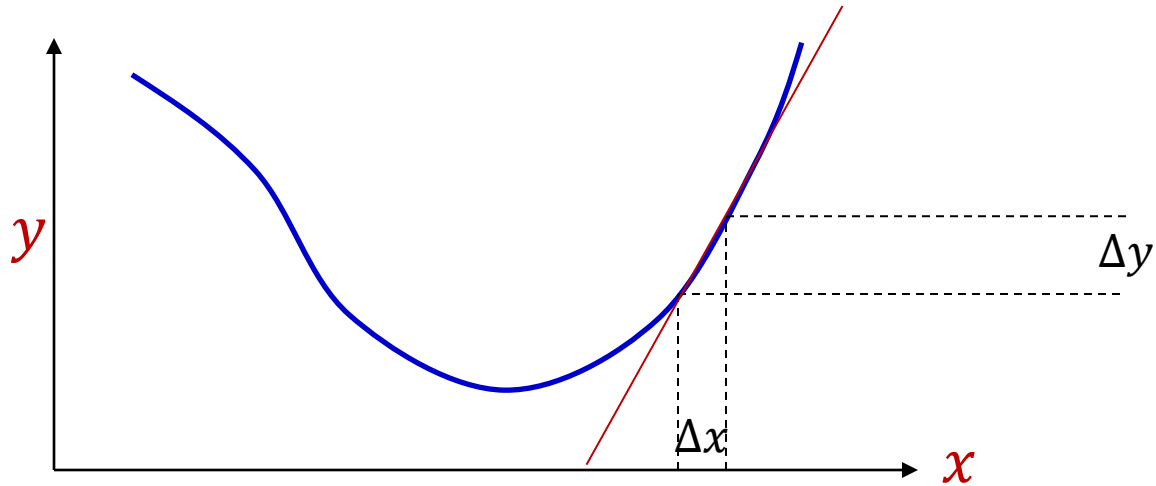
# Story so far

- We learn networks by “fitting” them to training instances drawn from a target function
- Learning networks of threshold-activation perceptrons requires solving a hard combinatorial-optimization problem
  - Because we cannot compute the influence of small changes to the parameters on the overall error
- Instead we use continuous activation functions with non-zero derivatives to enables us to estimate network parameters
  - This makes the output of the network differentiable w.r.t every parameter in the network
  - The *logistic* activation perceptron actually computes the *a posteriori* probability of the output given the input
- We define differentiable *divergence* between the output of the network and the desired output for the training instances
  - And a total error, which is the average divergence over all training instances
- We optimize network parameters to minimize this error
  - Empirical risk minimization
- This is an instance of function minimization



- **A CRASH COURSE ON FUNCTION OPTIMIZATION**
  - With an initial discussion of derivatives

# A brief note on derivatives..

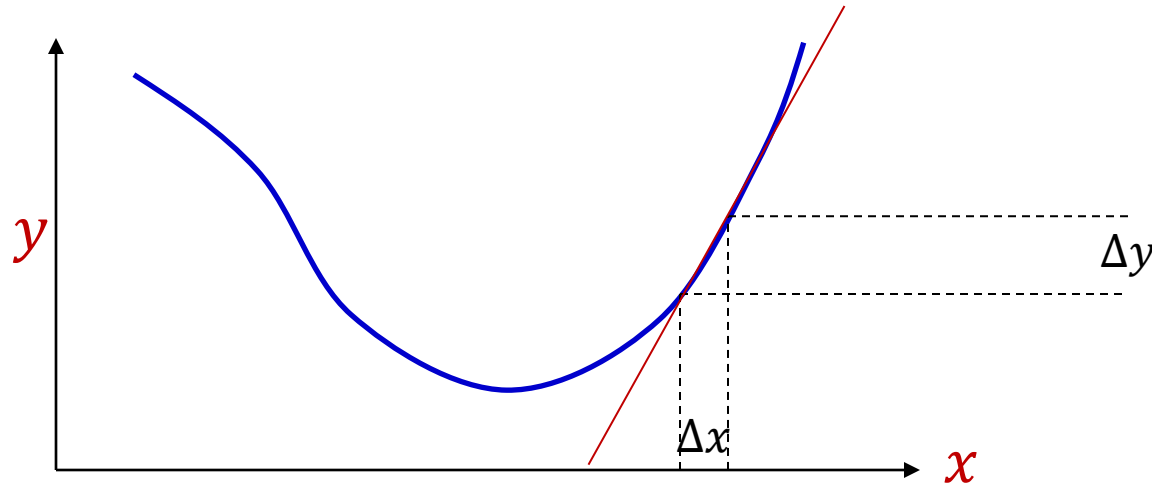


- A derivative of a function at any point tells us how much a minute increment to the *argument* of the function will increment the *value* of the function
  - For any  $y = f(x)$ , expressed as a multiplier  $\alpha$  to a tiny increment  $\Delta x$  to obtain the increments  $\Delta y$  to the output

$$\Delta y = \alpha \Delta x$$

- Based on the fact that at a fine enough resolution, any smooth, continuous function is locally linear at any point

# Scalar function of scalar argument



- When  $x$  and  $y$  are scalar

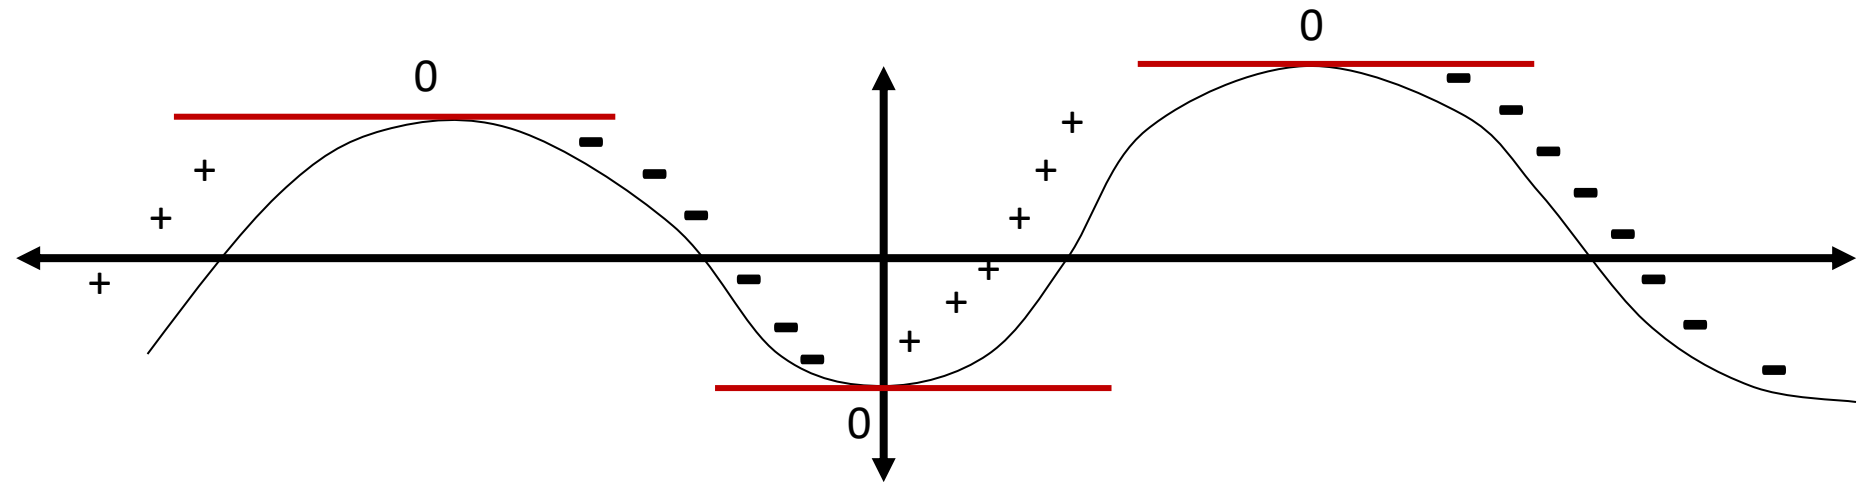
$$y = f(x)$$

- Derivative:

$$\Delta y = \alpha \Delta x$$

- Often represented (using somewhat inaccurate notation) as  $\frac{dy}{dx}$
- Or alternately (and more reasonably) as  $f'(x)$

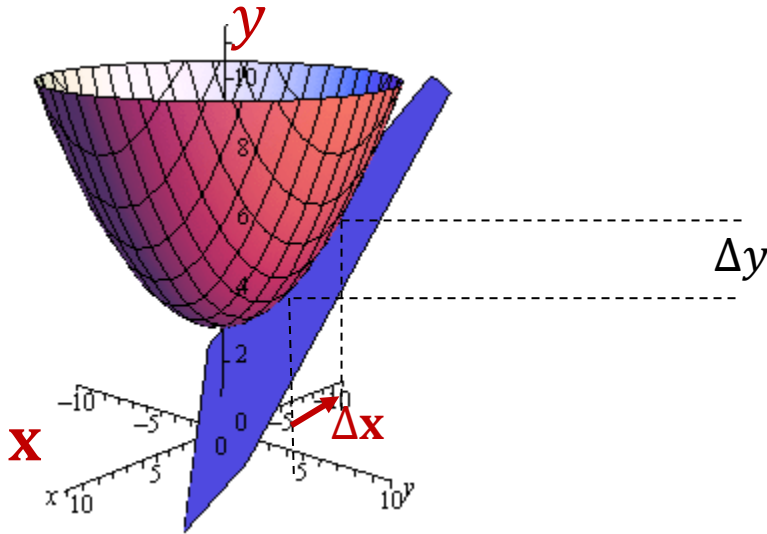
# Scalar function of scalar argument



- Derivative  $f'(x)$  is the *rate of change* of the function at  $x$ 
  - How fast it increases with increasing  $x$
  - The magnitude of  $f'(x)$  gives you the steepness of the curve at  $x$ 
    - Larger  $|f'(x)| \rightarrow$  the function is increasing or decreasing more rapidly
- It will be positive where a small increase in  $x$  results in an *increase* of  $f(x)$ 
  - Regions of positive slope
- It will be negative where a small increase in  $x$  results in a *decrease* of  $f(x)$ 
  - Regions of negative slope
- It will be 0 where the function is locally flat (neither increasing nor decreasing)

# Multivariate scalar function:

## Scalar function of *vector* argument



Note:  $\Delta \mathbf{x}$  is now a vector

$$\Delta \mathbf{x} = \begin{bmatrix} \Delta x_1 \\ \vdots \\ \Delta x_D \end{bmatrix}$$

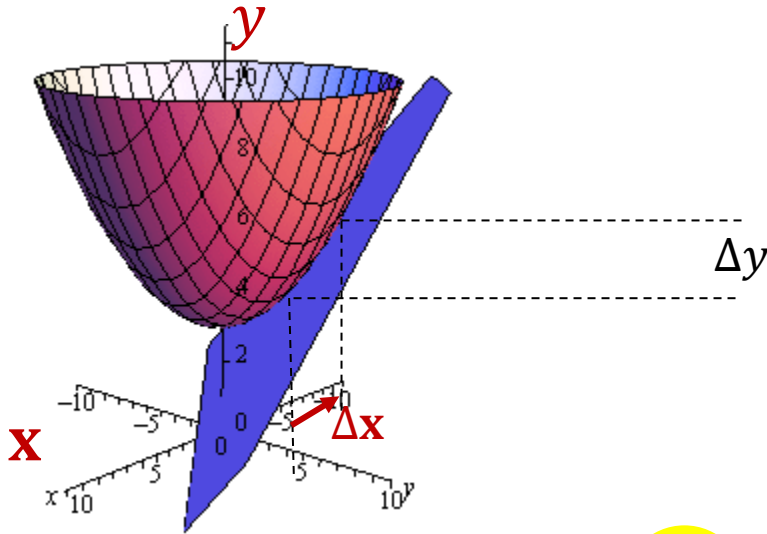
$$\Delta y = \alpha \Delta \mathbf{x}$$

- Giving us that  $\alpha$  is a row vector:  $\alpha = [\alpha_1 \quad \cdots \quad \alpha_D]$ 
$$\Delta y = \alpha_1 \Delta x_1 + \alpha_2 \Delta x_2 + \cdots + \alpha_D \Delta x_D$$
- The *partial* derivative  $\alpha_i$  gives us how  $y$  increments when *only*  $x_i$  is incremented
- Often represented as  $\frac{\partial y}{\partial x_i}$

$$\Delta y = \frac{\partial y}{\partial x_1} \Delta x_1 + \frac{\partial y}{\partial x_2} \Delta x_2 + \cdots + \frac{\partial y}{\partial x_D} \Delta x_D$$

# Multivariate scalar function:

## Scalar function of *vector* argument



Note:  $\Delta \mathbf{x}$  is now a vector

$$\Delta \mathbf{x} = \begin{bmatrix} \Delta x_1 \\ \vdots \\ \Delta x_D \end{bmatrix}$$

$$\Delta y = \nabla_{\mathbf{x}} y \Delta \mathbf{x}$$

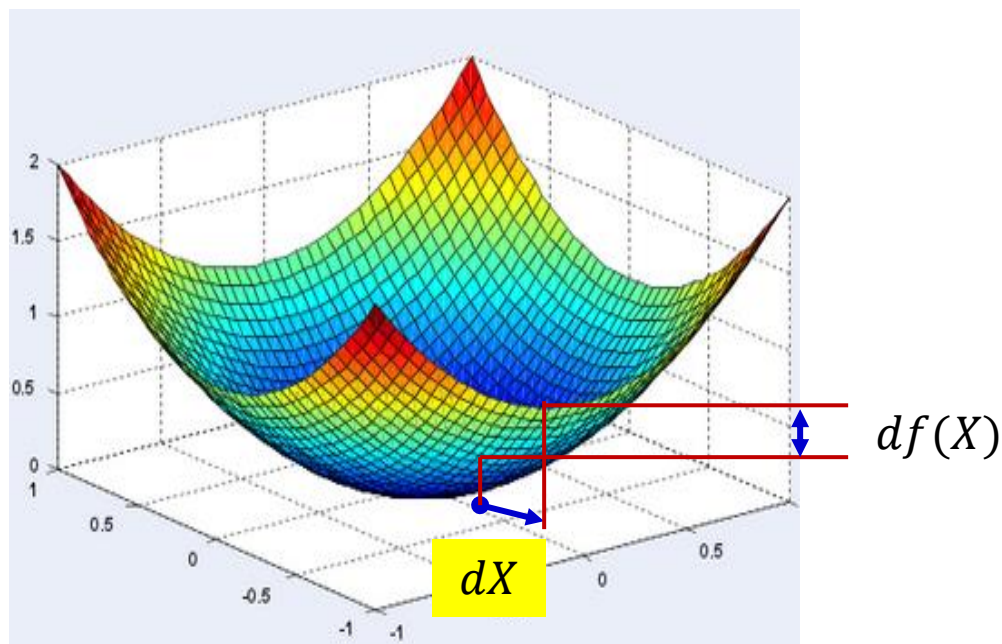
- Where

$$\nabla_{\mathbf{x}} y = \begin{bmatrix} \frac{\partial y}{\partial x_1} & \dots & \frac{\partial y}{\partial x_D} \end{bmatrix}$$

We will be using this symbol for vector and matrix derivatives

- You may be more familiar with the term “gradient” which is actually defined as the transpose of the derivative

# Gradient of a scalar function of a vector

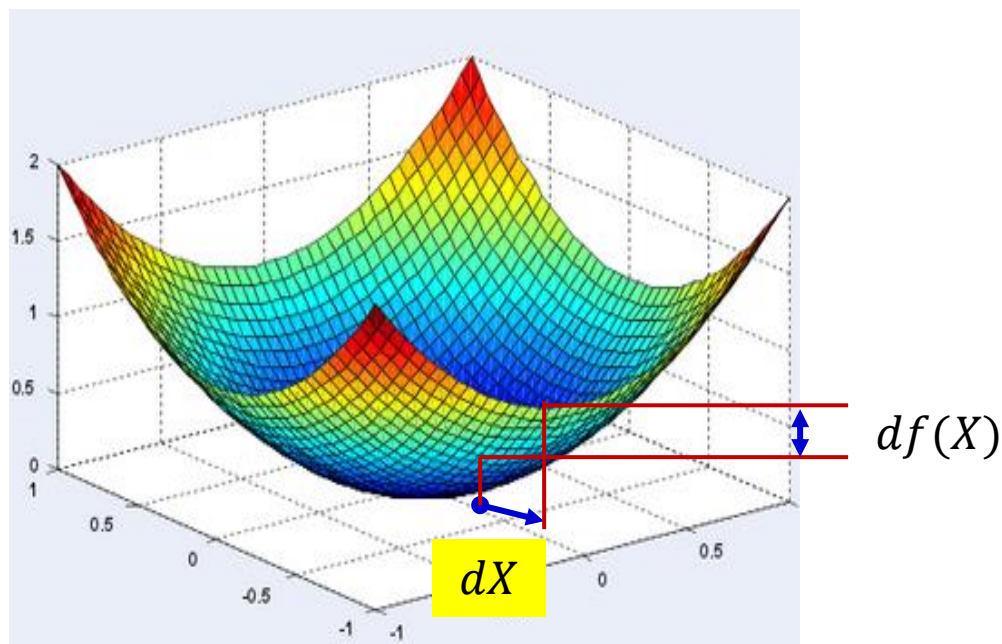


- The *derivative*  $\nabla_X f(X)$  of a scalar function  $f(X)$  of a multi-variate input  $X$  is a multiplicative factor that gives us the change in  $f(X)$  for tiny variations in  $X$

$$df(X) = \nabla_X f(X) dX$$

- $\nabla_X f(X) = \begin{bmatrix} \frac{\partial f(X)}{\partial x_1} & \frac{\partial f(X)}{\partial x_2} & \dots & \frac{\partial f(X)}{\partial x_n} \end{bmatrix}$
- The **gradient** is the transpose of the derivative  $\nabla_X f(X)^T$ 
  - A column vector of the same dimensionality as  $X$

# Gradient of a scalar function of a vector



- The *derivative*  $\nabla_X f(X)$  of a scalar function  $f(X)$  of a multi-variate input  $X$  is a multiplicative factor that gives us the change in  $f(X)$  for tiny variations in  $X$

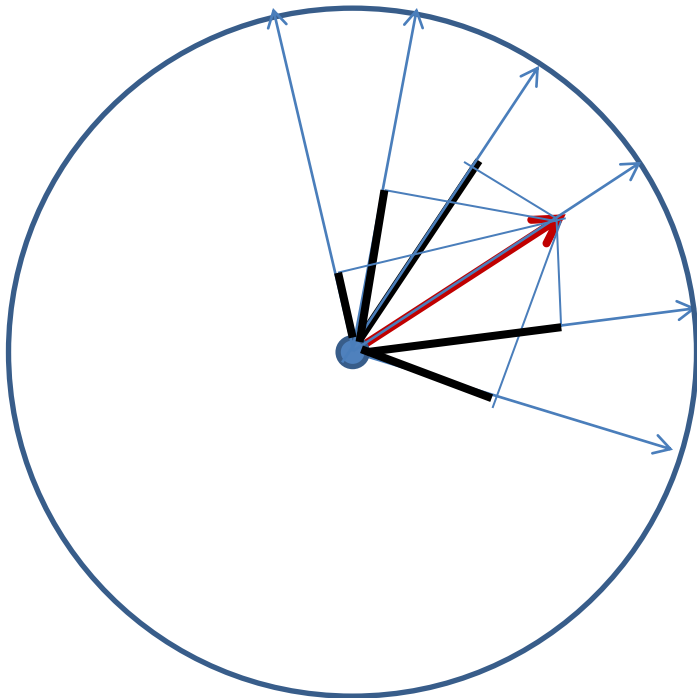
$$df(X) = \nabla_X f(X) dX$$

$$- \nabla_X f(X) = \left[ \frac{\partial f(X)}{\partial x_1} \quad \frac{\partial f(X)}{\partial x_2} \quad \dots \quad \frac{\partial f(X)}{\partial x_n} \right]$$

This is a vector inner product. To understand its behavior let's consider a well-known property of inner products



# A well-known vector property

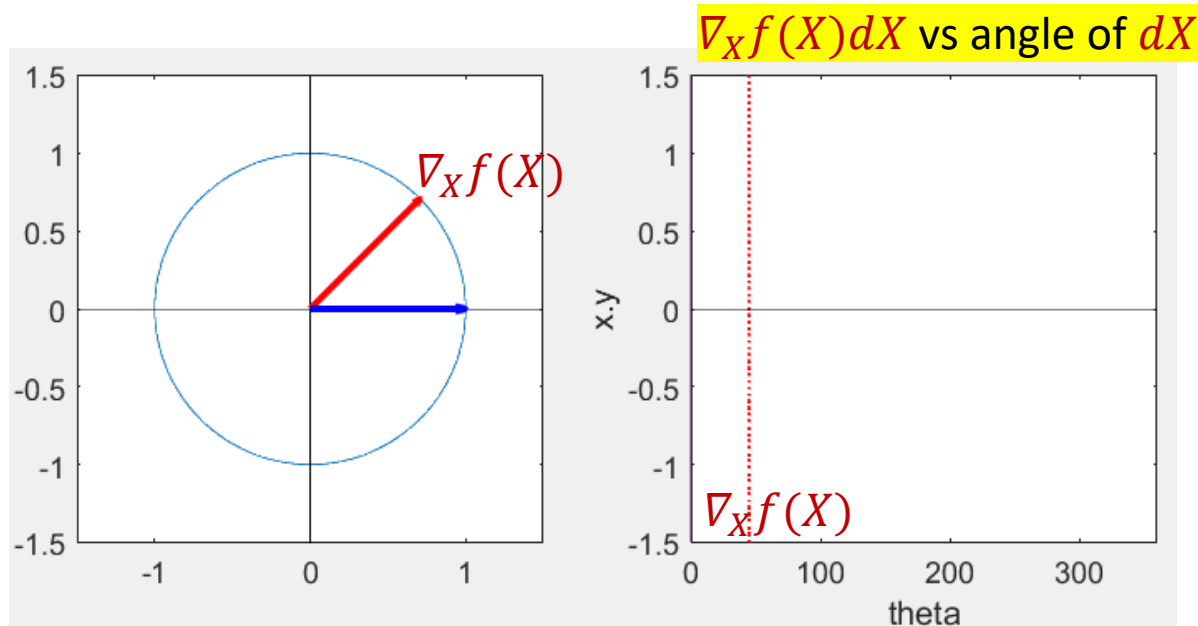


$$\mathbf{u}^T \mathbf{v} = |\mathbf{u}| |\mathbf{v}| \cos \theta$$

- The inner product between two vectors of fixed lengths is maximum when the two vectors are aligned
  - i.e. when  $\theta = 0$

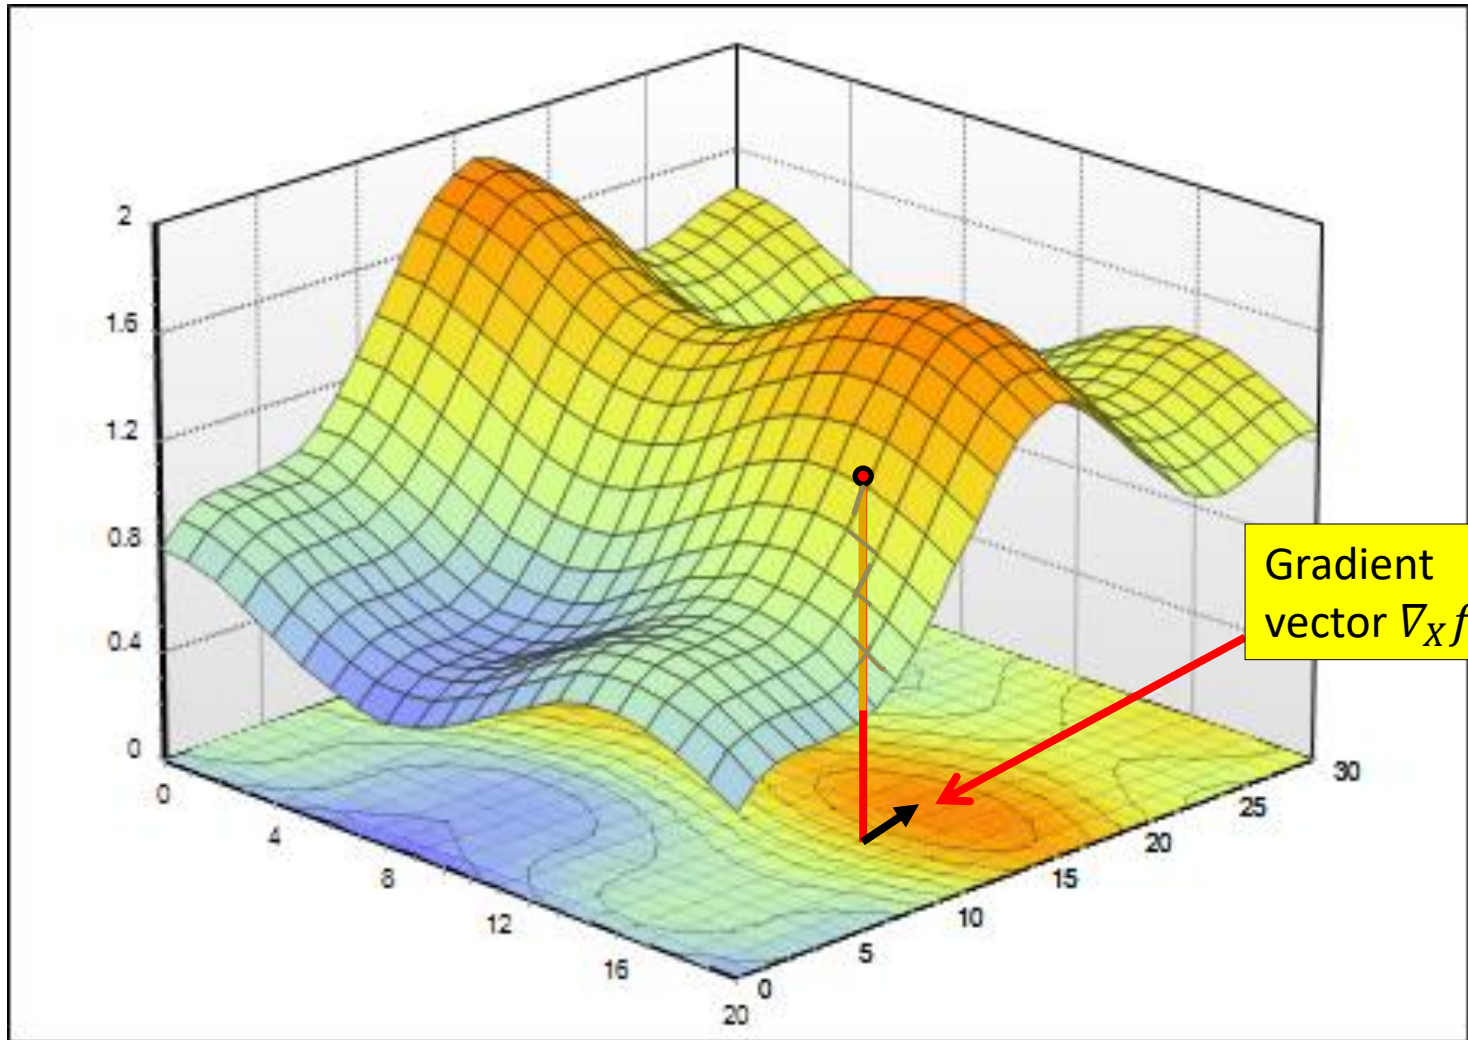
# Properties of Gradient

Blue arrow  
is  $dX$

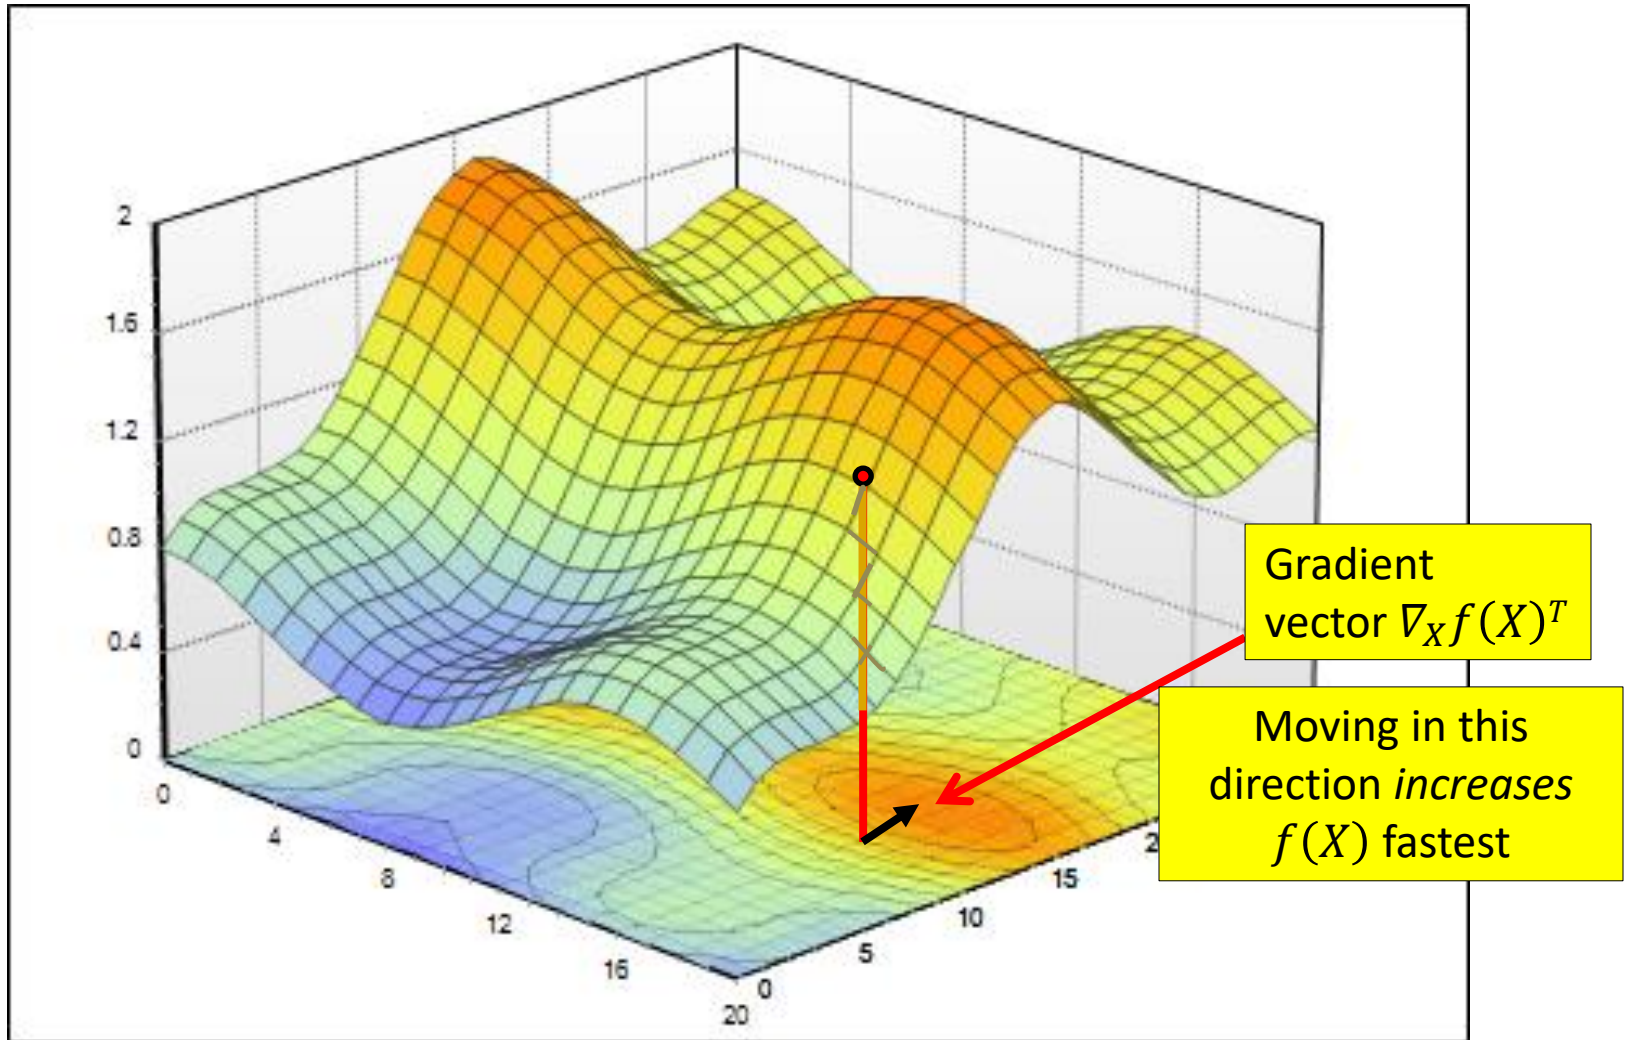


- $df(X) = \nabla_X f(X)dX$
- For an increment  $dX$  of any given length  $df(X)$  is max if  $dX$  is aligned with  $\nabla_X f(X)^T$ 
  - The function  $f(X)$  increases most rapidly if the input increment  $dX$  is exactly in the direction of  $\nabla_X f(X)^T$
- The gradient is the direction of fastest increase in  $f(X)$

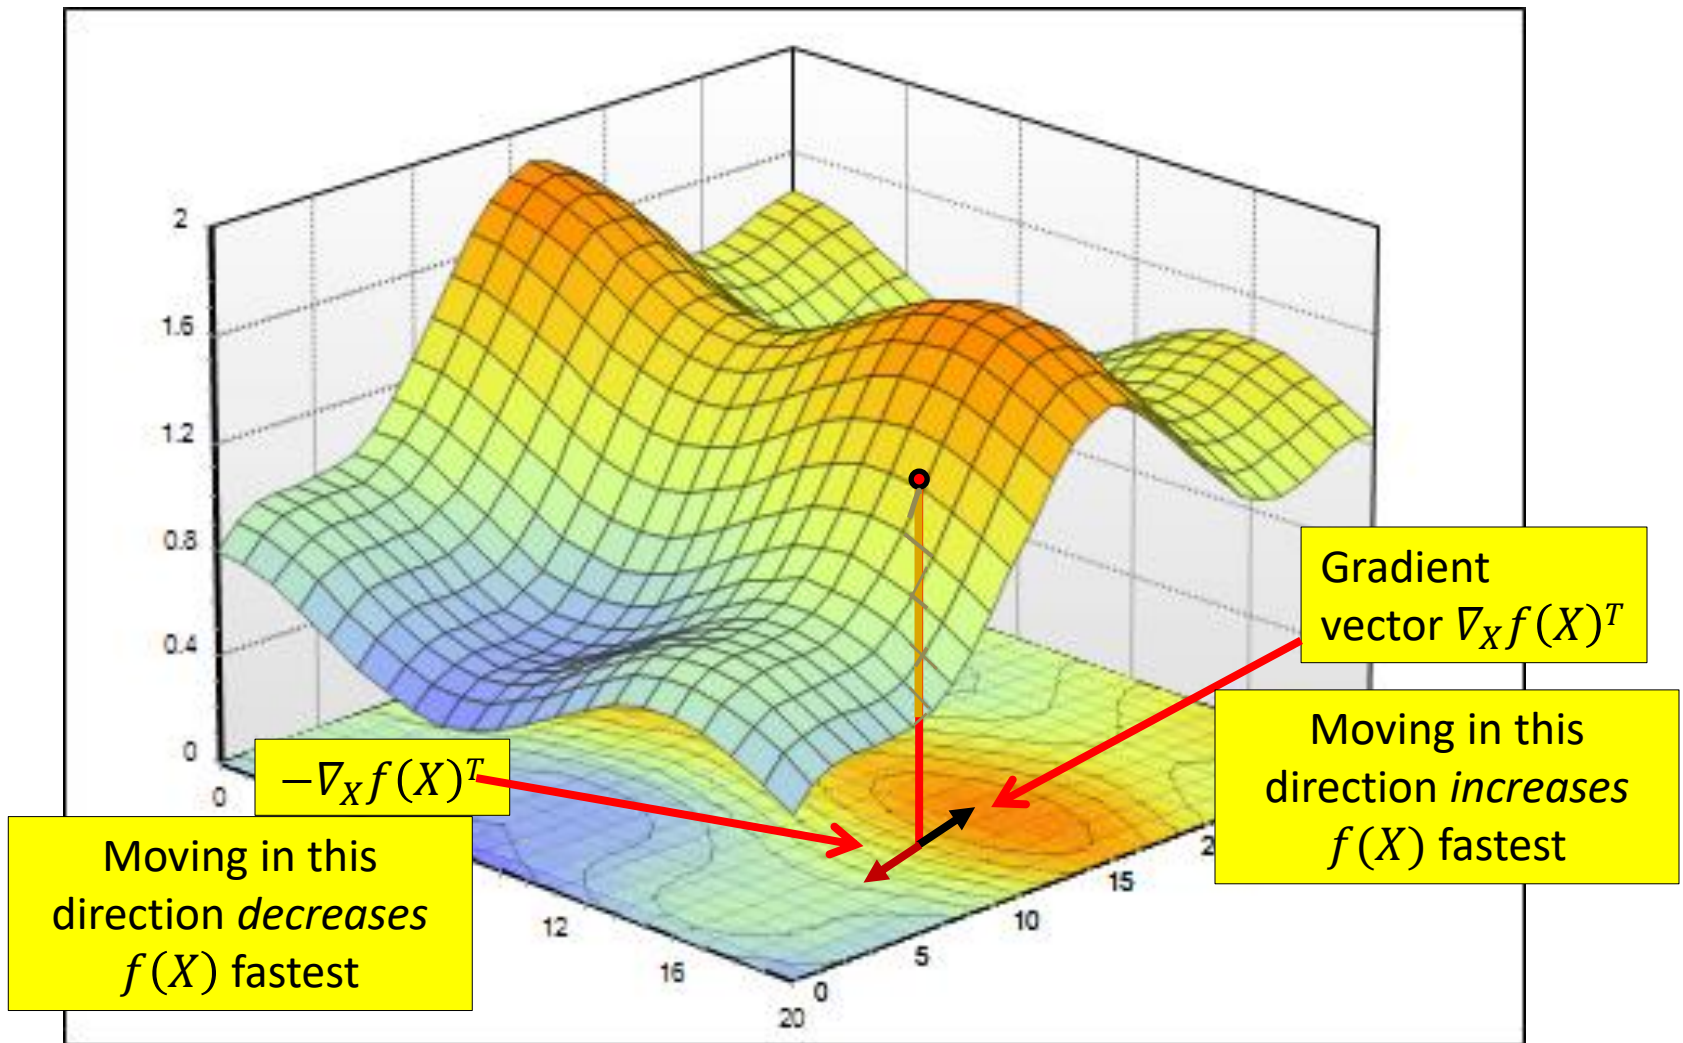
# Gradient



# Gradient

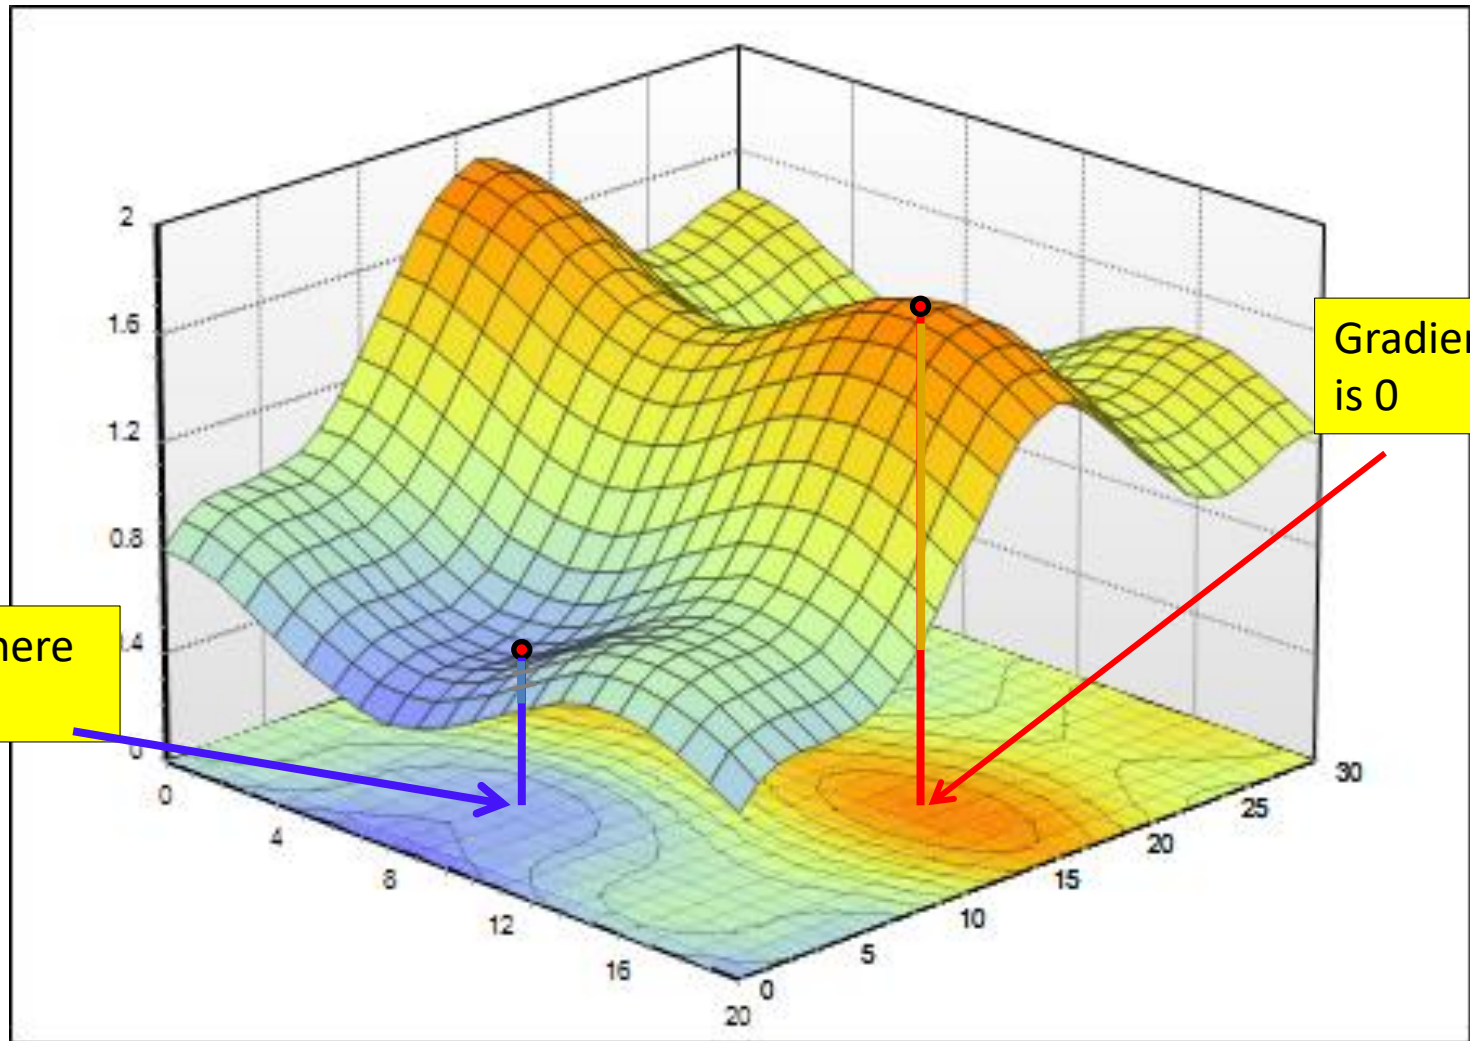


# Gradient





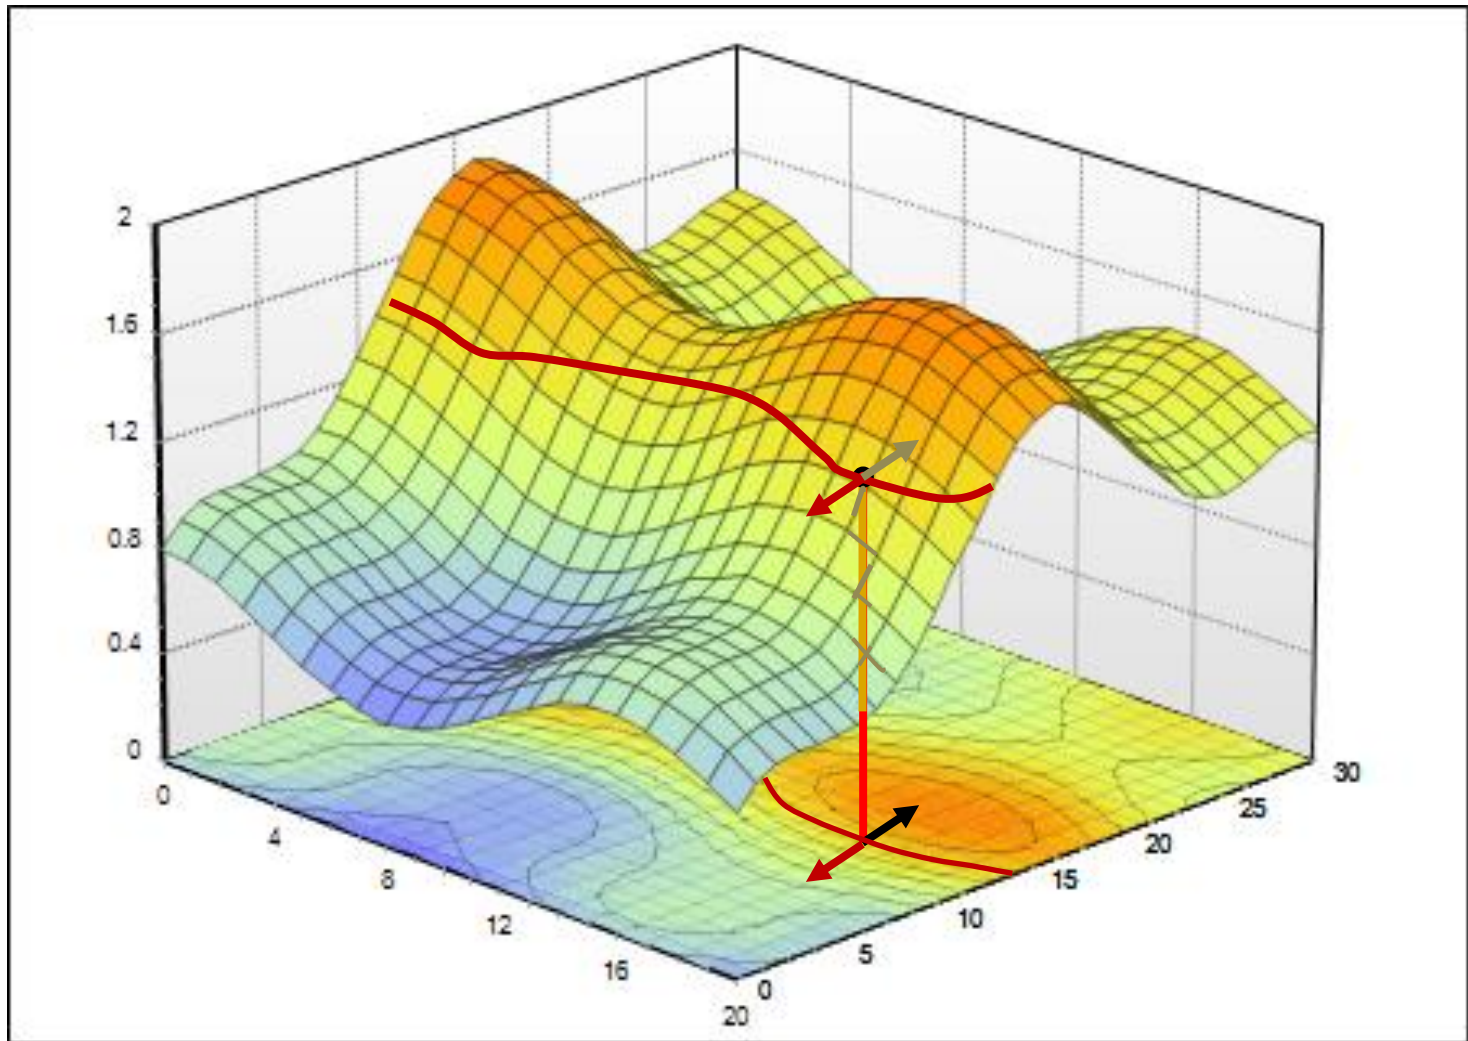
# Gradient



Gradient here  
is 0

Gradient here  
is 0

# Properties of Gradient: 2



- The gradient vector  $\nabla_x f(X)^T$  is perpendicular to the level curve

# The Hessian

- The Hessian of a function  $f(x_1, x_2, \dots, x_n)$  is given by the second derivative

$$\nabla_x^2 f(x_1, \dots, x_n) := \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdot & \cdot & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdot & \cdot & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdot & \cdot & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$



# Next up

- Continuing on function optimization
- Gradient descent to train neural networks
- A.K.A. Back propagation